

Deep Learning with Keras and R

Feb 15, 2018

RajivShah.com

rshah@pobox.com

 github.com/rajshah4/image_keras

 rajcs4

Hotdog!



Share

No Thanks

Not hotdog!



Share

No Thanks

decryptr



generate startup names

CALLER

A SIMPLE, AND INTUITIVE VISUAL INSTRUCTION, BASED ON 56+ PHOTO PROPERTIES.

DEARDHEAR

BUILDS WORKFLOW & COLLABORATION SOFTWARE FOR SMALL AND MID-MARKET BUSINESSES.

R8VUTURE HEALTH

PROVIDES NATIVE VIDEO ADVERTISING EXPERIENCE FOR KIDS, COUPLES AND ATTORNEYS.

STRICWITEA

ENABLES LOCAL GOVERNMENT TO ACCEPT AND CURATE DIGITAL CONTENT.

HIMY

IF YOUR STARTUP IS LOOKING FOR A SHARE OF THEIR TEMPORARY STAFFING NEEDS.

HYPART

WE'RE A B2B MARKETPLACE FOR INSURANCE SECURITIZATION.

LATT

TABLETS ON RESTAURANT TABLES, SO GUESTS CAN ORDER, PAY AND PLAY GAMES FROM THEIR SMARTPHONE IN LESS THAN 5 MINUTES.

AIYDLRMSS

FLOYD IS HEROKU FOR MACHINE LEARNING TO HELP YOU GET FEEDBACK AND ITERATE ON THEIR DESIGNS.

KIPPI

WE BUILD CROWDFUNDING SITES THAT LETS WOMEN RENT DRESSES FROM EACH OTHER.

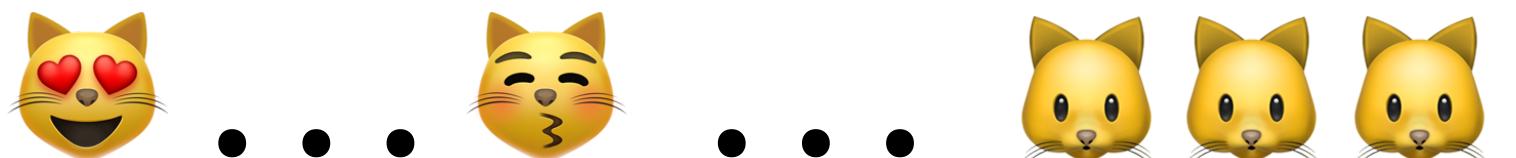
TATT

CCG EMPOWERS CLINICIANS TO MAKE PR SCALE IN THE REAL WORLD.

goals

History of Keras for R

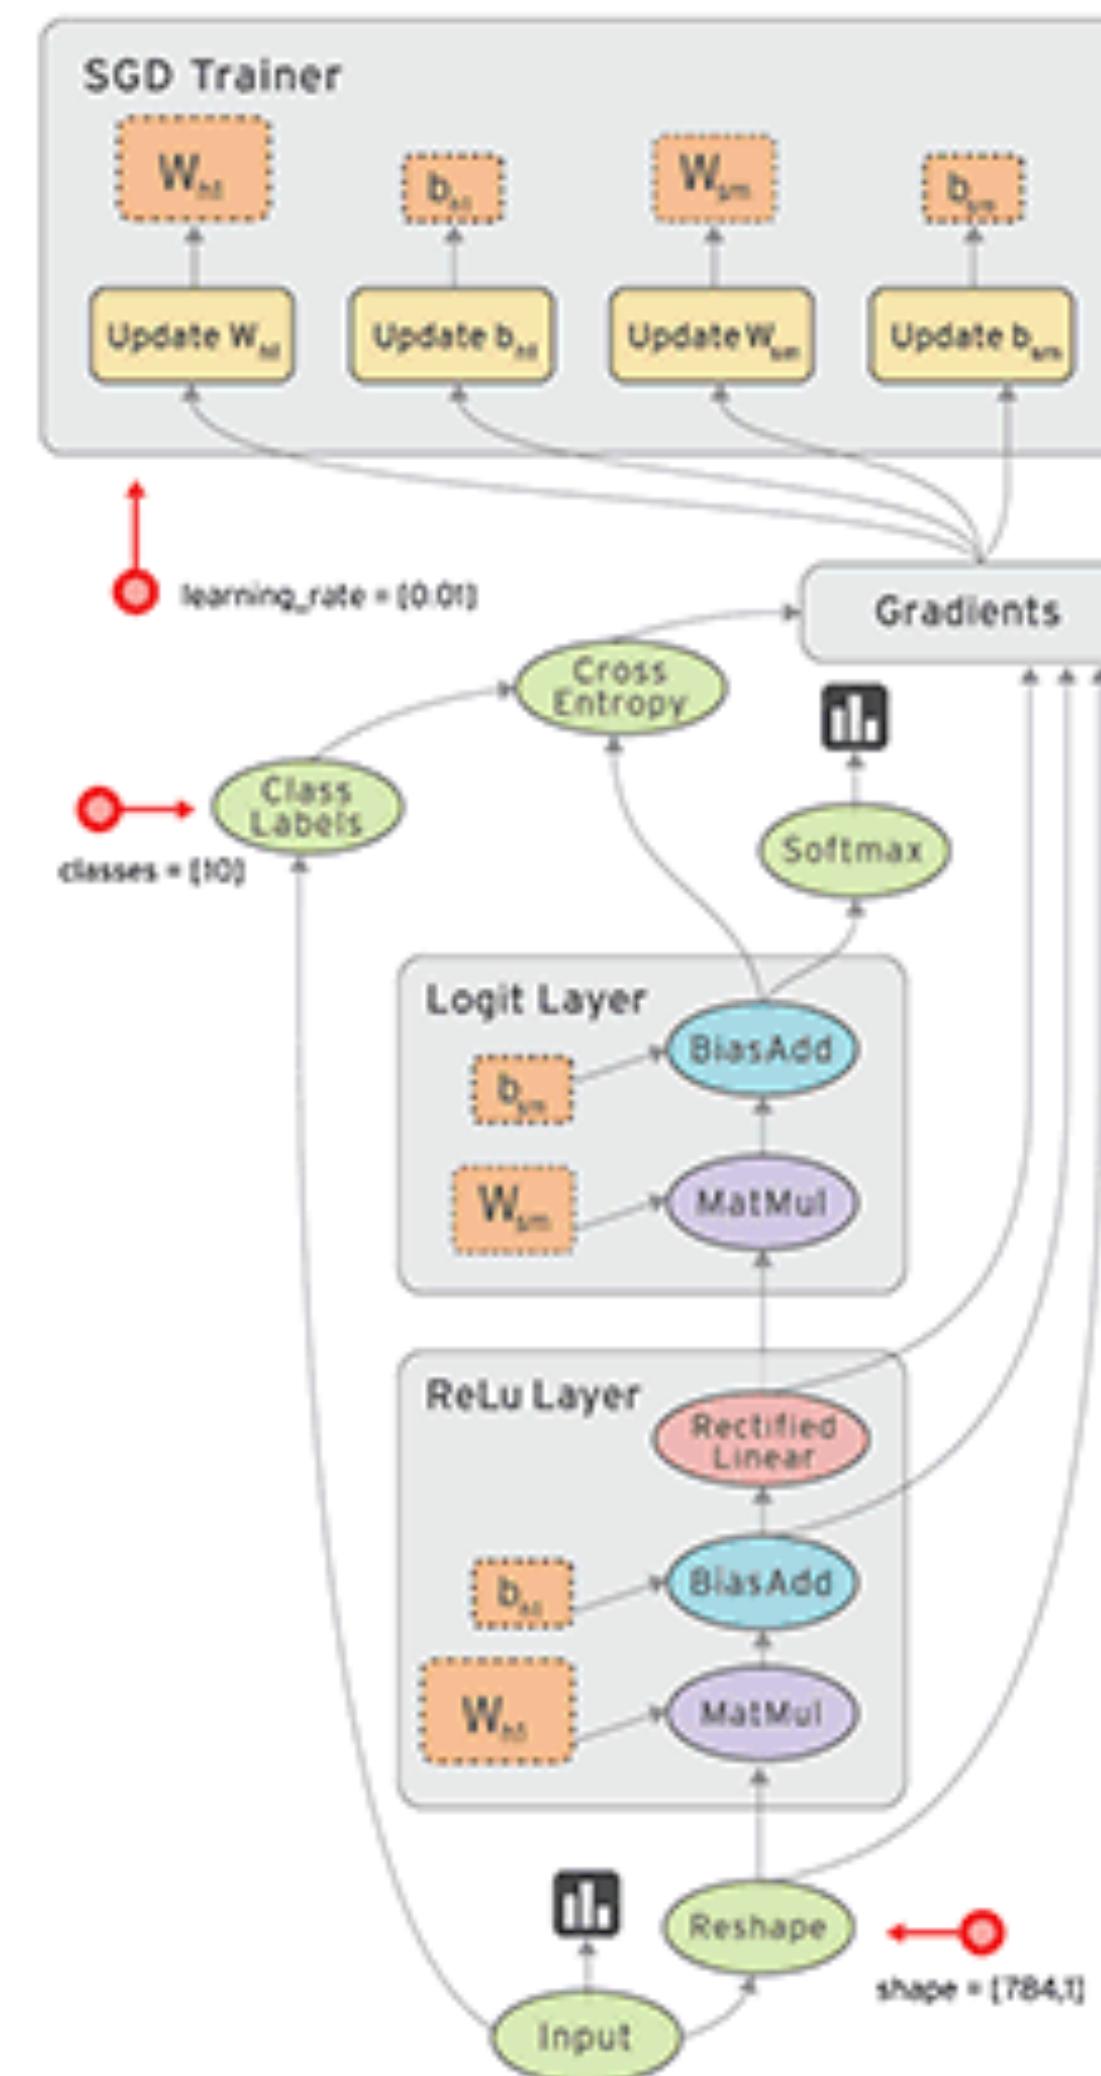
Use Keras for:
Telling  from 



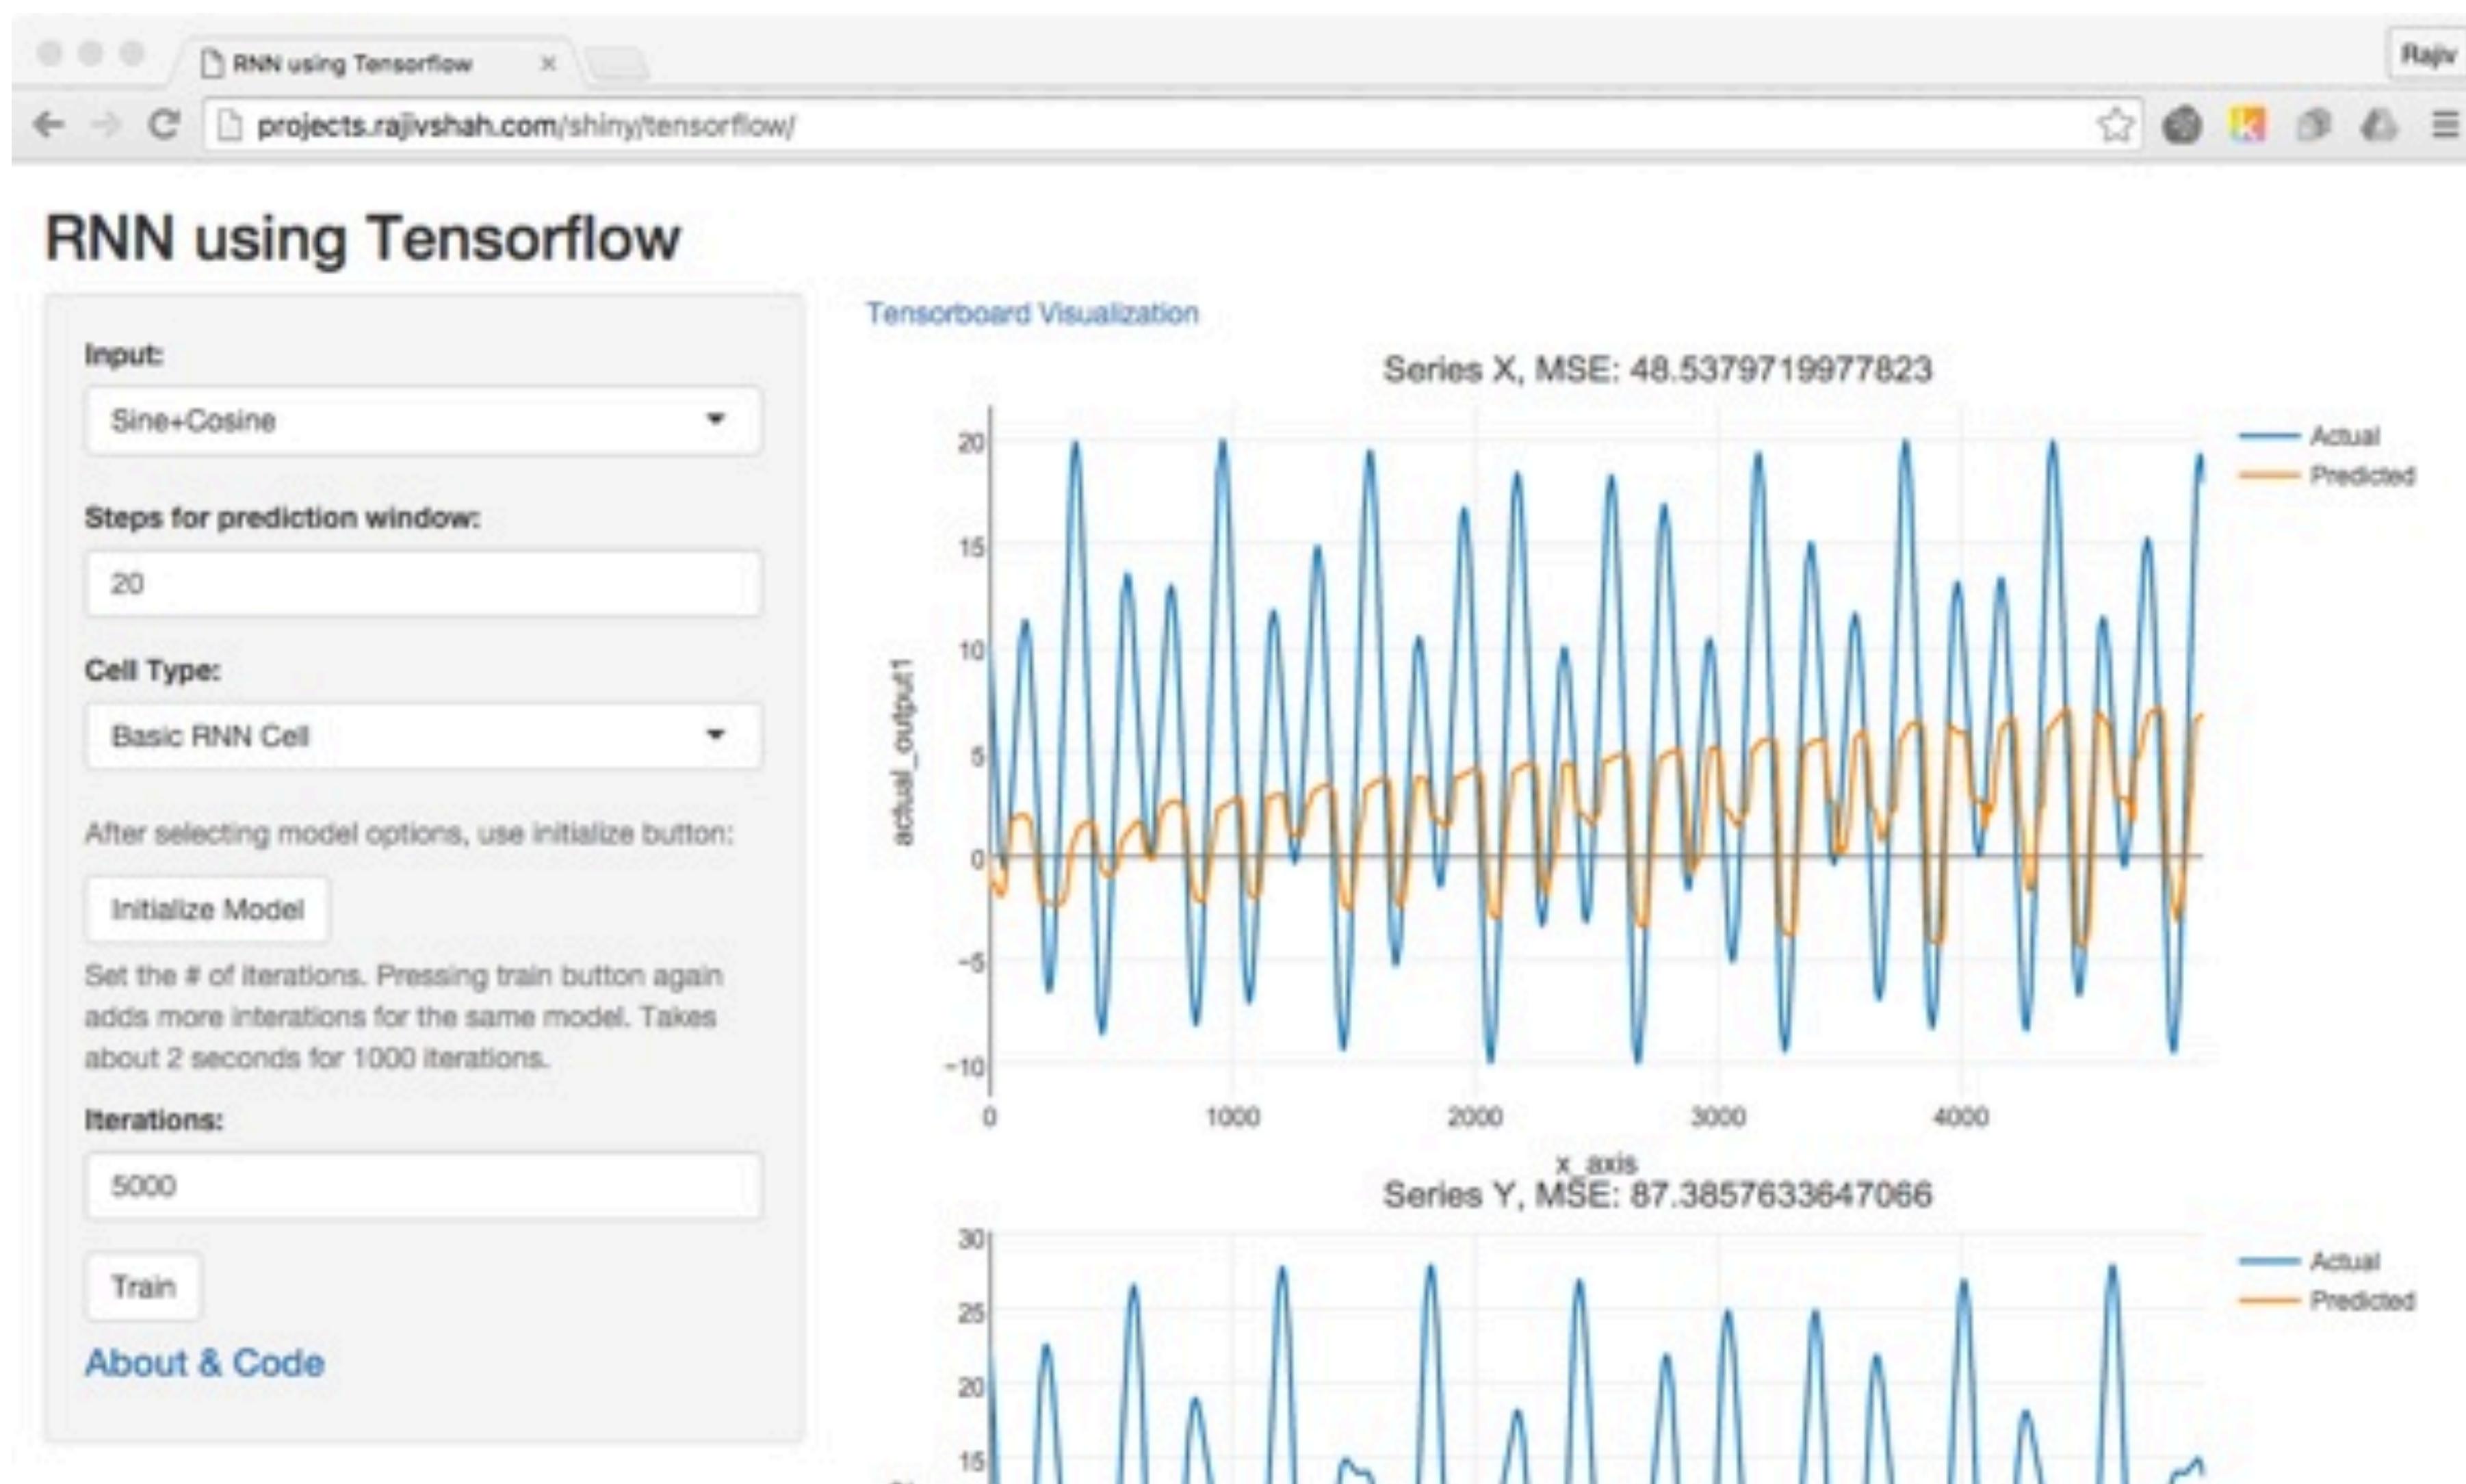
Generating new 

Python & R interfaces

Tensorflow



rPython (c. 2010)



rPython (c. 2010)

```
iw <- as.numeric(input$input_wave)-1
python.assign( "input_wave", iw )
python.assign( "prlag", input$prlag )
python.exec(
    lag = prlag
    def get_sample():
        global angle1, angle2
        angle1 += 2*pi/float(frequency1)
        angle2 += 2*pi/float(frequency2)
        angle1 %= 2*pi
        angle2 %= 2*pi
        return array([array([
            5 + 5*sin(angle1) + 10*cos(angle2)**input_wave,
            7 + 7*sin(angle2)**input_wave + 14*cos(angle1)])])
    )
```

Tensorflow for R

(Oct. 2016)

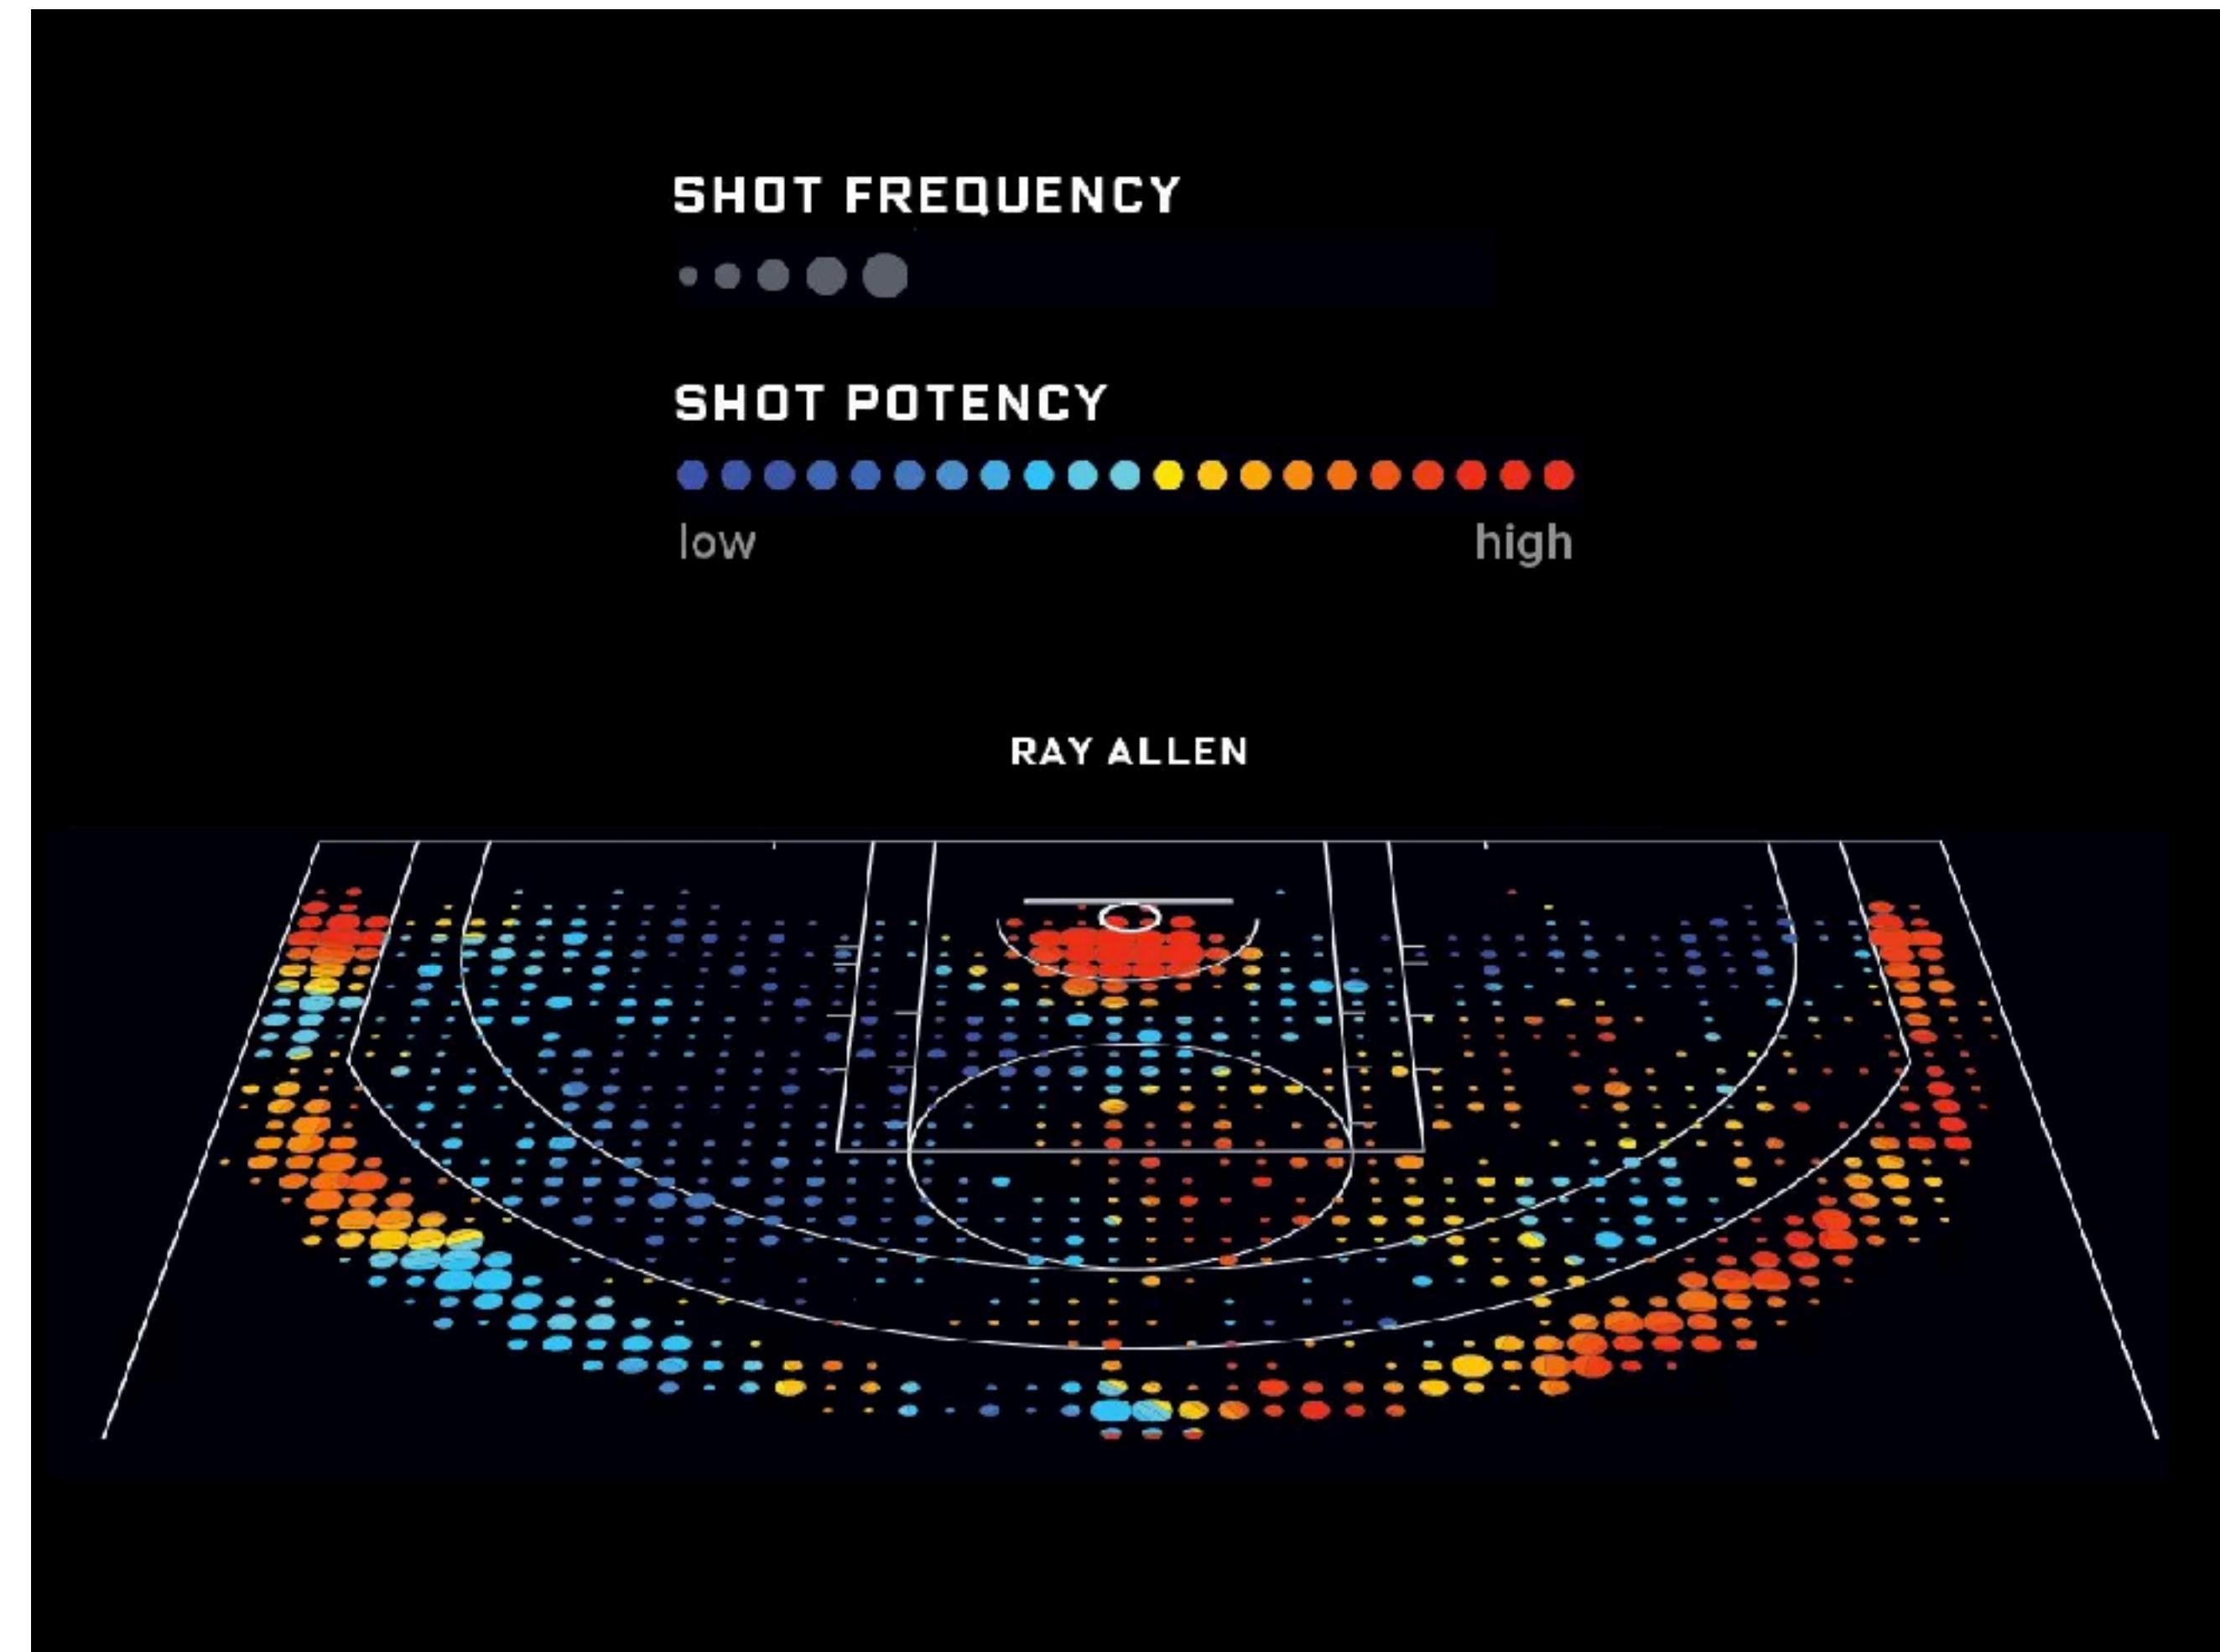
```
cross_entropy <- tf$reduce_mean(-tf$reduce_sum(y_ * tf$log(y_conv),  
reduction_indices=1L))  
  
train_step <- tf$train$AdamOptimizer(1e-4)$minimize(cross_entropy)  
  
correct_prediction <- tf$equal(tf$argmax(y_conv, 1L), tf$argmax(y_, 1L))  
  
accuracy <- tf$reduce_mean(tf$cast(correct_prediction, tf$float32))  
sess$run(tf$global_variables_initializer())
```

Reticulate (Feb. 2017)

The `reticulate` package provides a very clean & concise interface bridge between R and Python which makes it handy to work with modules that have yet to be ported to R (going native is always better when you can do it). This post shows how to use `reticulate` to create parquet files directly from R using `reticulate` as a bridge to the `pyarrow` module, which has the ability to natively create parquet files.

spacyr: an R wrapper for spaCy

Reticulate (Feb. 2017)



Reticulate (Feb. 2017)

```
library(reticulate)

reticulate::import('goldsberry') -> g
df <- g $ PlayerList(Season = '2016-17') $ players() %>%
bind_rows
```

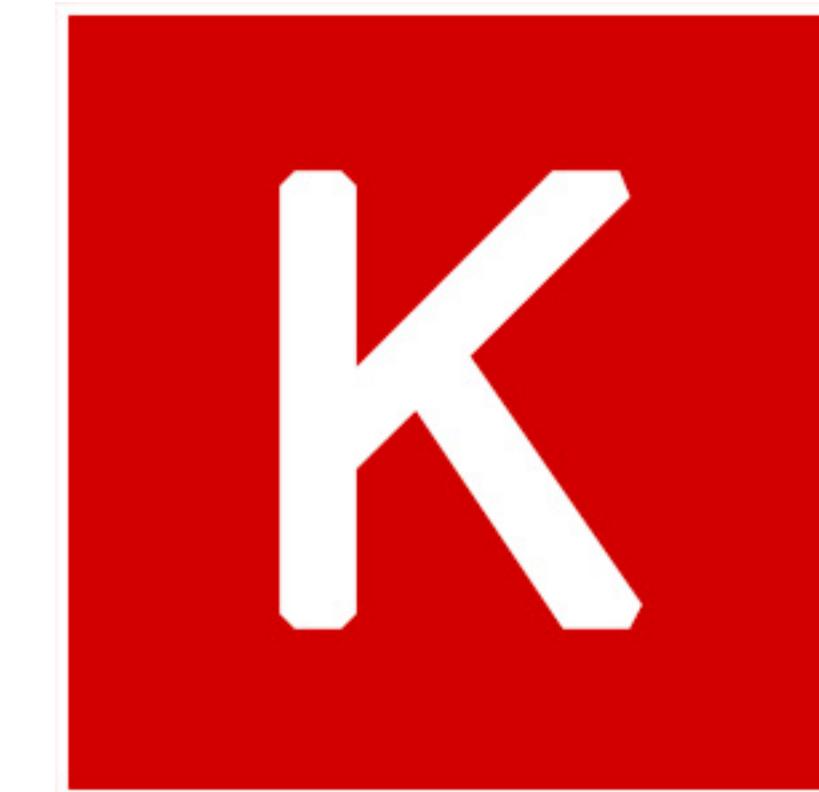


Room for both

Keras

High Level Neural Network API

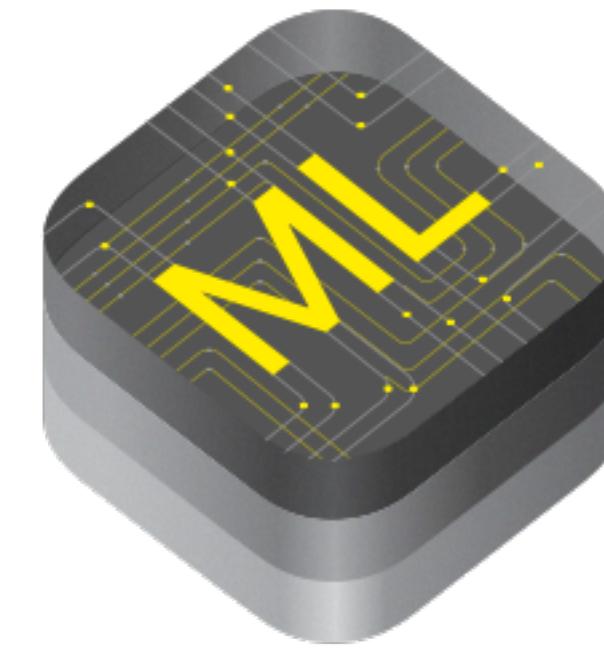
theano



Microsoft
CNTK



TensorFlow



mxnet

the book



Deep Learning with **R**

François Chollet
with J. J. Allaire

MANNING



faq: keras/r



Install: library(keras)

Works with GPU

**You can set locations of python/tensorflow
– extensive instructions for installing/
configuring tensorflow**

Approximately the same speed as python

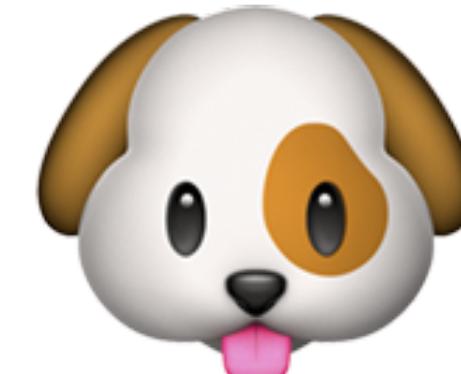
Awesome work by J.J. Allaire

Telling  **from** 

Telling from



In 2013 - 82.7%
CNN - 99%

goals:  **from** 

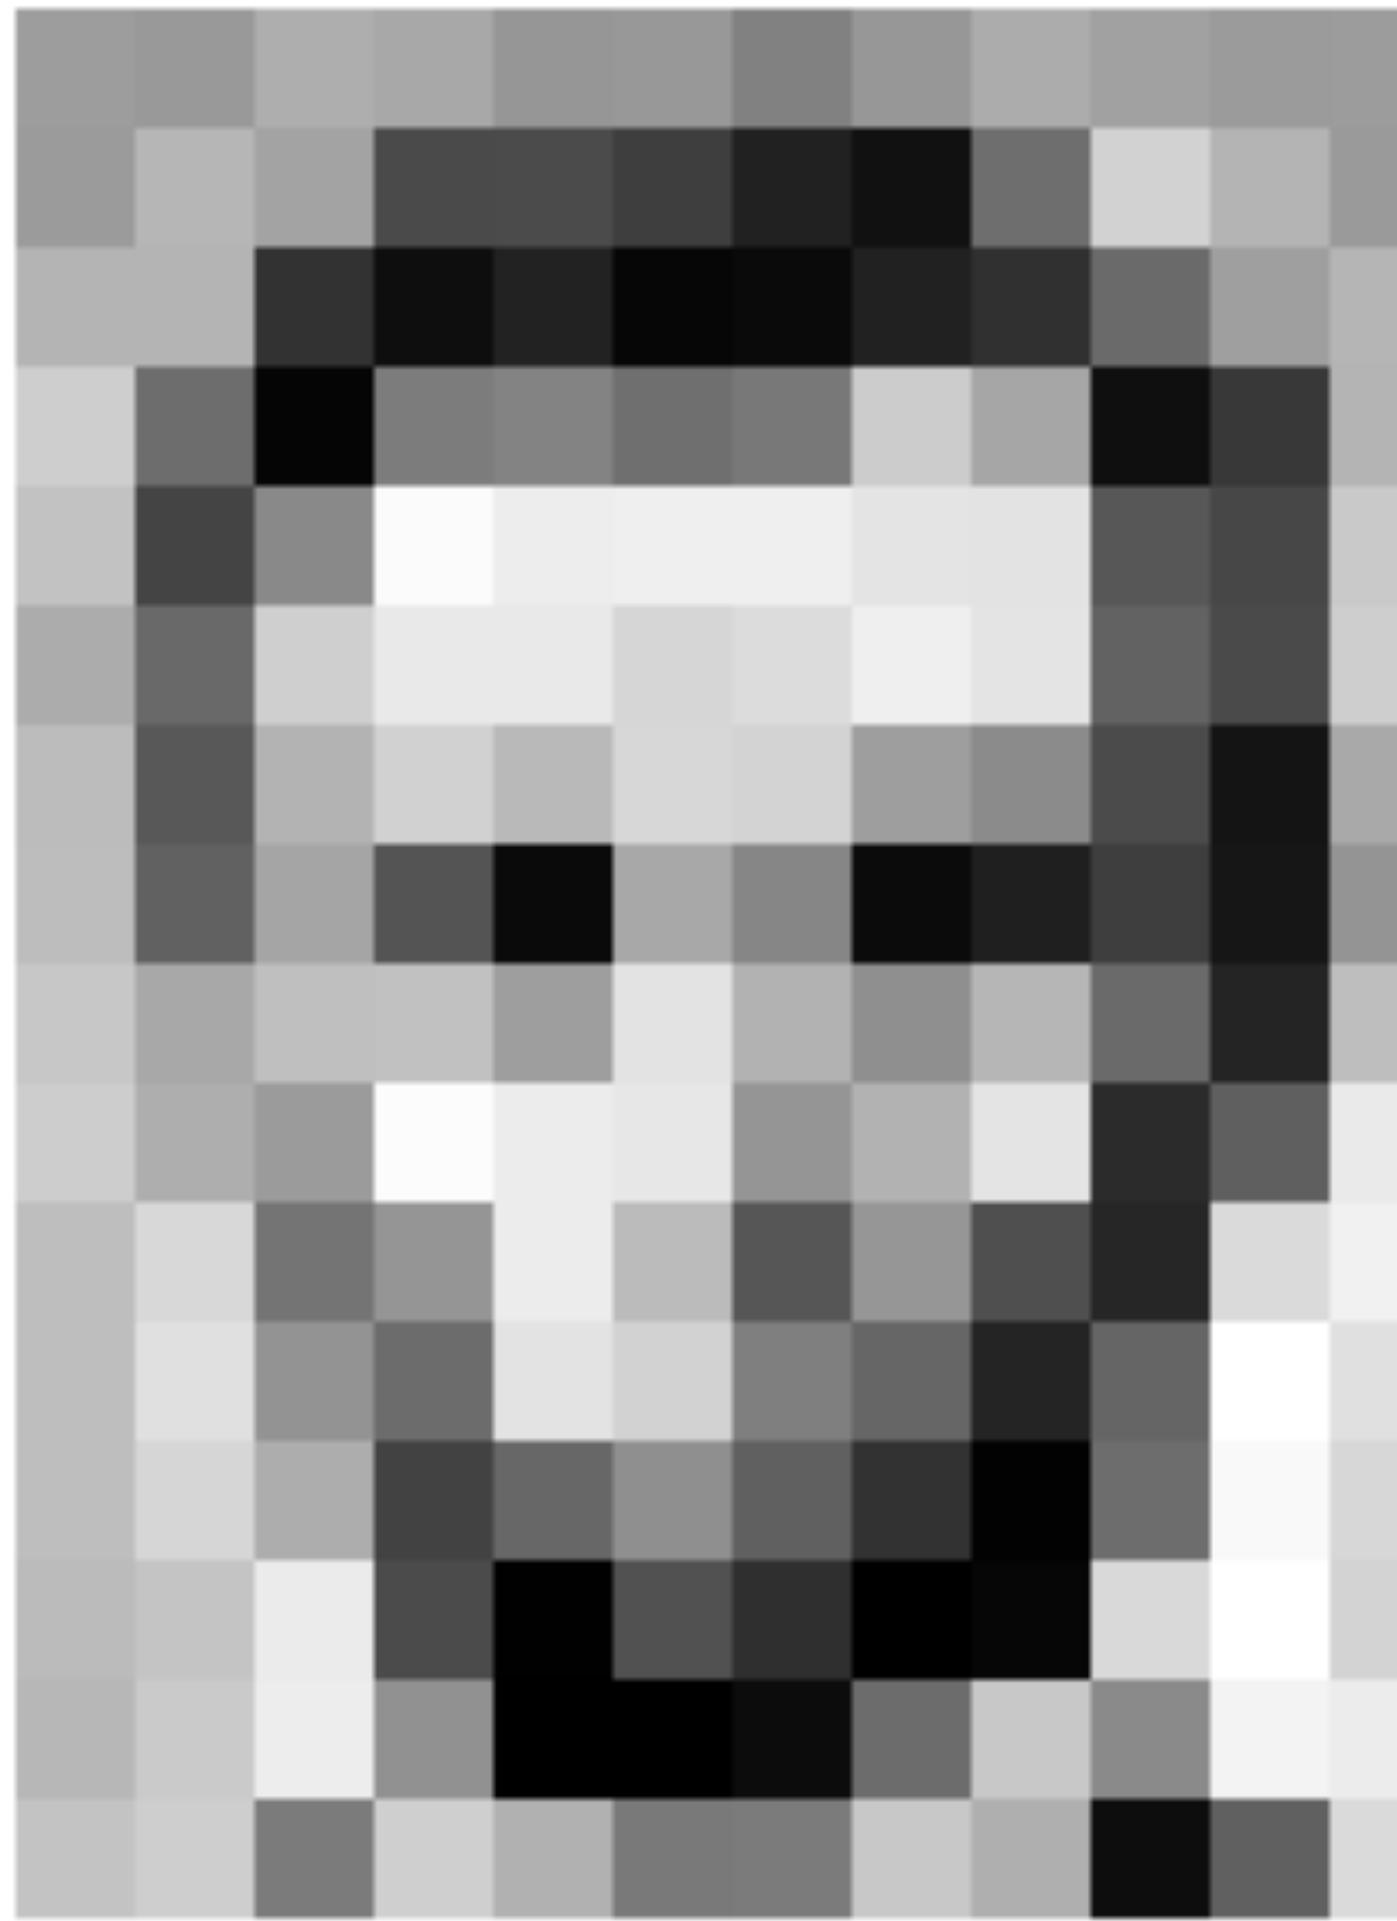
Build a simple convolutional neural network

Augment data

Use a pretrained convolutional neural network

Use transfer learning (fine tuning a pretrained network)

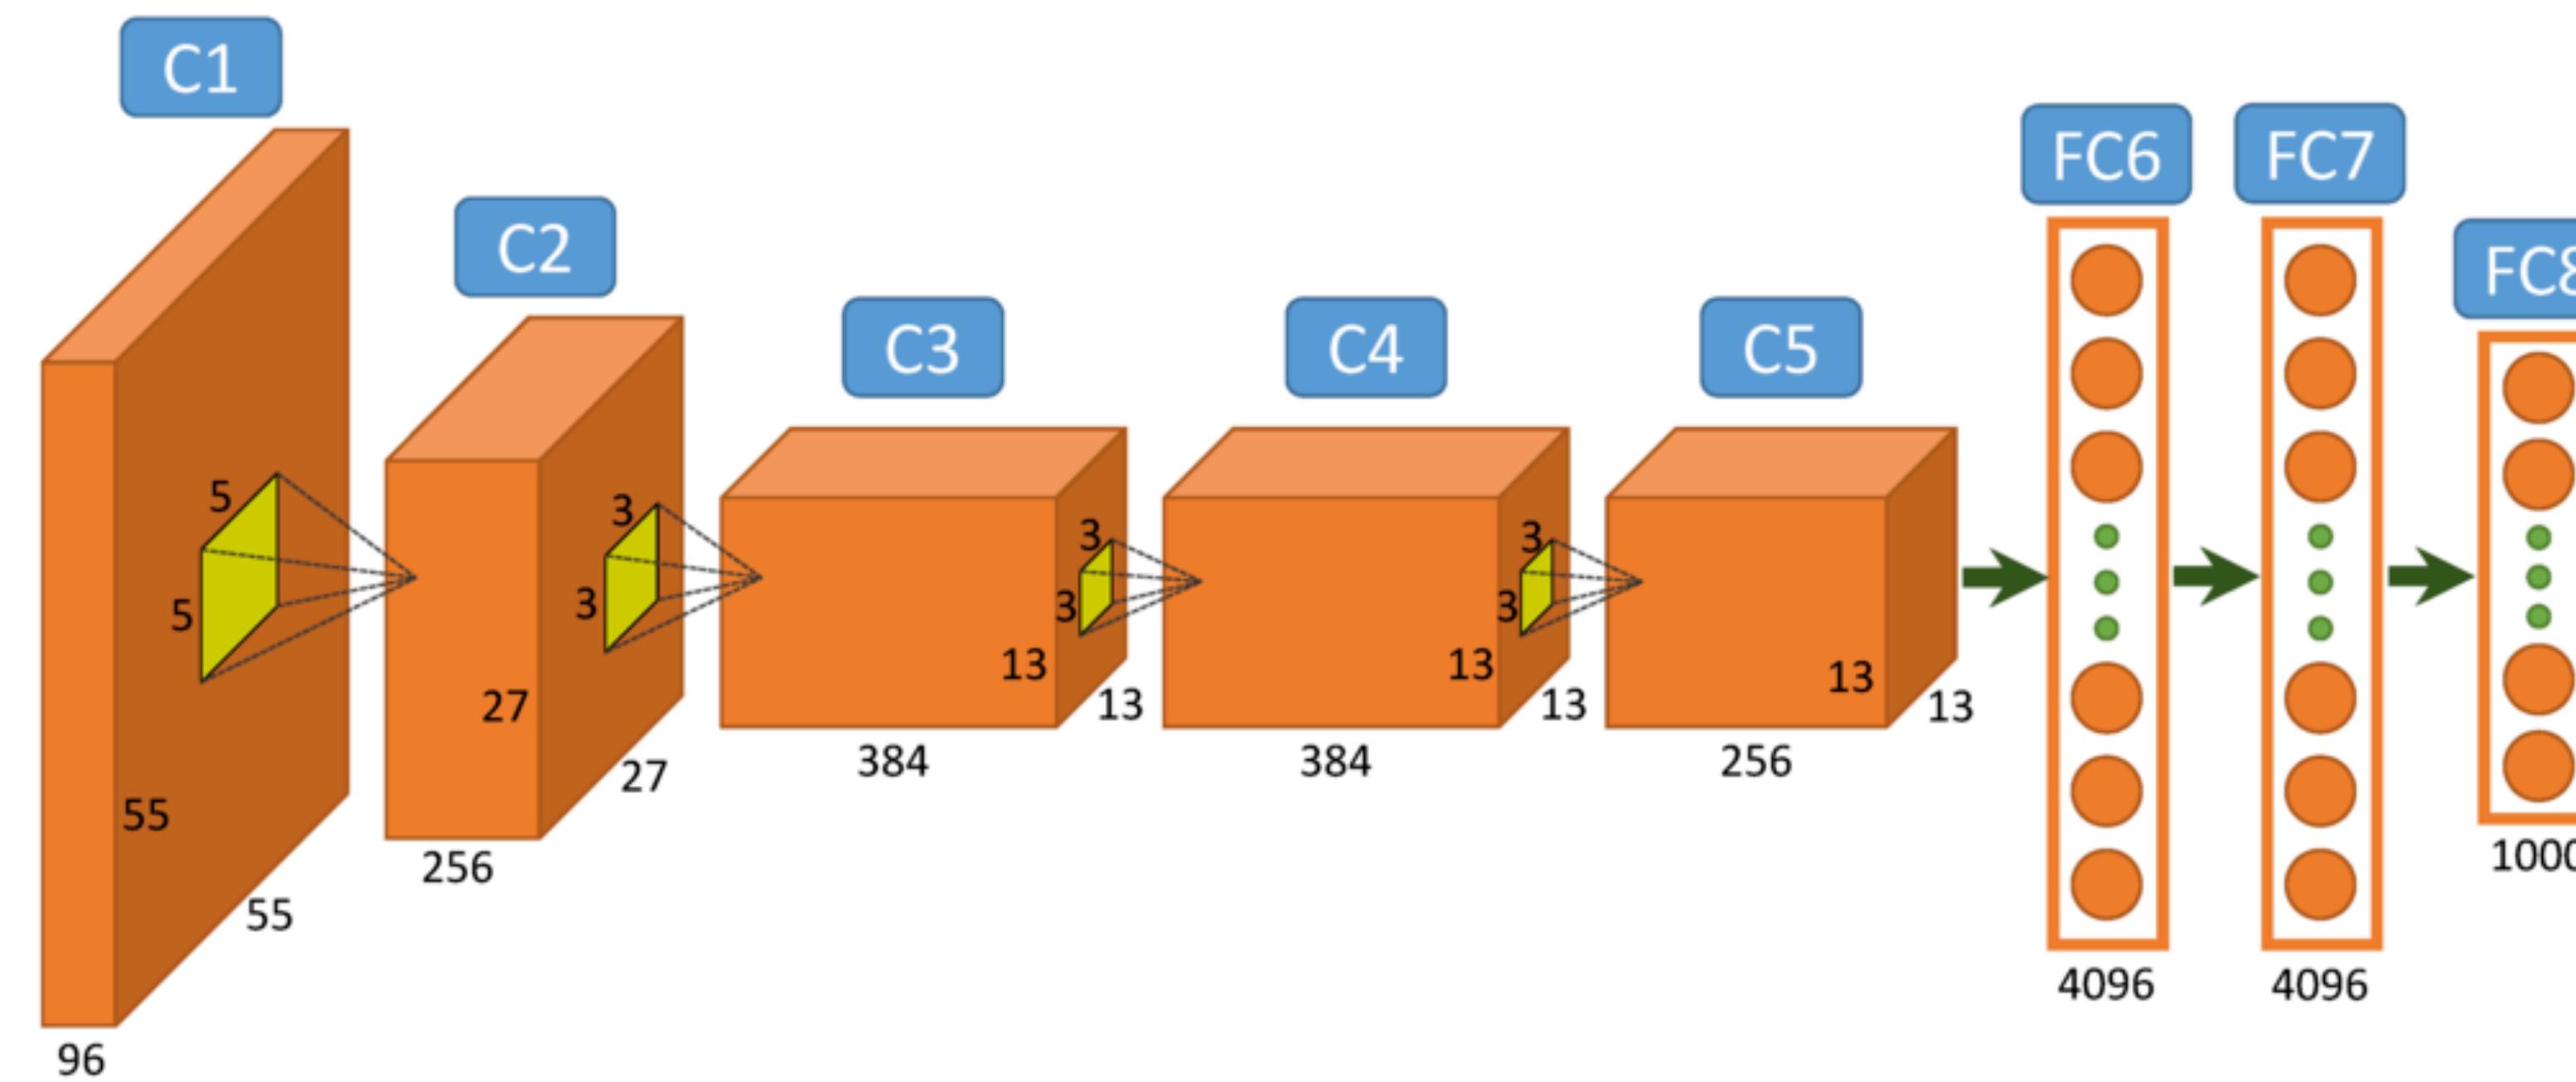
Computer vision:



157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	105	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	105	35	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	35	101	255	224
190	214	173	65	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

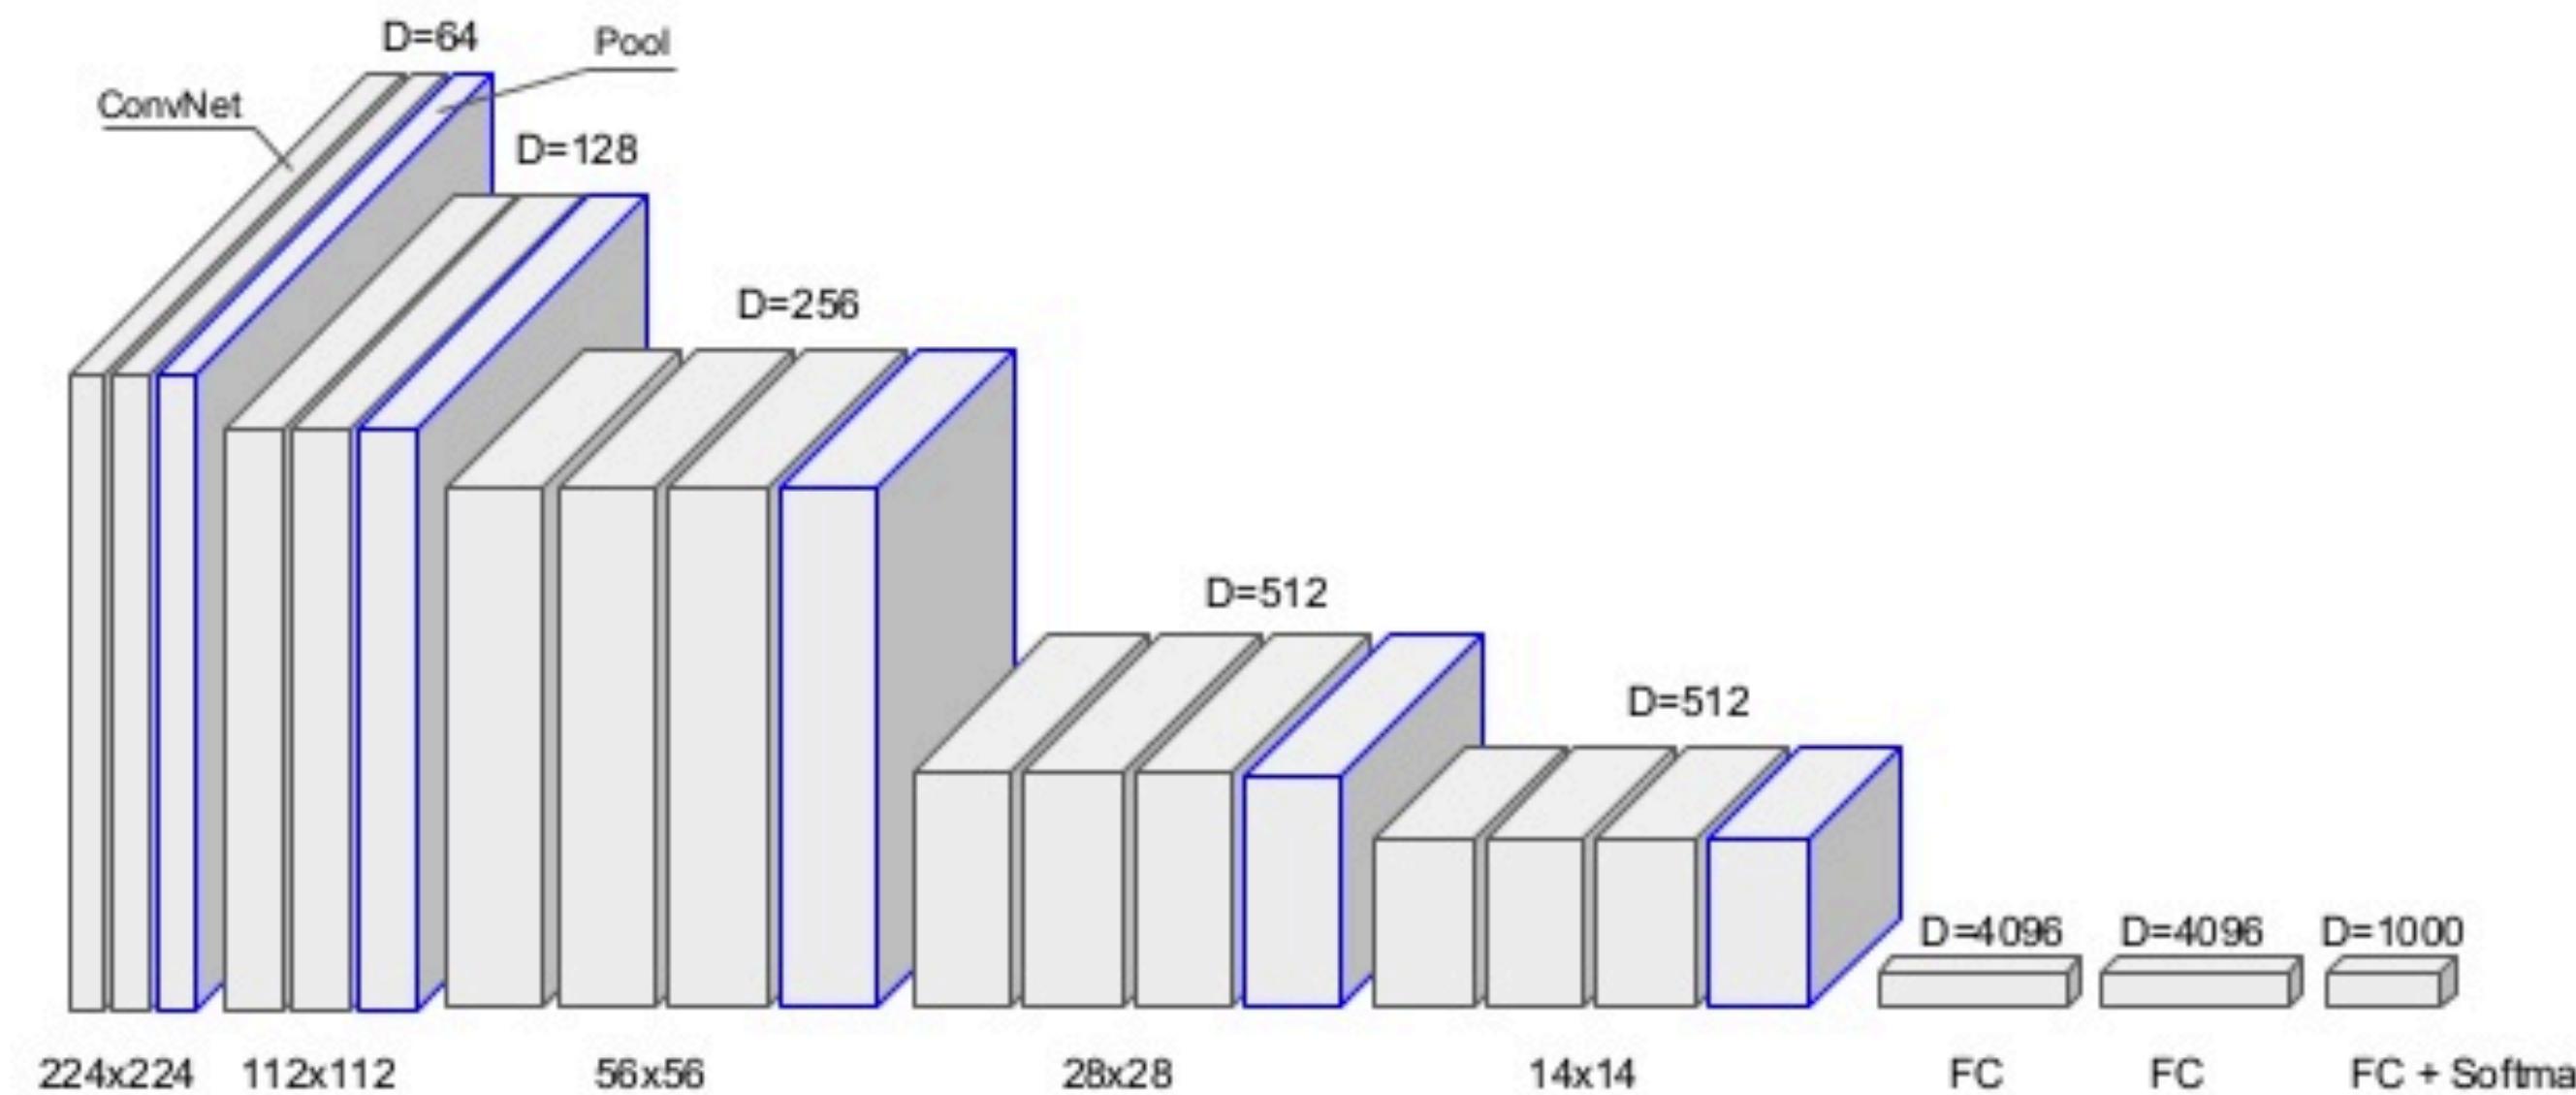
157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	237	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	102	35	101	255	224
190	214	173	65	103	143	95	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

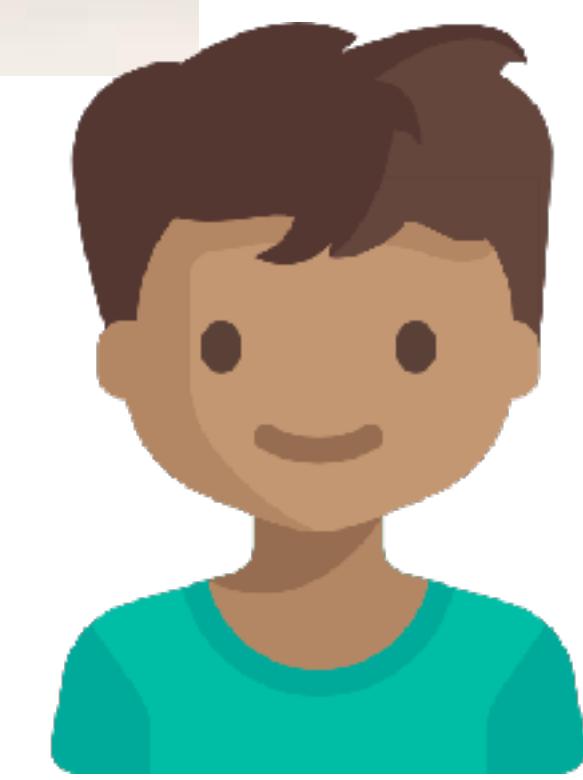
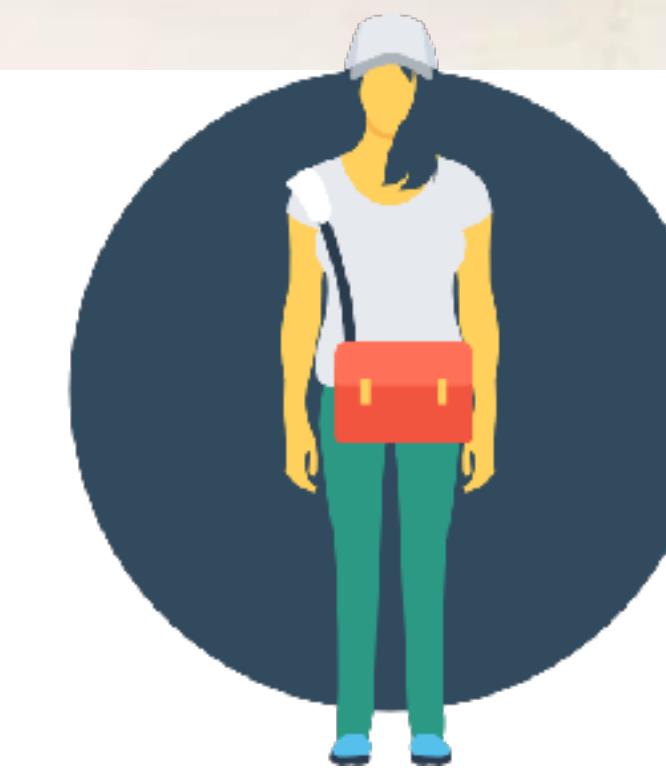
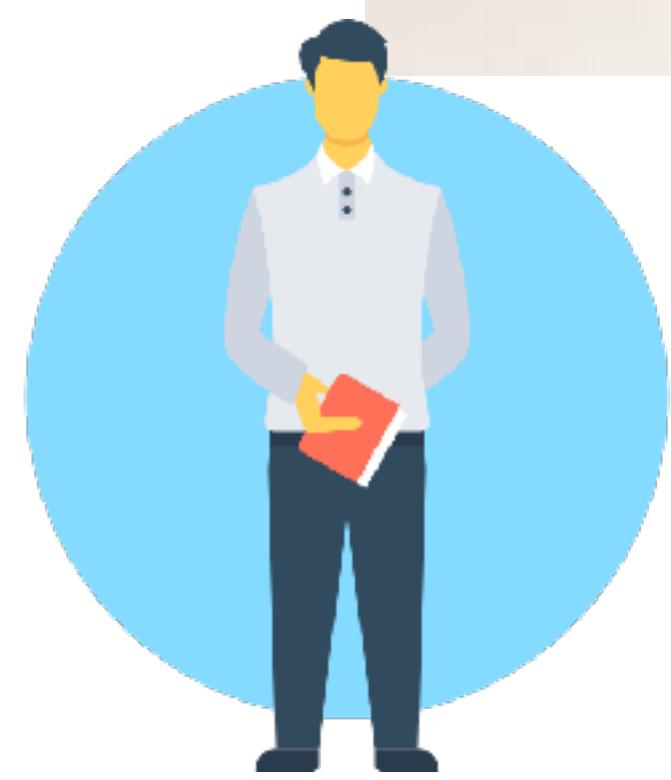
Alexnet architecture

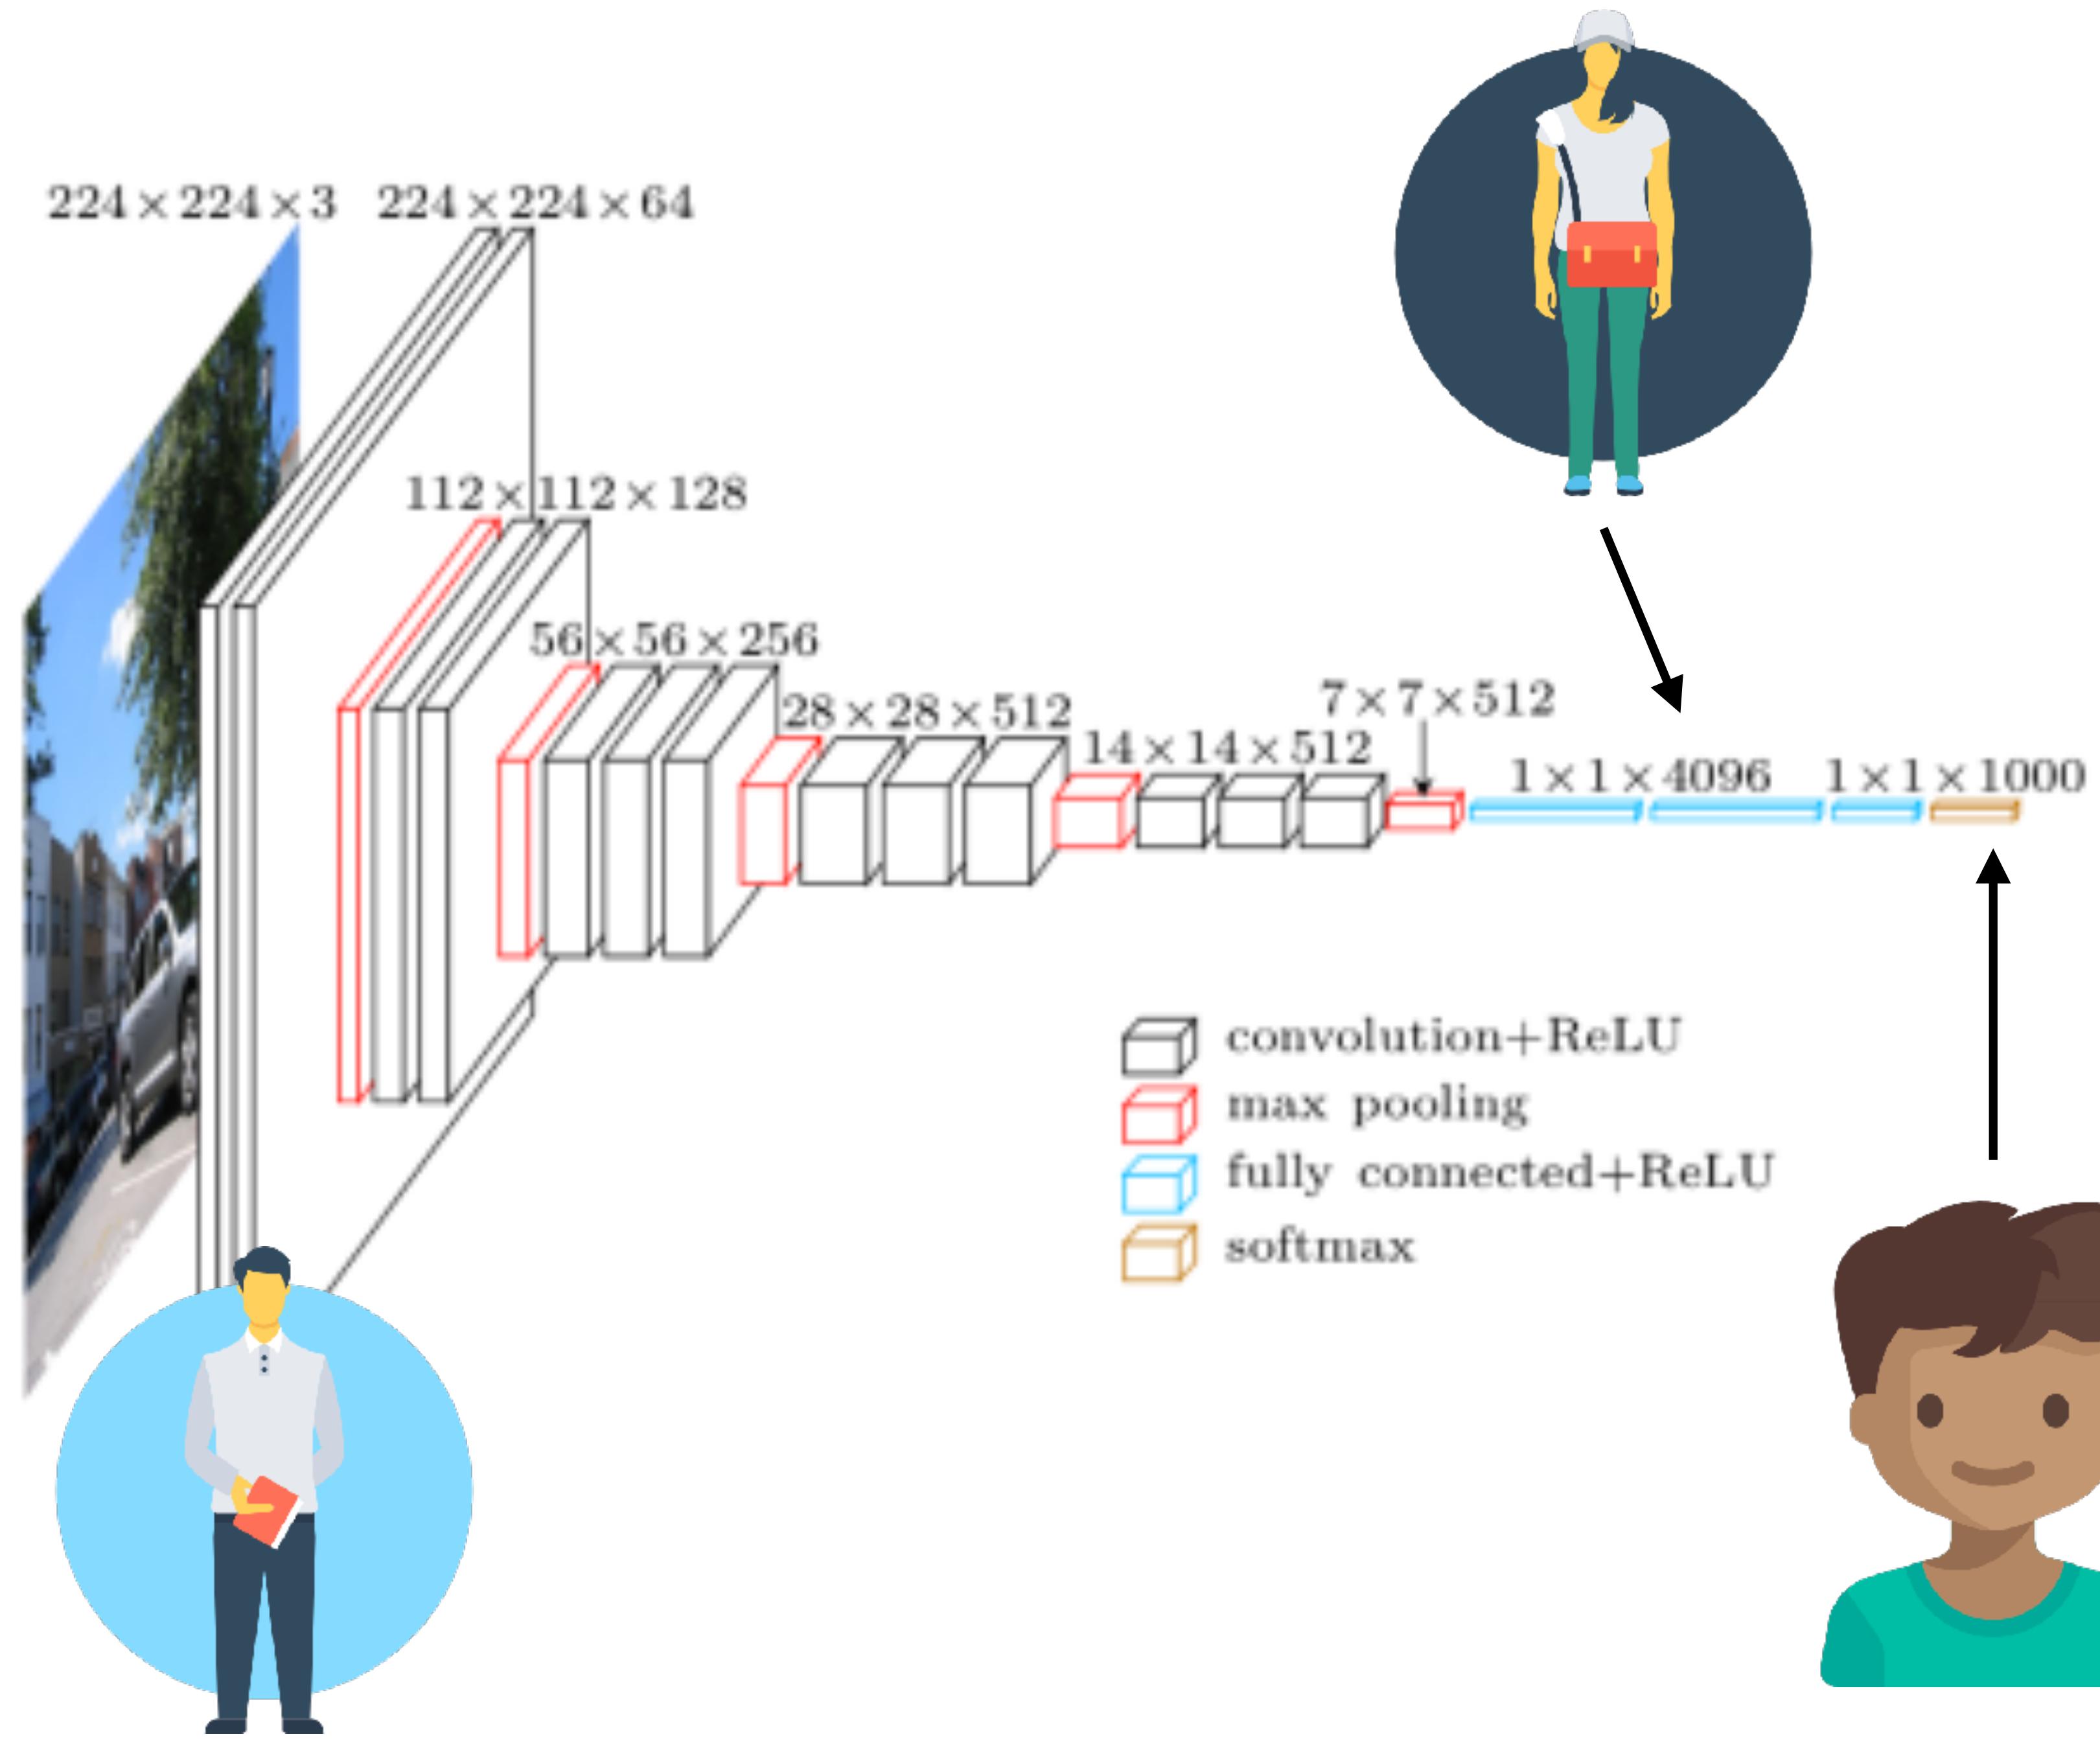


VGG architecture

Classical CNN topology - VGGNet (2013)







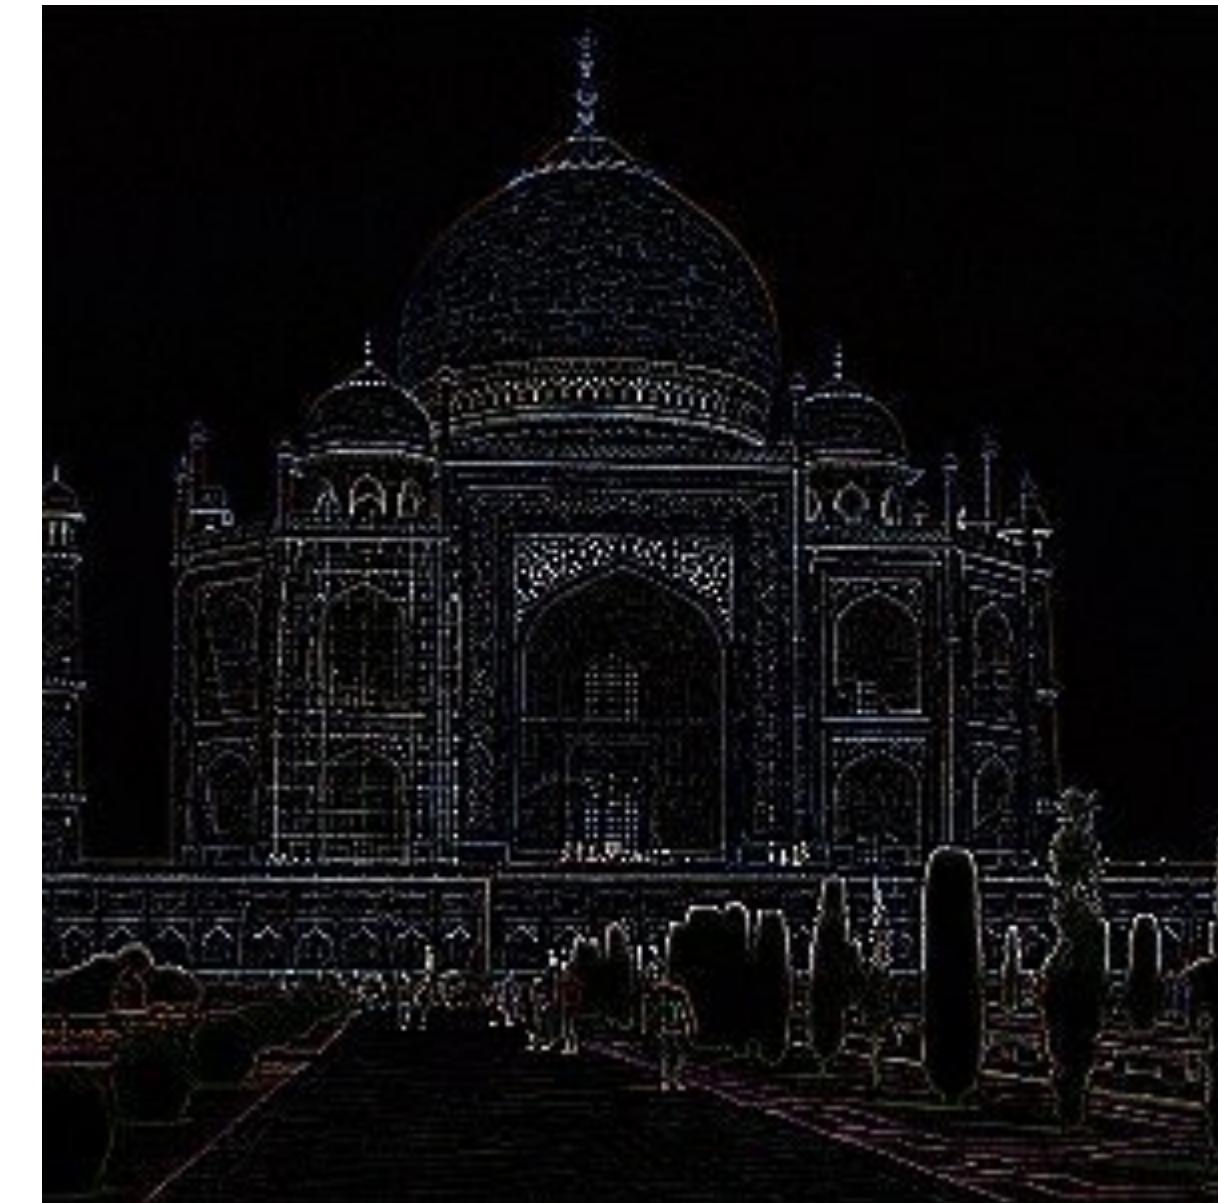
understanding convolution

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature



Simple convolutional neural network

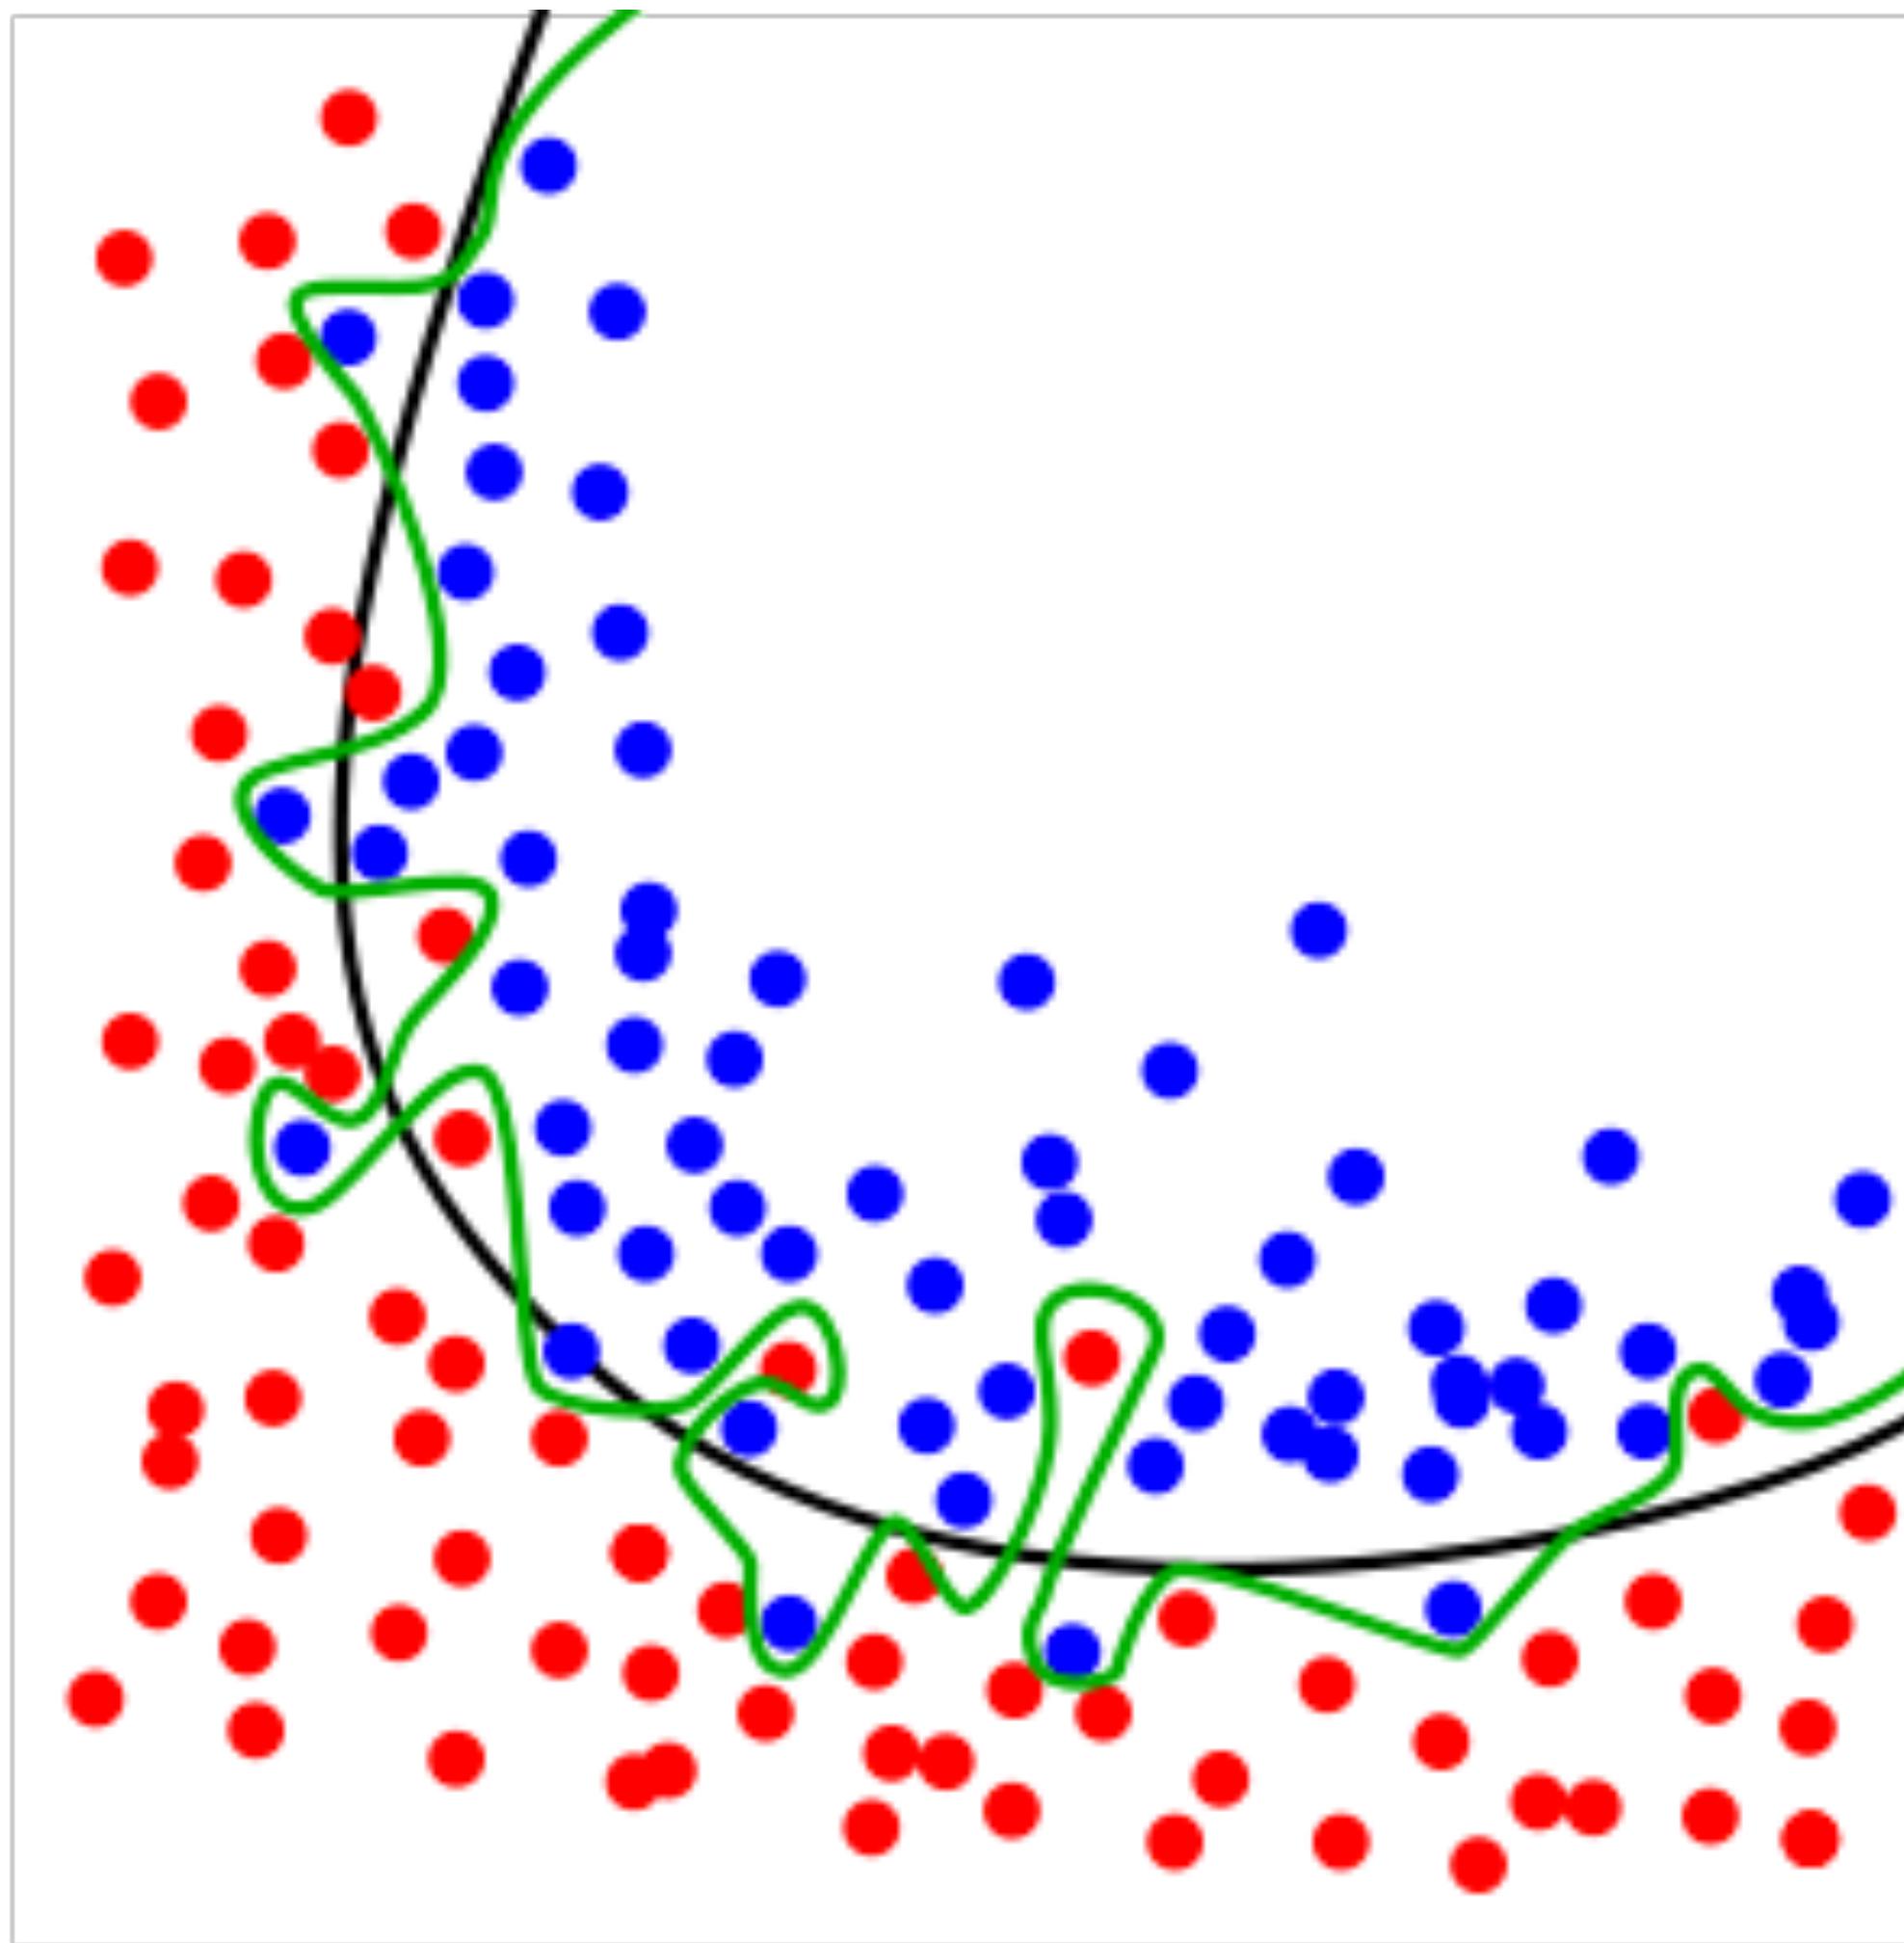
```
model %>%
  layer_conv_2d(filter = 32, kernel_size = c(3,3), input_shape = c(img_width,
img_height, 3)) %>%
  layer_activation("relu") %>%
  layer_max_pooling_2d(pool_size = c(2,2)) %>%
  
  layer_conv_2d(filter = 32, kernel_size = c(3,3)) %>%
  layer_activation("relu") %>%
  layer_max_pooling_2d(pool_size = c(2,2)) %>%
  
  layer_conv_2d(filter = 64, kernel_size = c(3,3)) %>%
  layer_activation("relu") %>%
  layer_max_pooling_2d(pool_size = c(2,2)) %>%
  
  layer_flatten() %>%
  layer_dense(64) %>%
  layer_activation("relu") %>%
  layer_dropout(0.5) %>%
  layer_dense(1) %>%
  layer_activation("sigmoid")
```

**augmenting
data**

overfitting (tanks)



overfitting



augmentation

Data augmentation for improving the model

By applying random transformation to our train set, we artificially enhance our dataset with new unseen images.



`shear_range`: Rotate image by 0.2

`zoom_range`: Zoom the image by 0.2

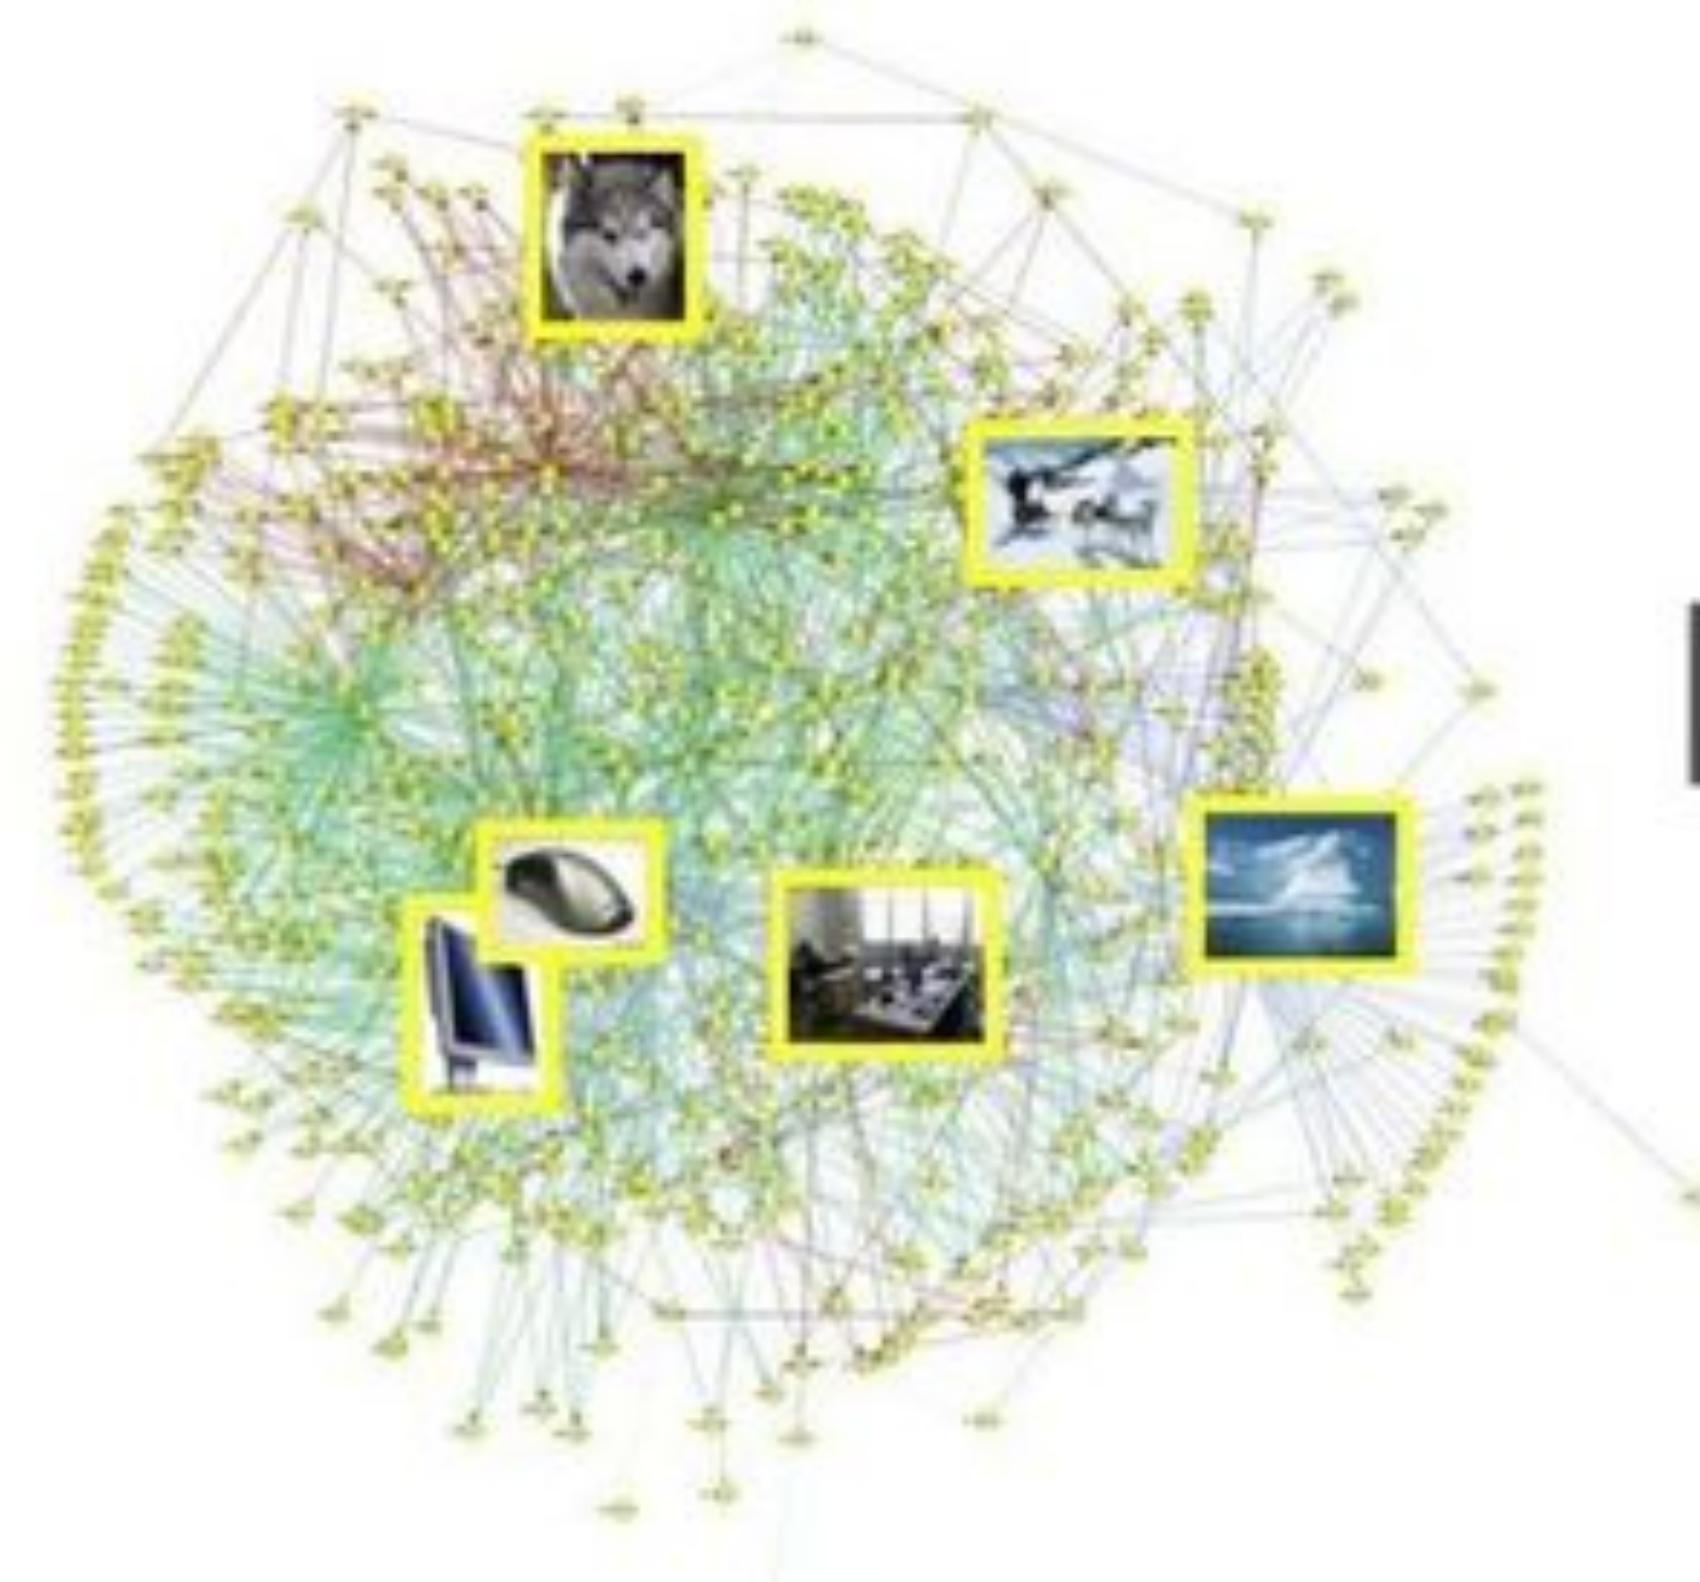
`horizontal_flip`: Randomly flip inputs horizontally

`rescale`: rescaling factor (1./255)

To augment data:

```
augment <- imageDataGenerator(rescale=1./255,  
                           shear_range=0.2,  
                           zoom_range=0.2,  
                           horizontal_flip=TRUE)
```

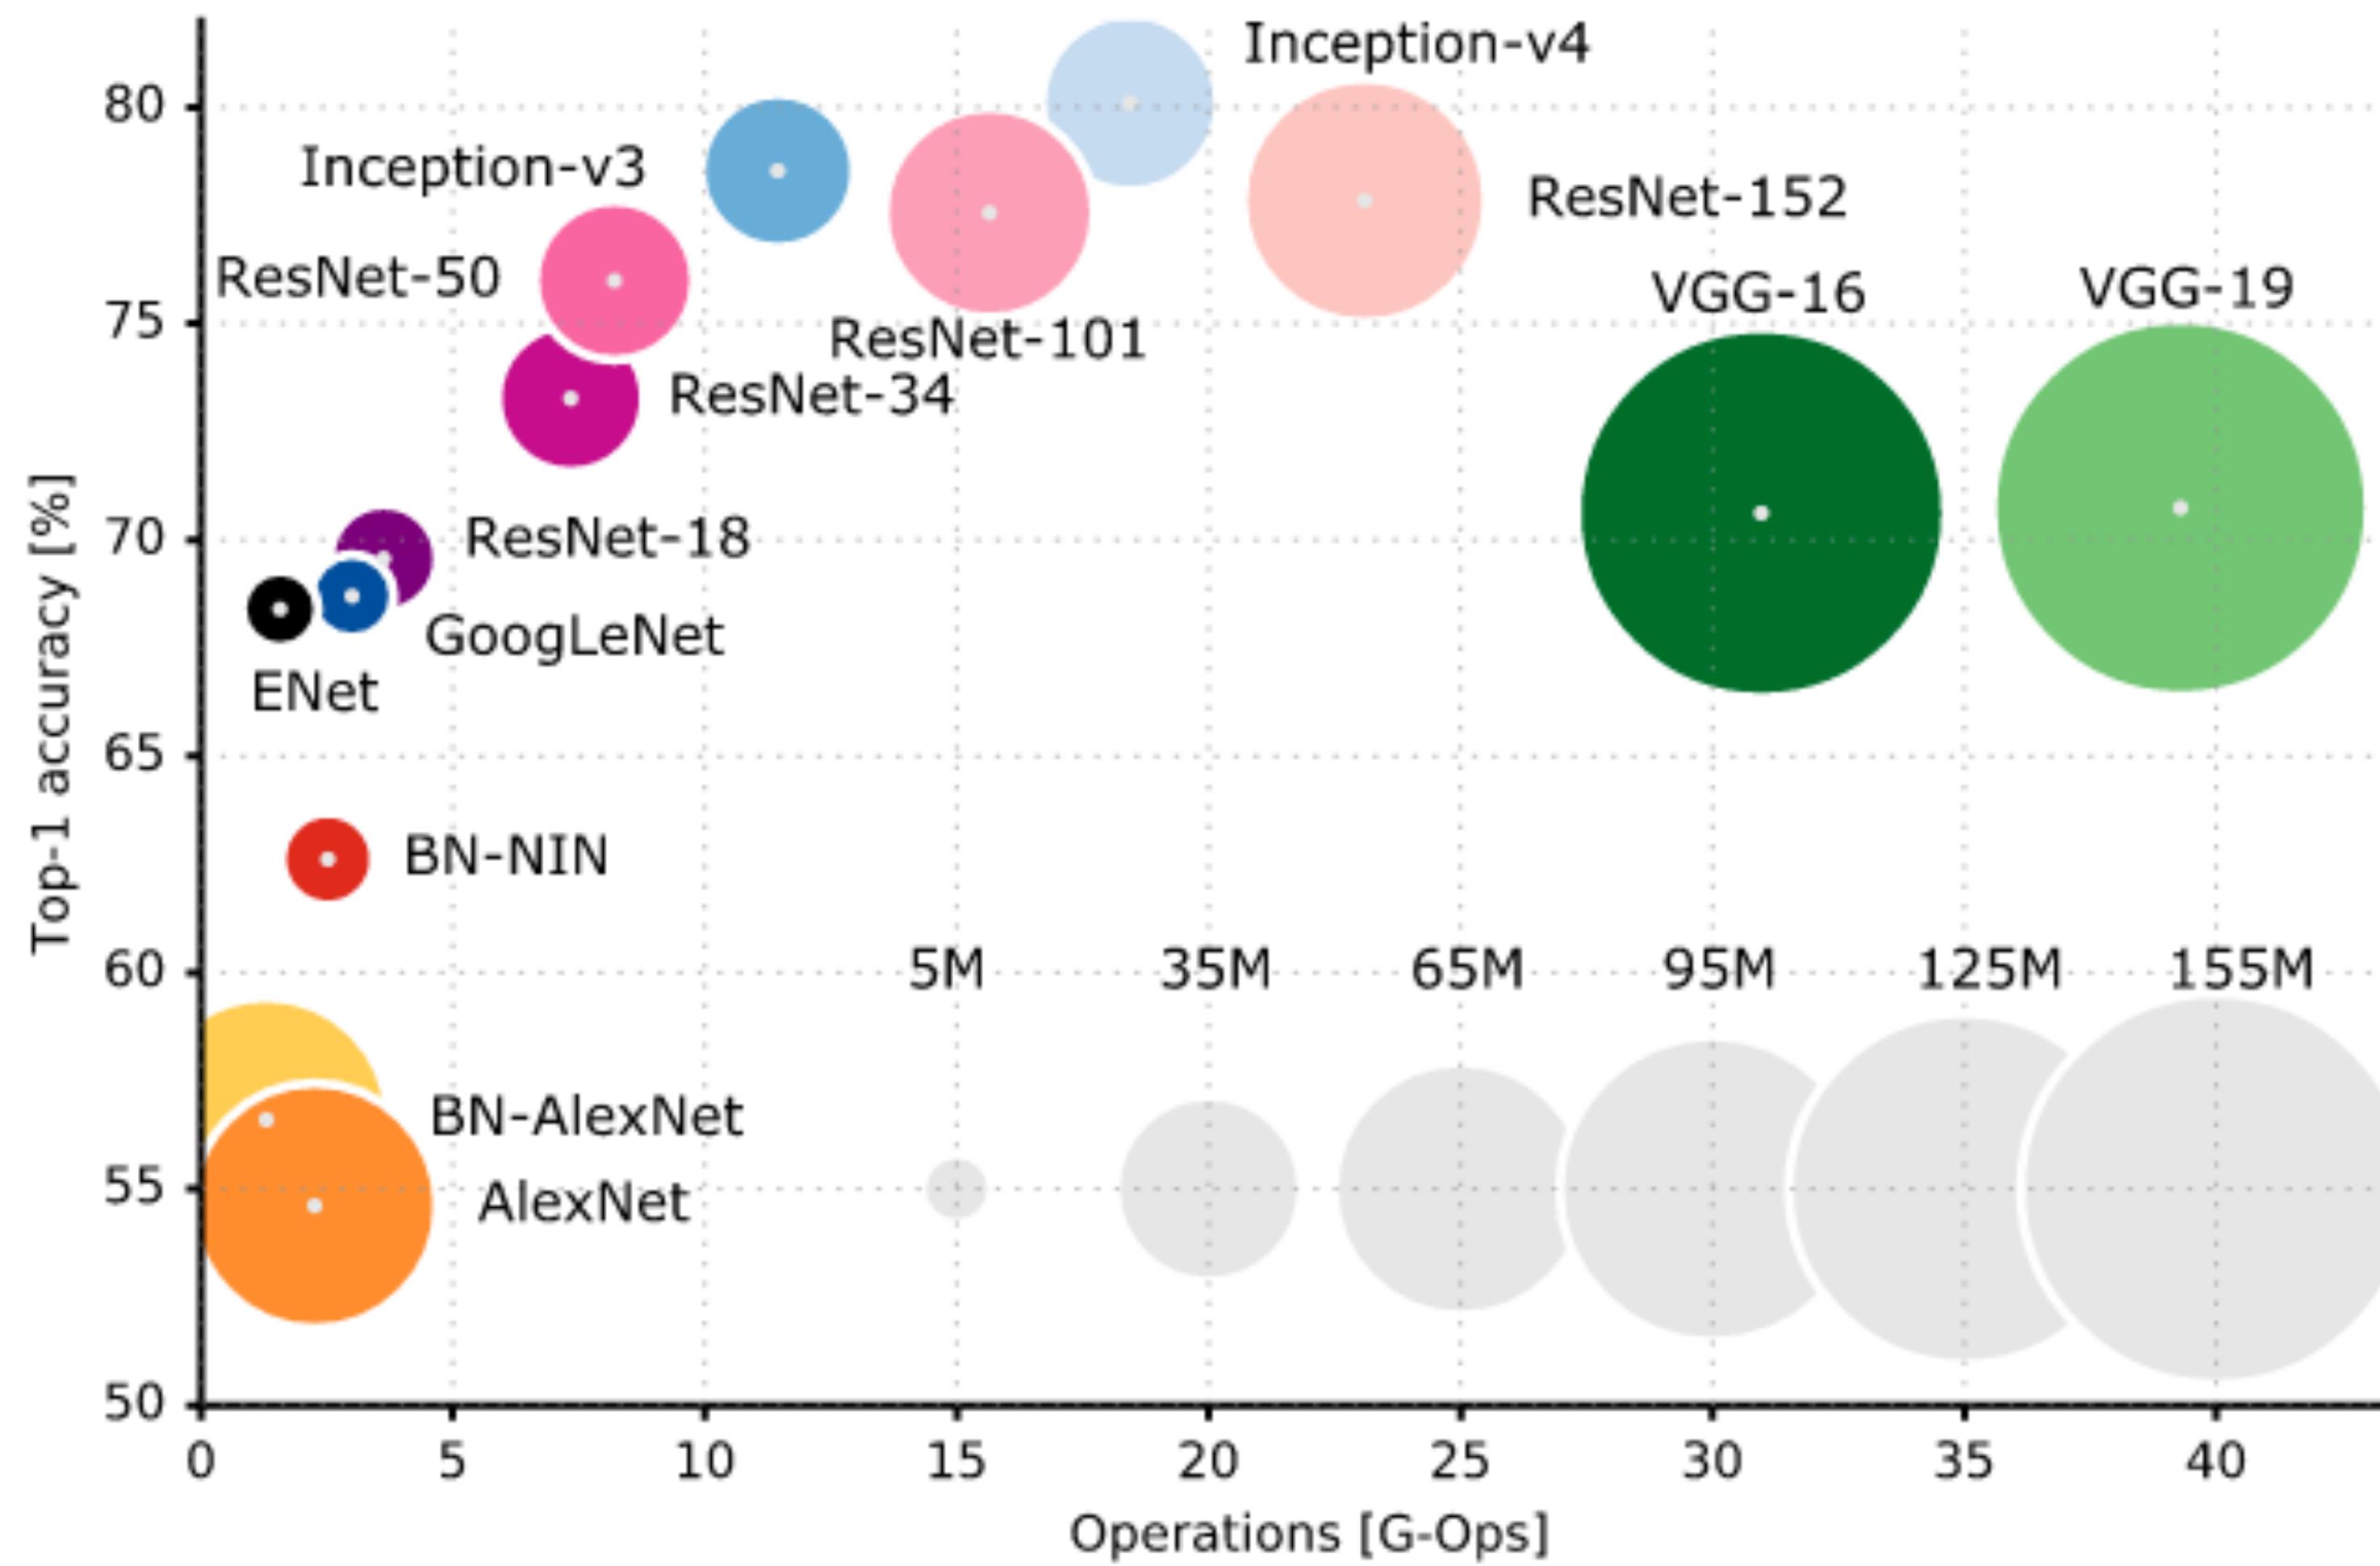
**using a pre-
trained network**



IM_▲GENET

**138 GB, 14 million
images**

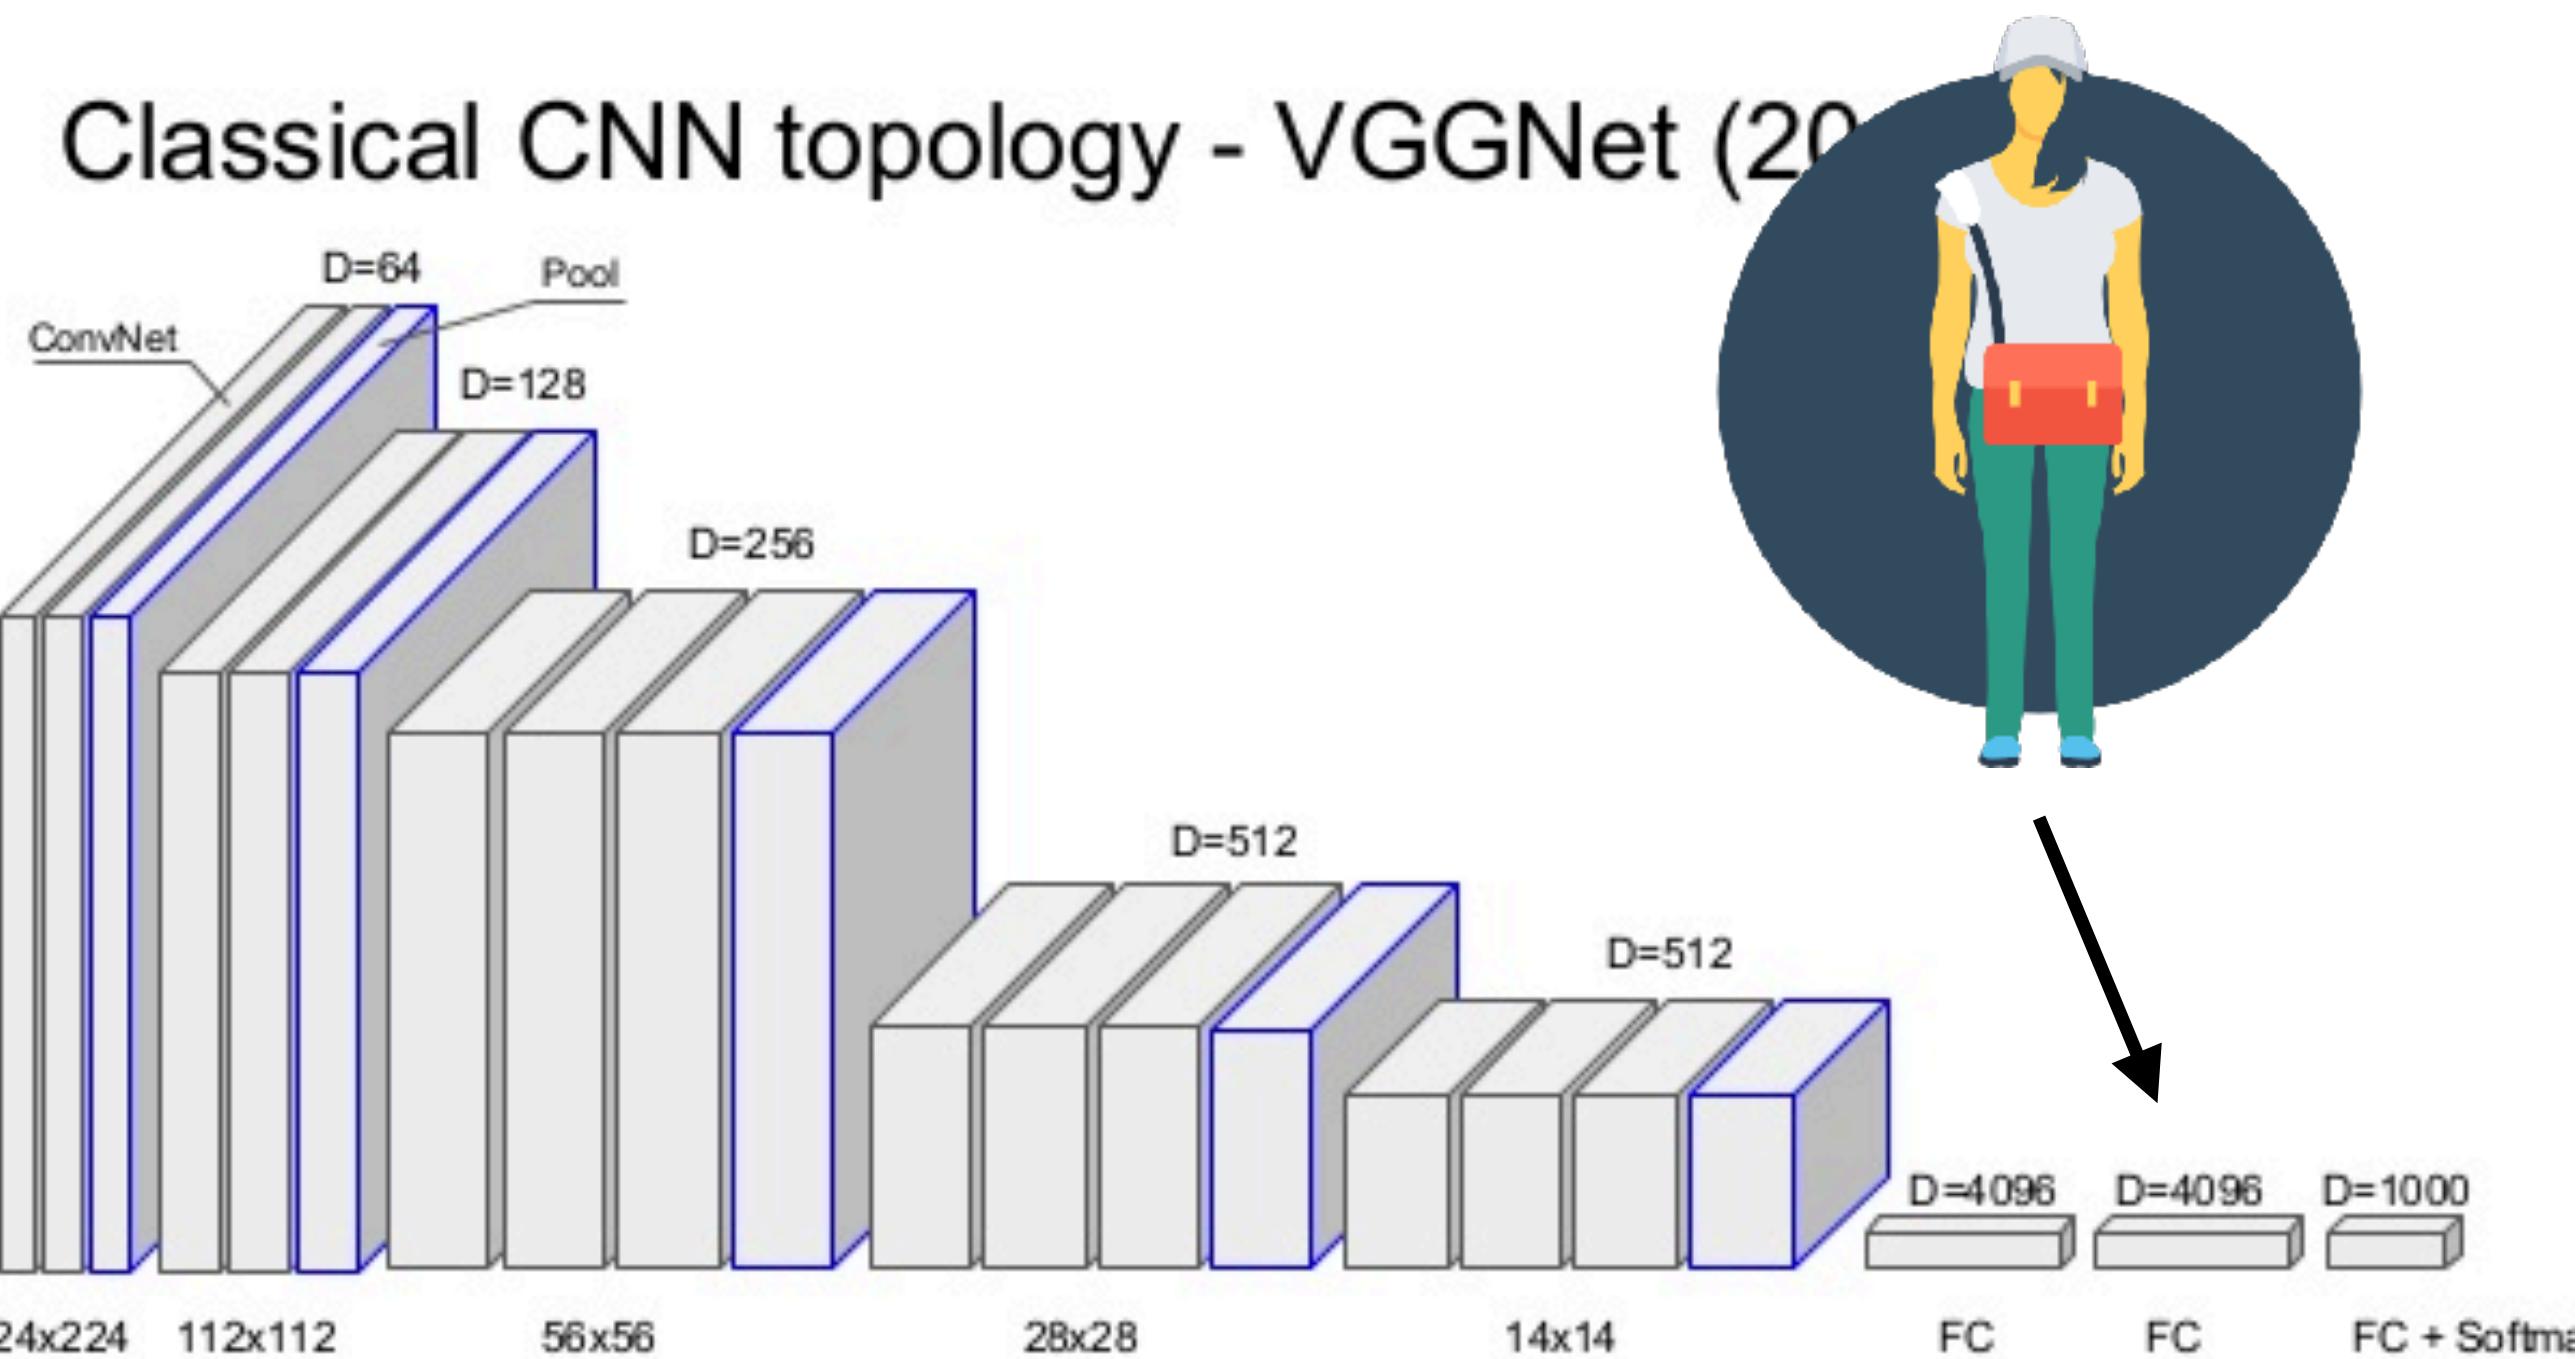
pre-trained networks



Load a pretrained network:

```
model_vgg <- application_vgg16(include_top = FALSE, weights = "imagenet")
```

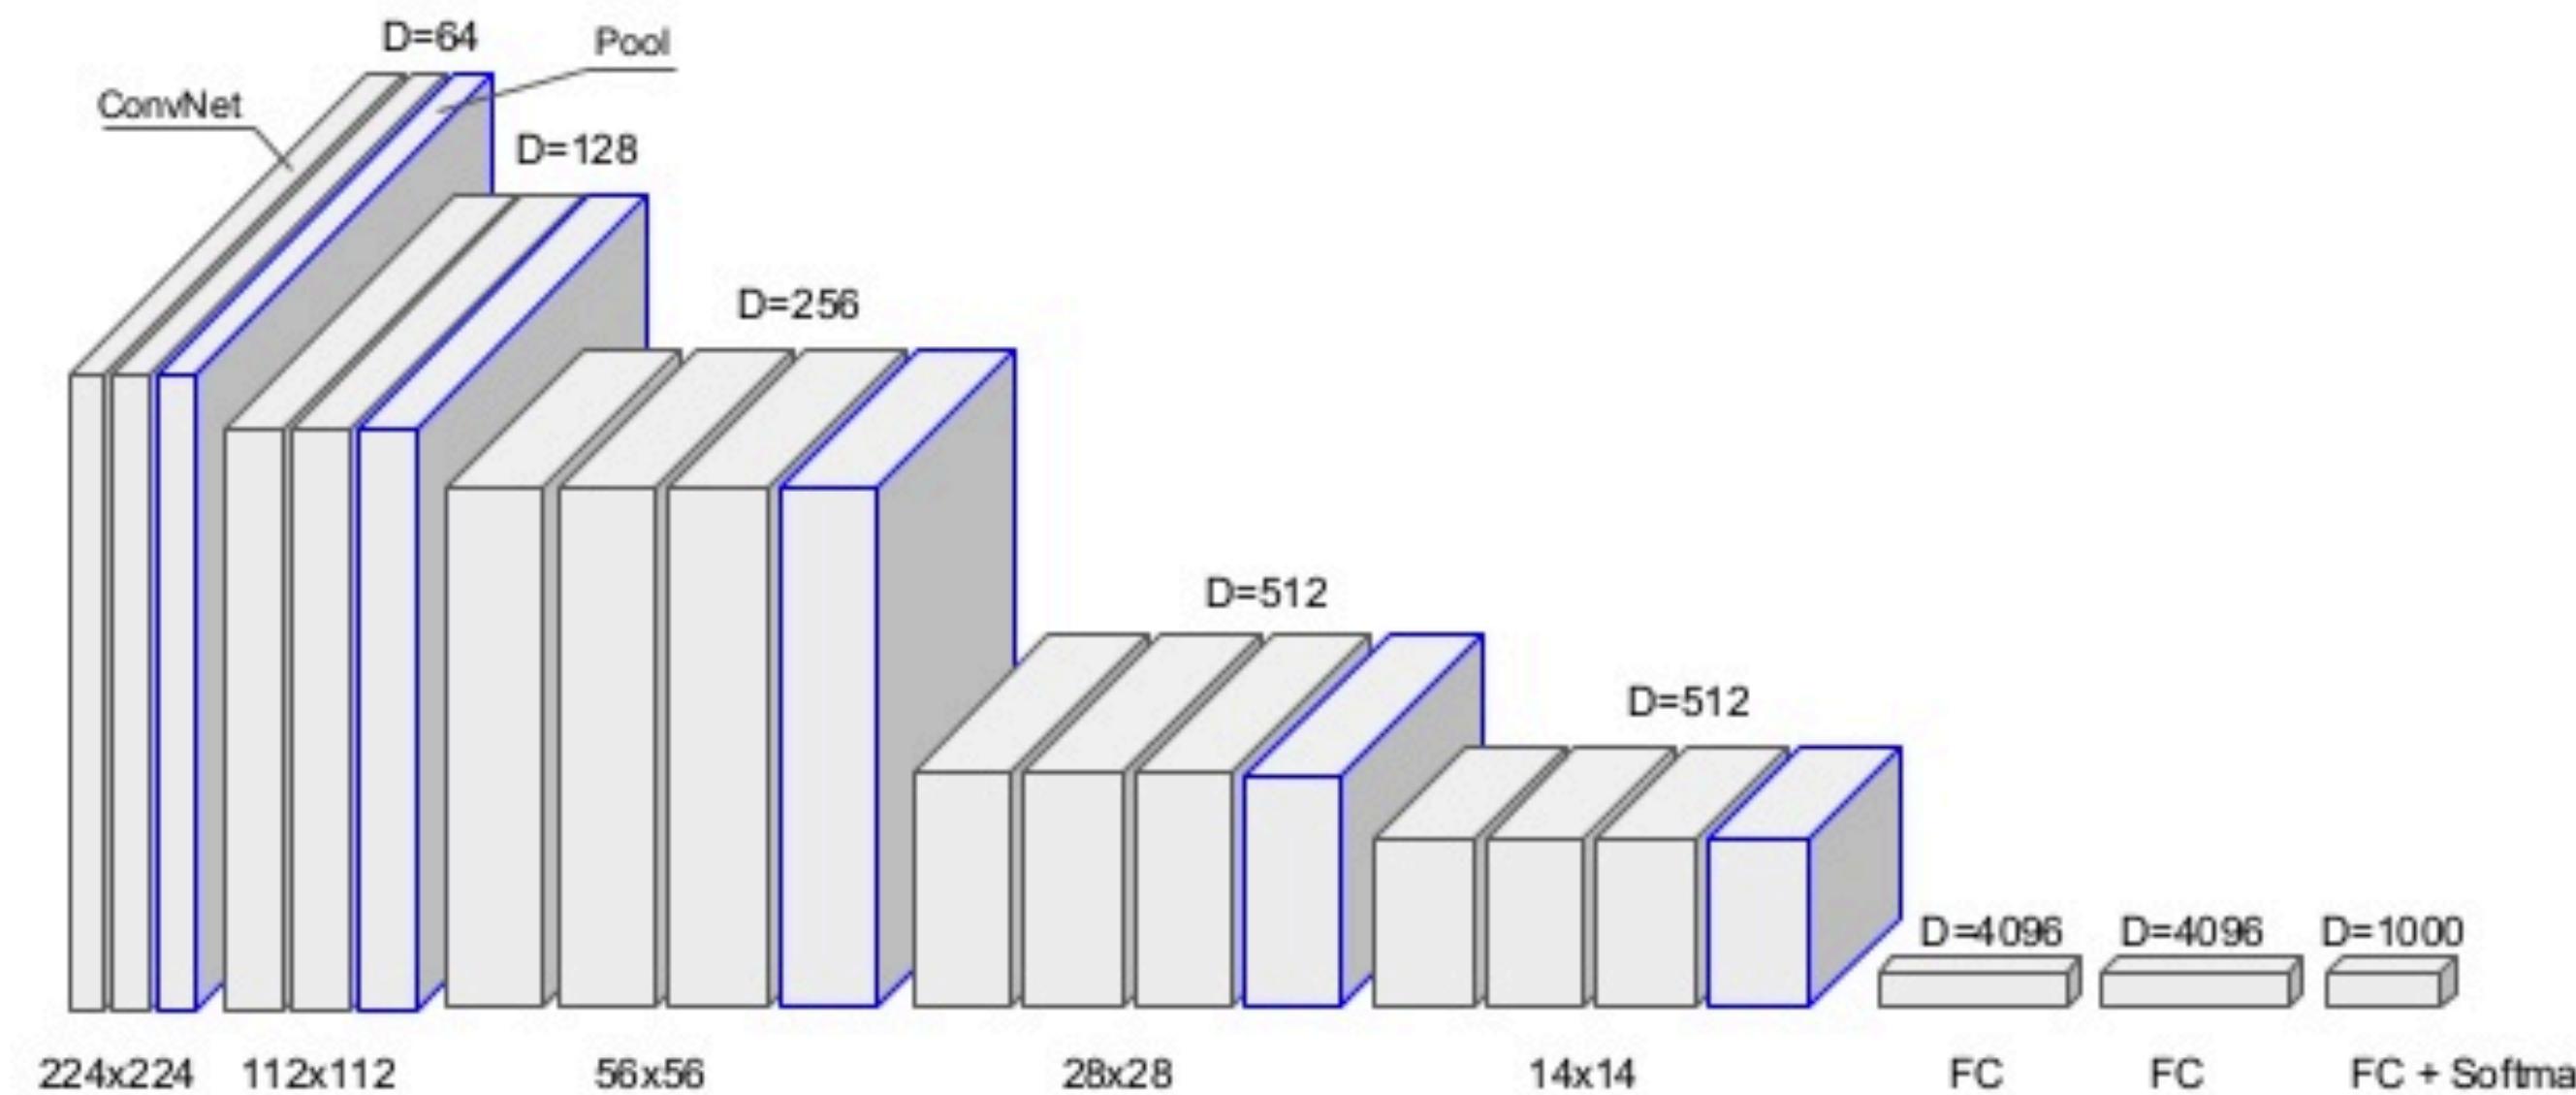
VGG architecture



**fine tuning a
pre-trained network**

VGG architecture

Classical CNN topology - VGGNet (2013)



Train only last few layers:

```
for (layer in model_ft$layers[1:16])  
  layer$trainable <- FALSE
```

Complete Notebook

https://github.com/rajshah4/image_keras

The screenshot shows a web browser window with the title bar 'Image Classification Done Sim...'. The address bar contains the URL 'htmlpreview.github.io/?https://github.com/rajshah4/image_keras/blob/master/Rnotebook.nb.html'. The main content area displays a Jupyter Notebook page. The title is 'Image Classification Done Simply Using Keras' by Rajiv Shah. It includes a note about finding code and data at https://github.com/rajshah4/image_keras, a reference to Francois Fleuret's tutorial, and a link to a GitHub repository. The page is divided into sections: 'Introduction', 'Data', and 'Loading Tensorflow and Keras'. The 'Introduction' section discusses building an image classifier using Keras from a few hundred to thousand pictures per class, mentioning training from scratch, using pre-trained networks, and fine-tuning. It also lists features like fit_generator and ImageDataGenerator. The 'Data' section describes a folder structure for training, validation, and test sets, each containing dog and cat images. The 'Loading Tensorflow and Keras' section is partially visible at the bottom.

Image Classification Done Simply Using Keras

by Rajiv Shah

You can find the code and data for this notebook at: https://github.com/rajshah4/image_keras

The work here is based on the tutorial by Francois Fleuret <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html> and the workbook by Guillaume Domici <https://github.com/gggdominic/keras-workshop>

Introduction

This tutorial presents several ways to build an image classifier using keras from just a few hundred or thousand pictures for each class.

We will go over the following options:

- training a small network from scratch (as a baseline)
- using the bottleneck features of a pre-trained network
- fine-tuning the top layers of a pre-trained network

This will lead us to cover the following Keras features:

- fit_generator for training Keras a model using data generators
- ImageDataGenerator for real-time data augmentation
- layer freezing and model fine-tuning
- ...and more.

Data

The github repo includes about 3000 images for this model. The original Kaggle dataset is much larger. The purpose of this demo is to show how you can build models with smaller size datasets. You can improve this model by using more data, which is available at: <https://www.kaggle.com/c/dogs-vs-cats/data>

The recommended folder structure is:

```
data/
  train/
    dogs/ #~ 1024 pictures
      dog001.jpg
      dog002.jpg
      ...
    cats/ #~ 1024 pictures
      cat001.jpg
      cat002.jpg
      ...
  validation/
    dogs/ #~ 416 pictures
      dog001.jpg
      dog002.jpg
      ...
    cats/ #~ 416 pictures
      cat001.jpg
      cat002.jpg
      ...
```

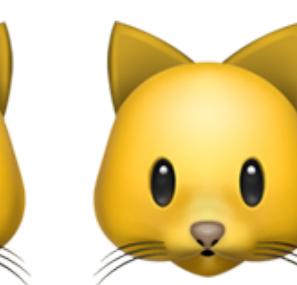
Loading Tensorflow and Keras



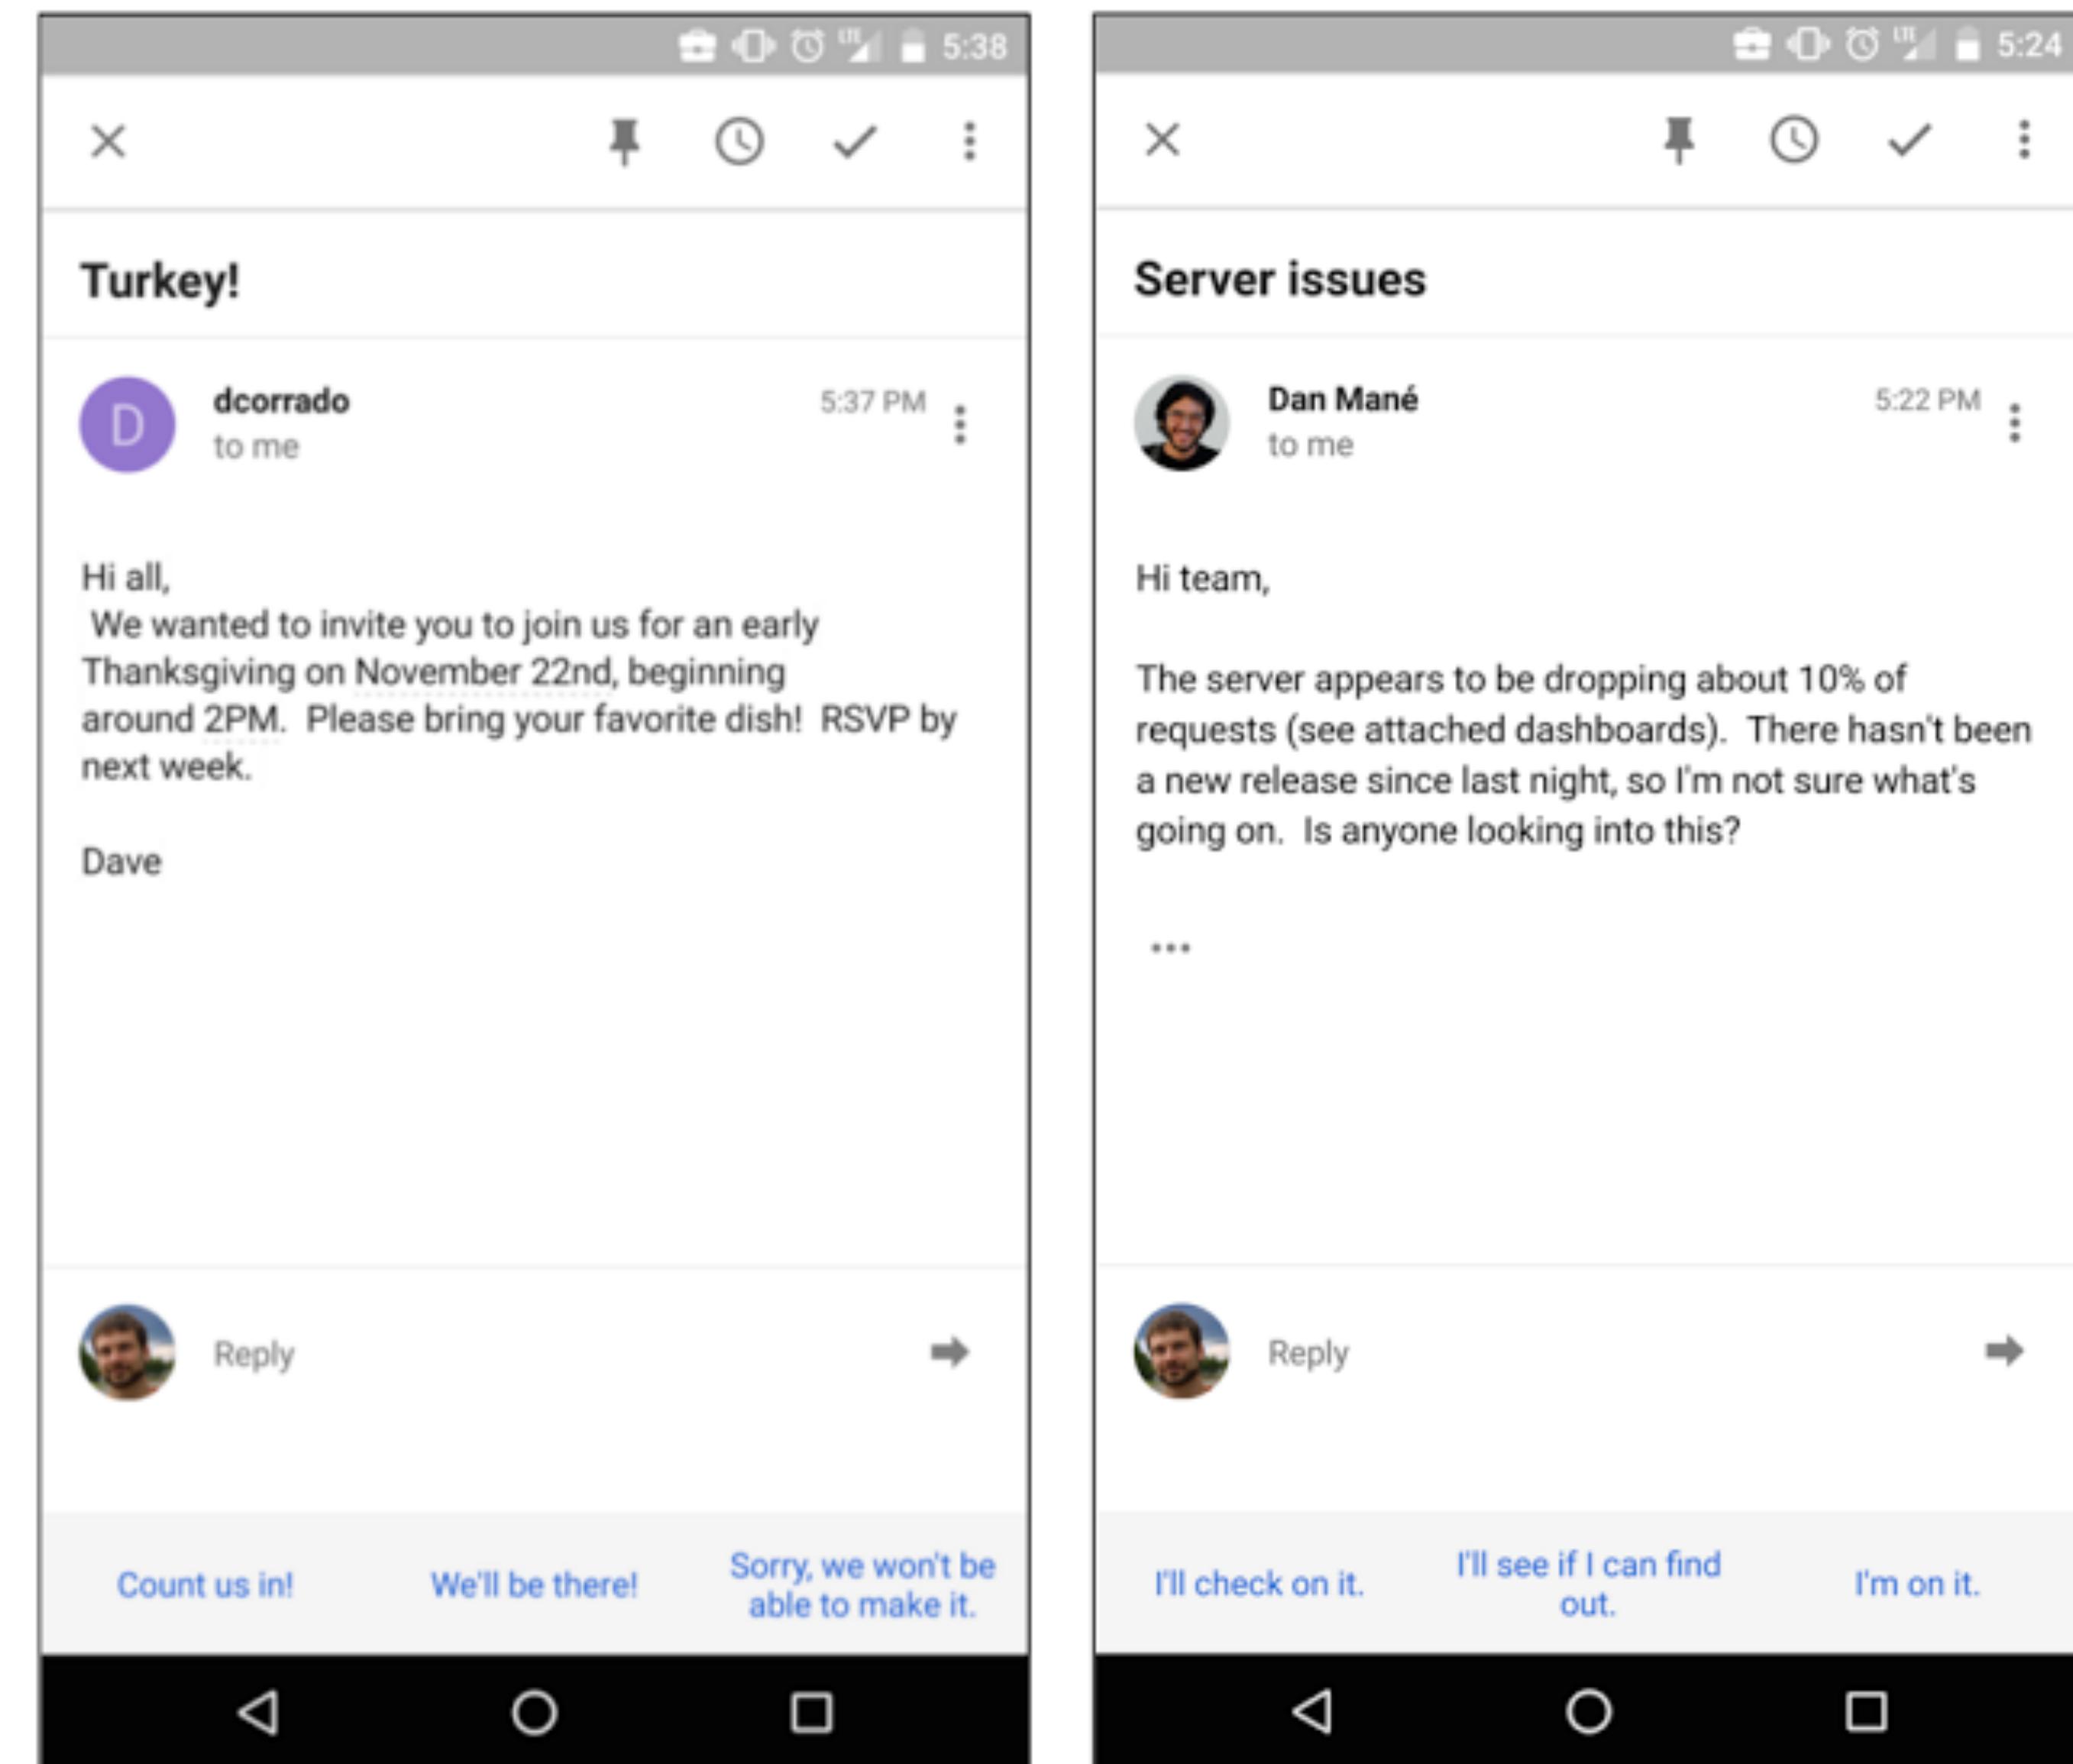
• • •



• • •



smart reply



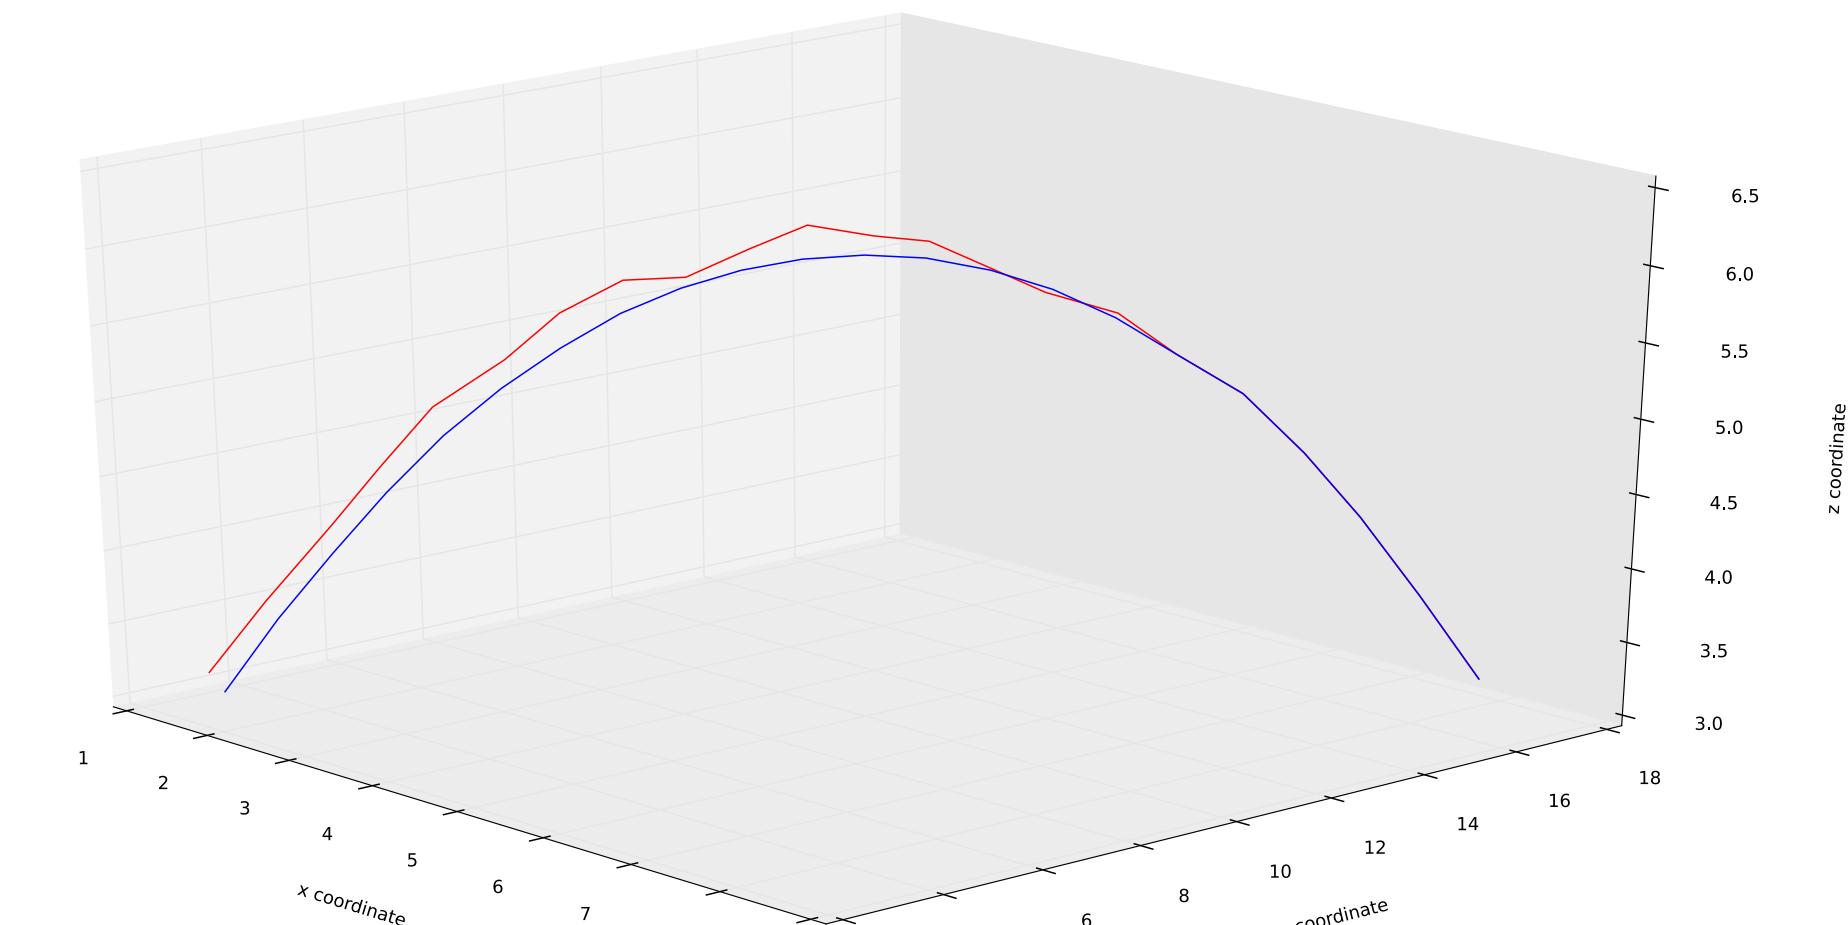
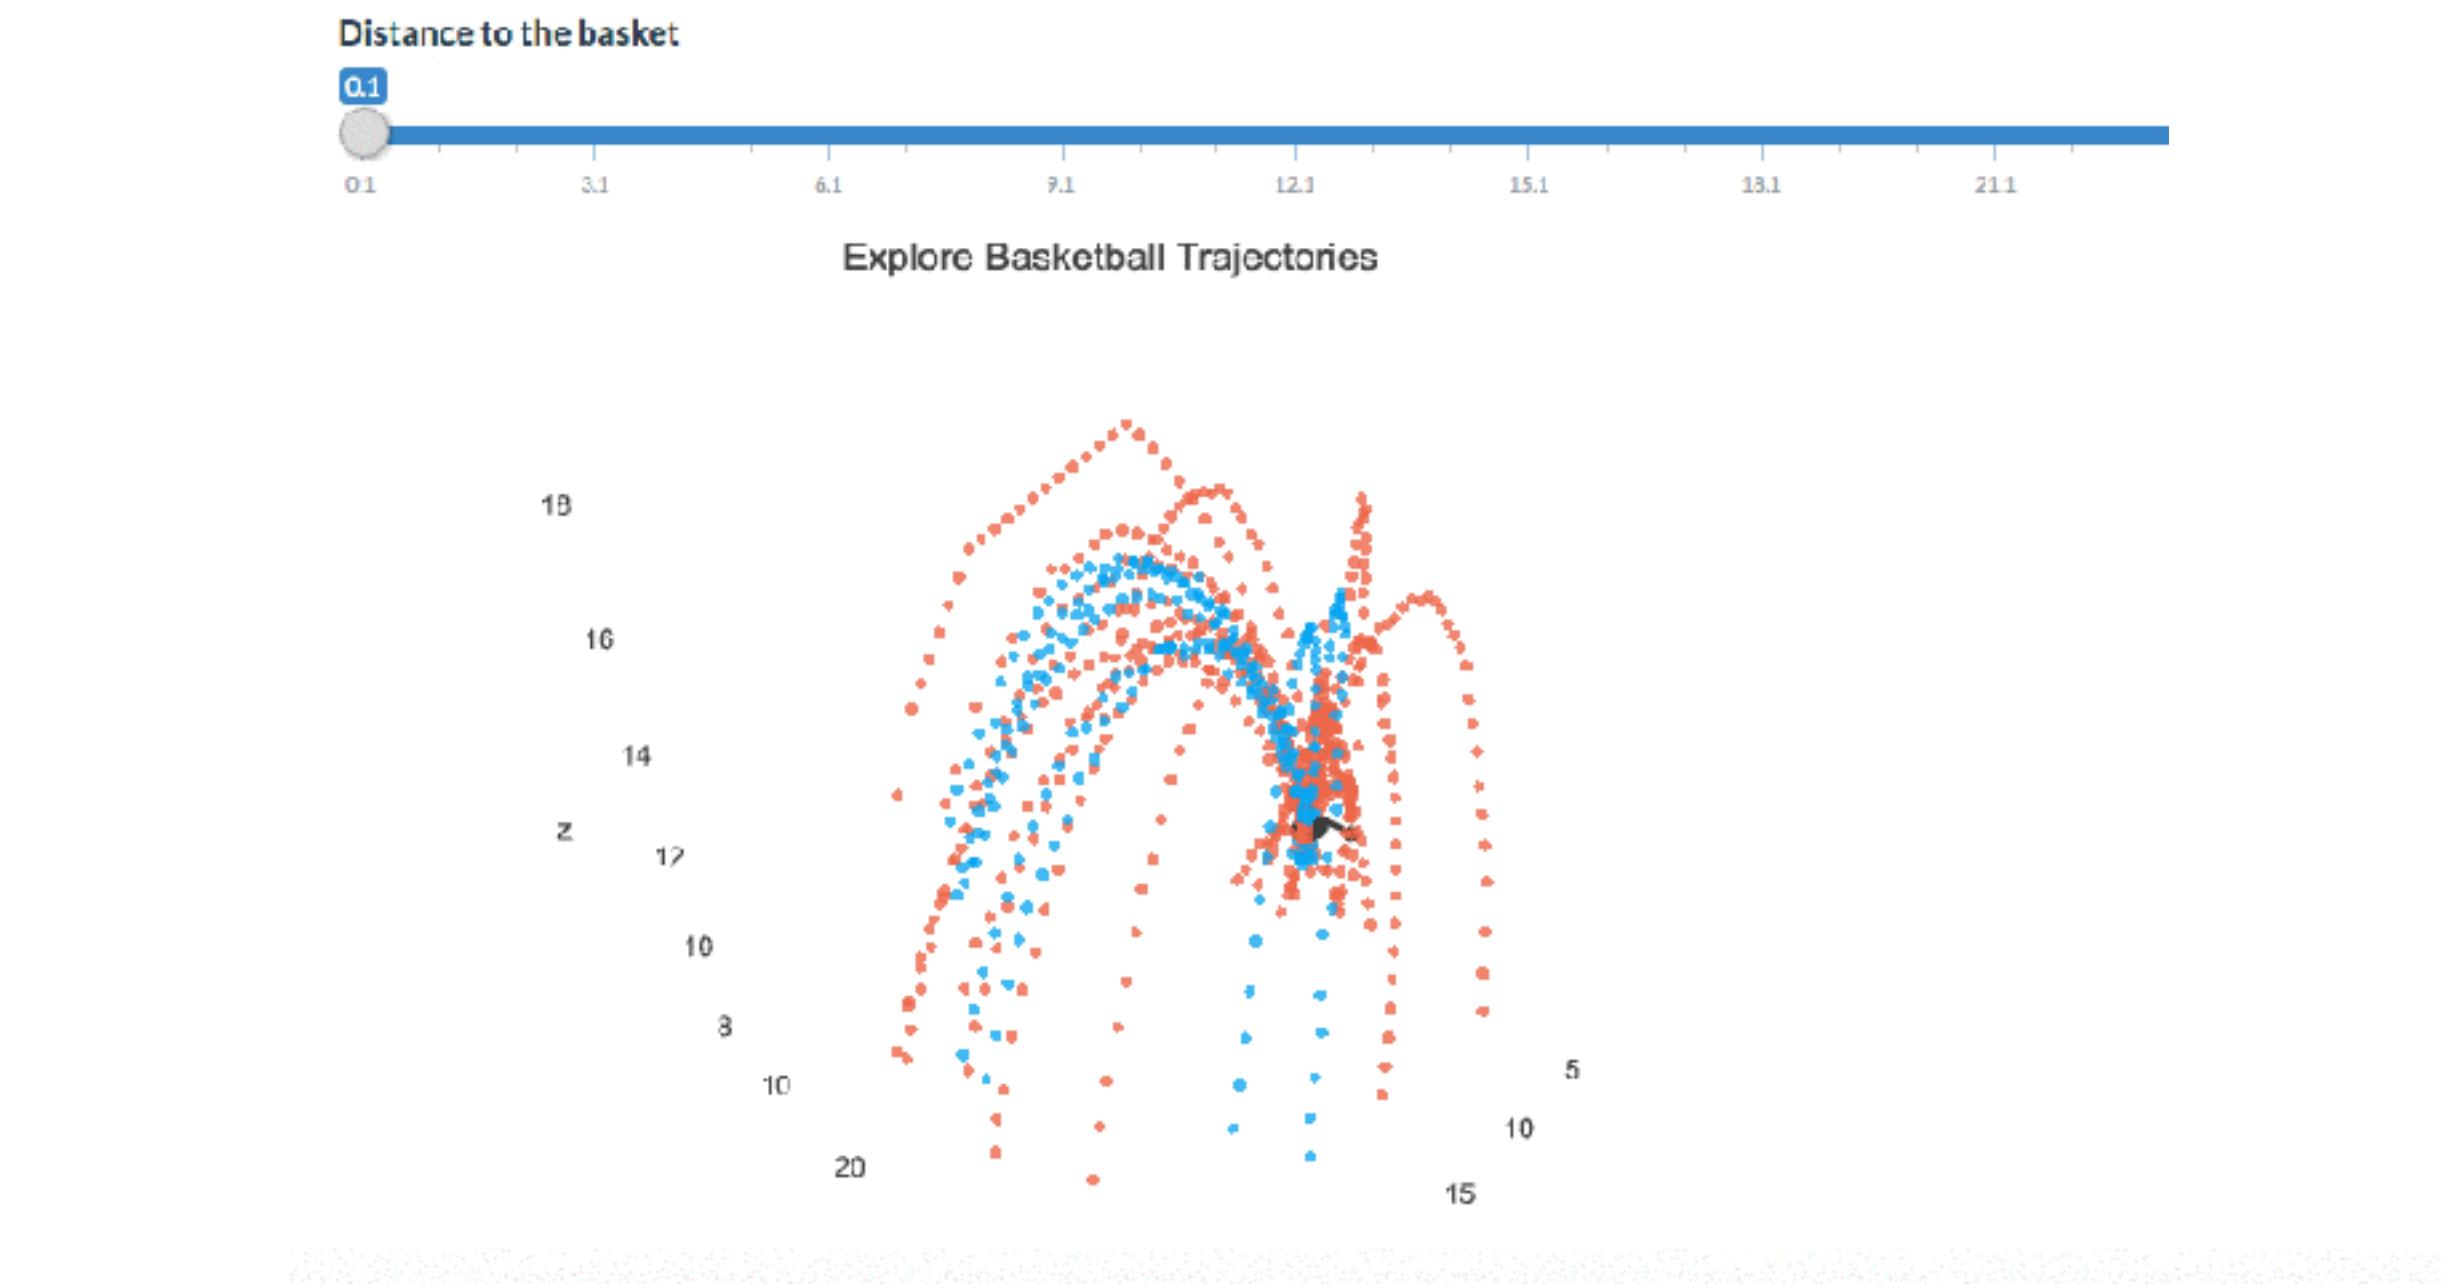
demo: handwriting

Play/Pause Clear Length of prediction 20 Variation¹ 0.1

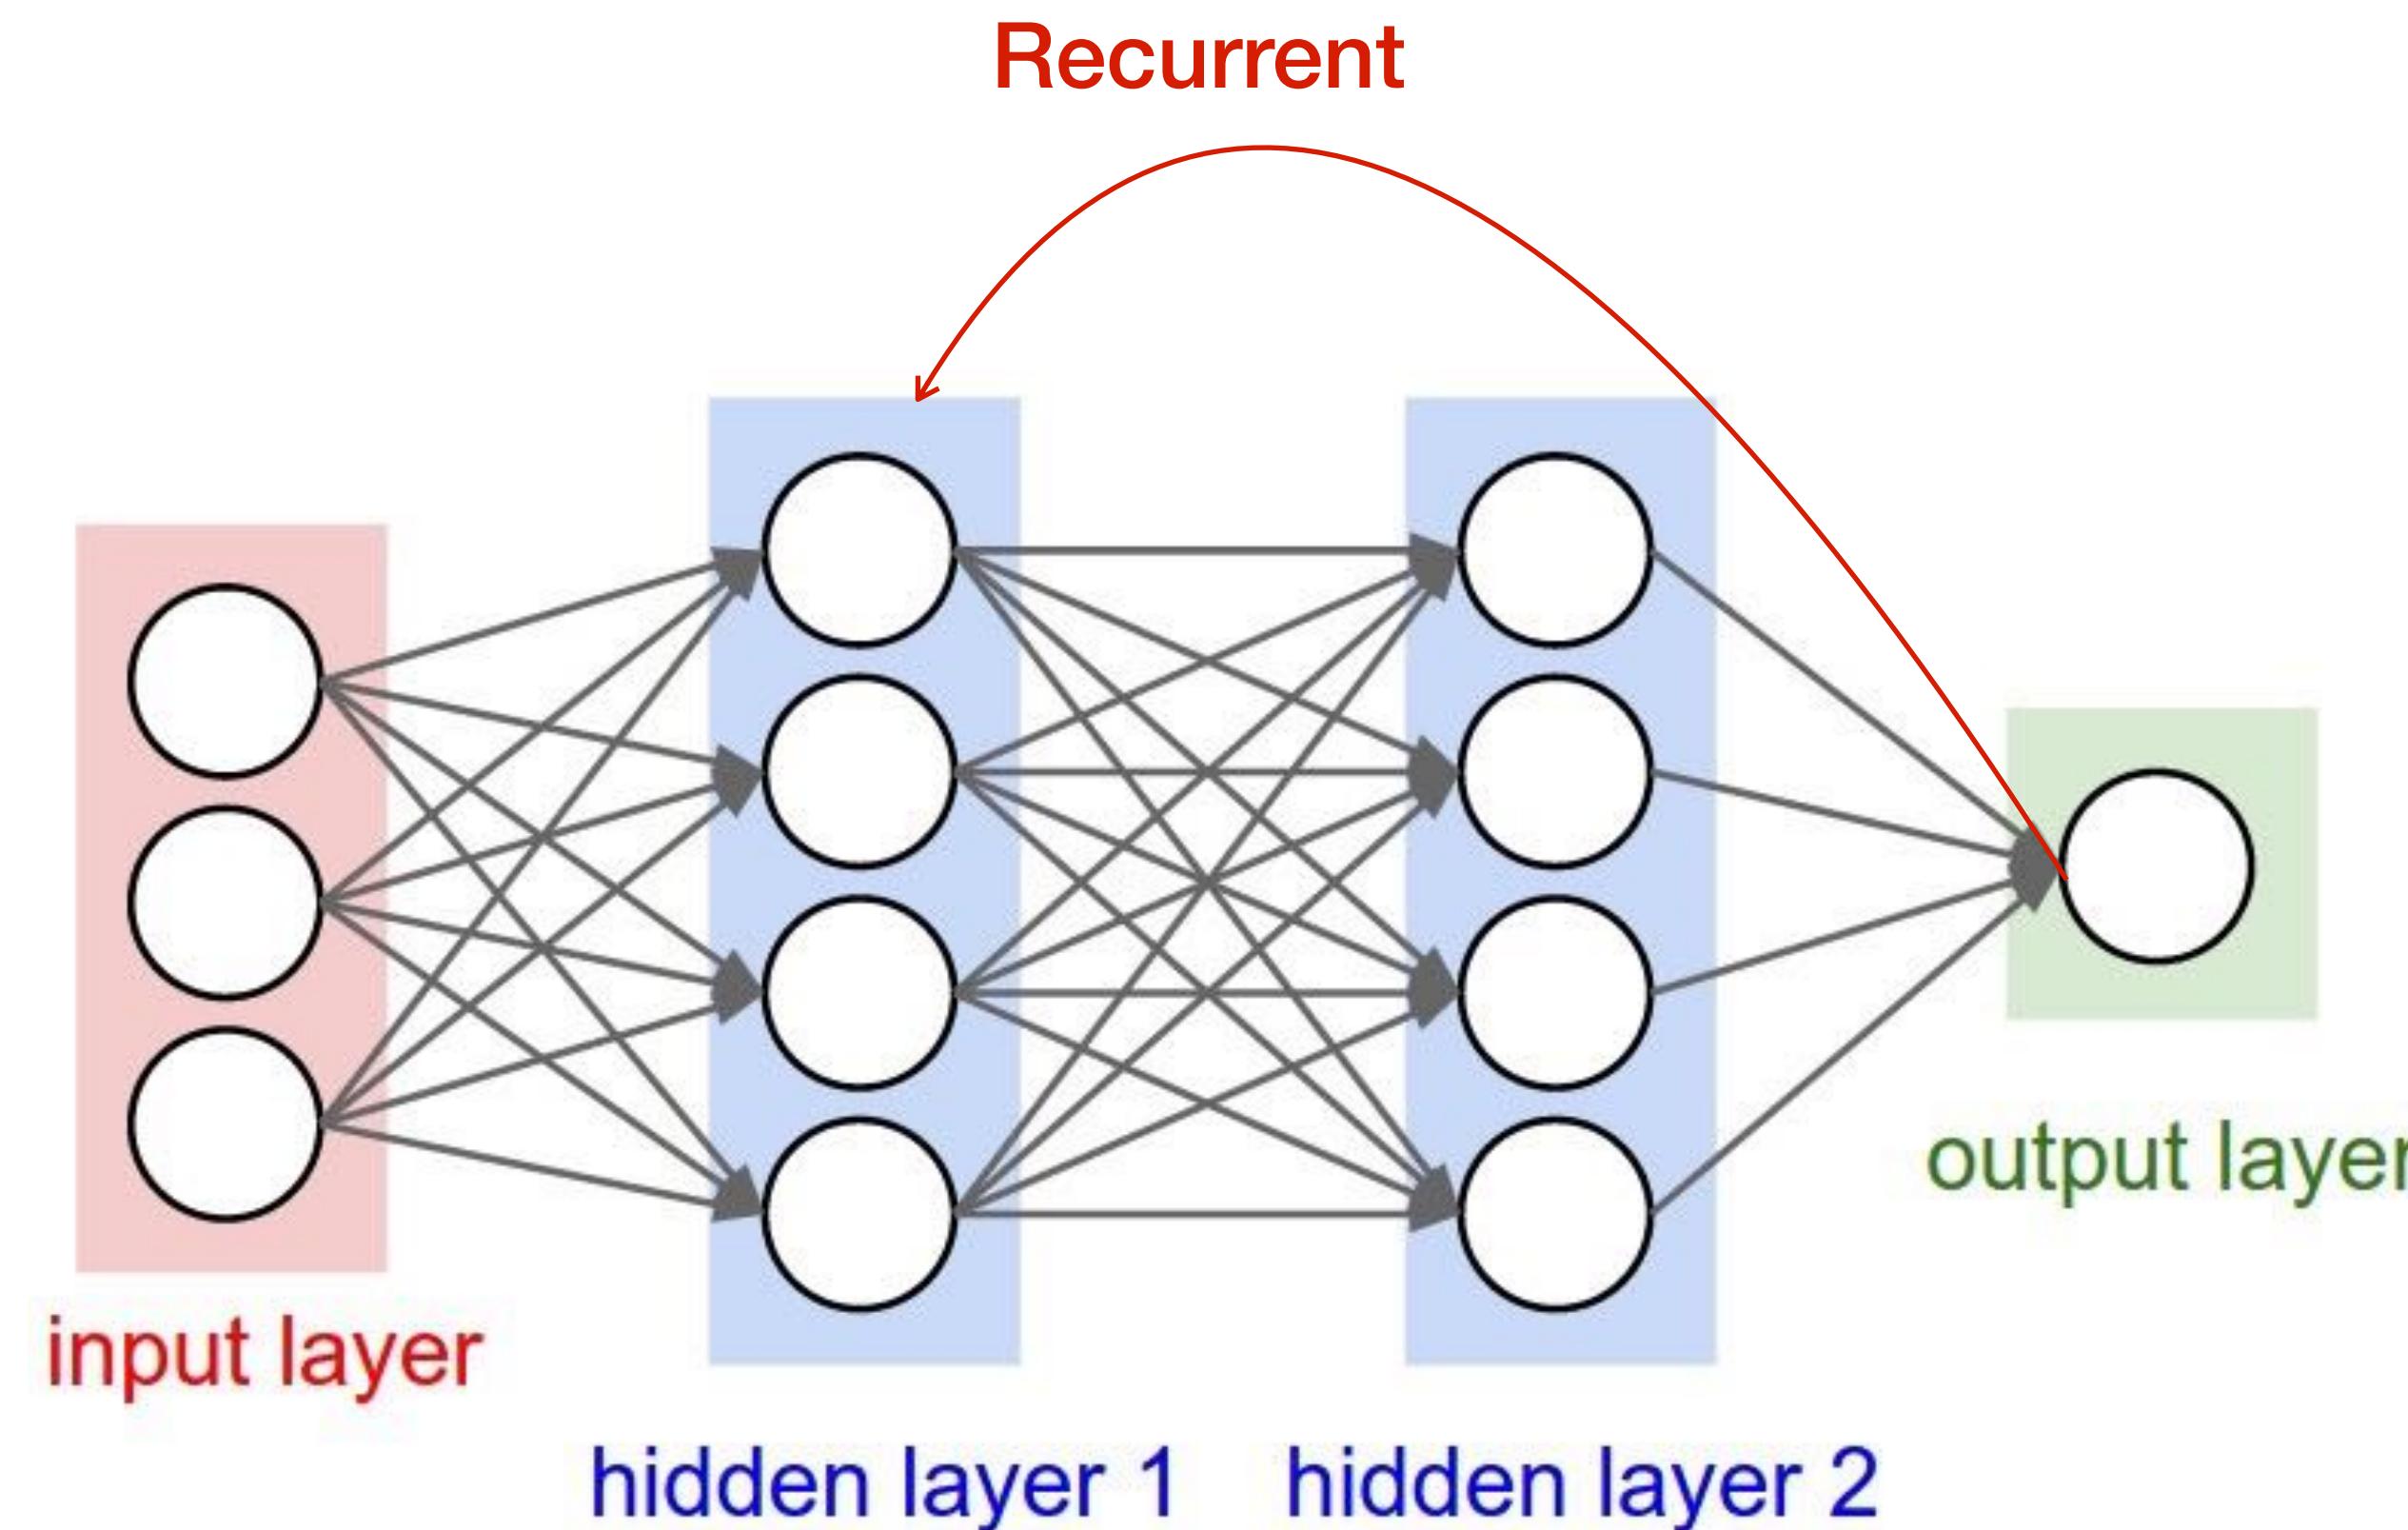
Start writing here and the network will try to generate new strokes in your style

http://distill.pub/2016/handwriting/

demo: basketball trajectories



recurrent neural network



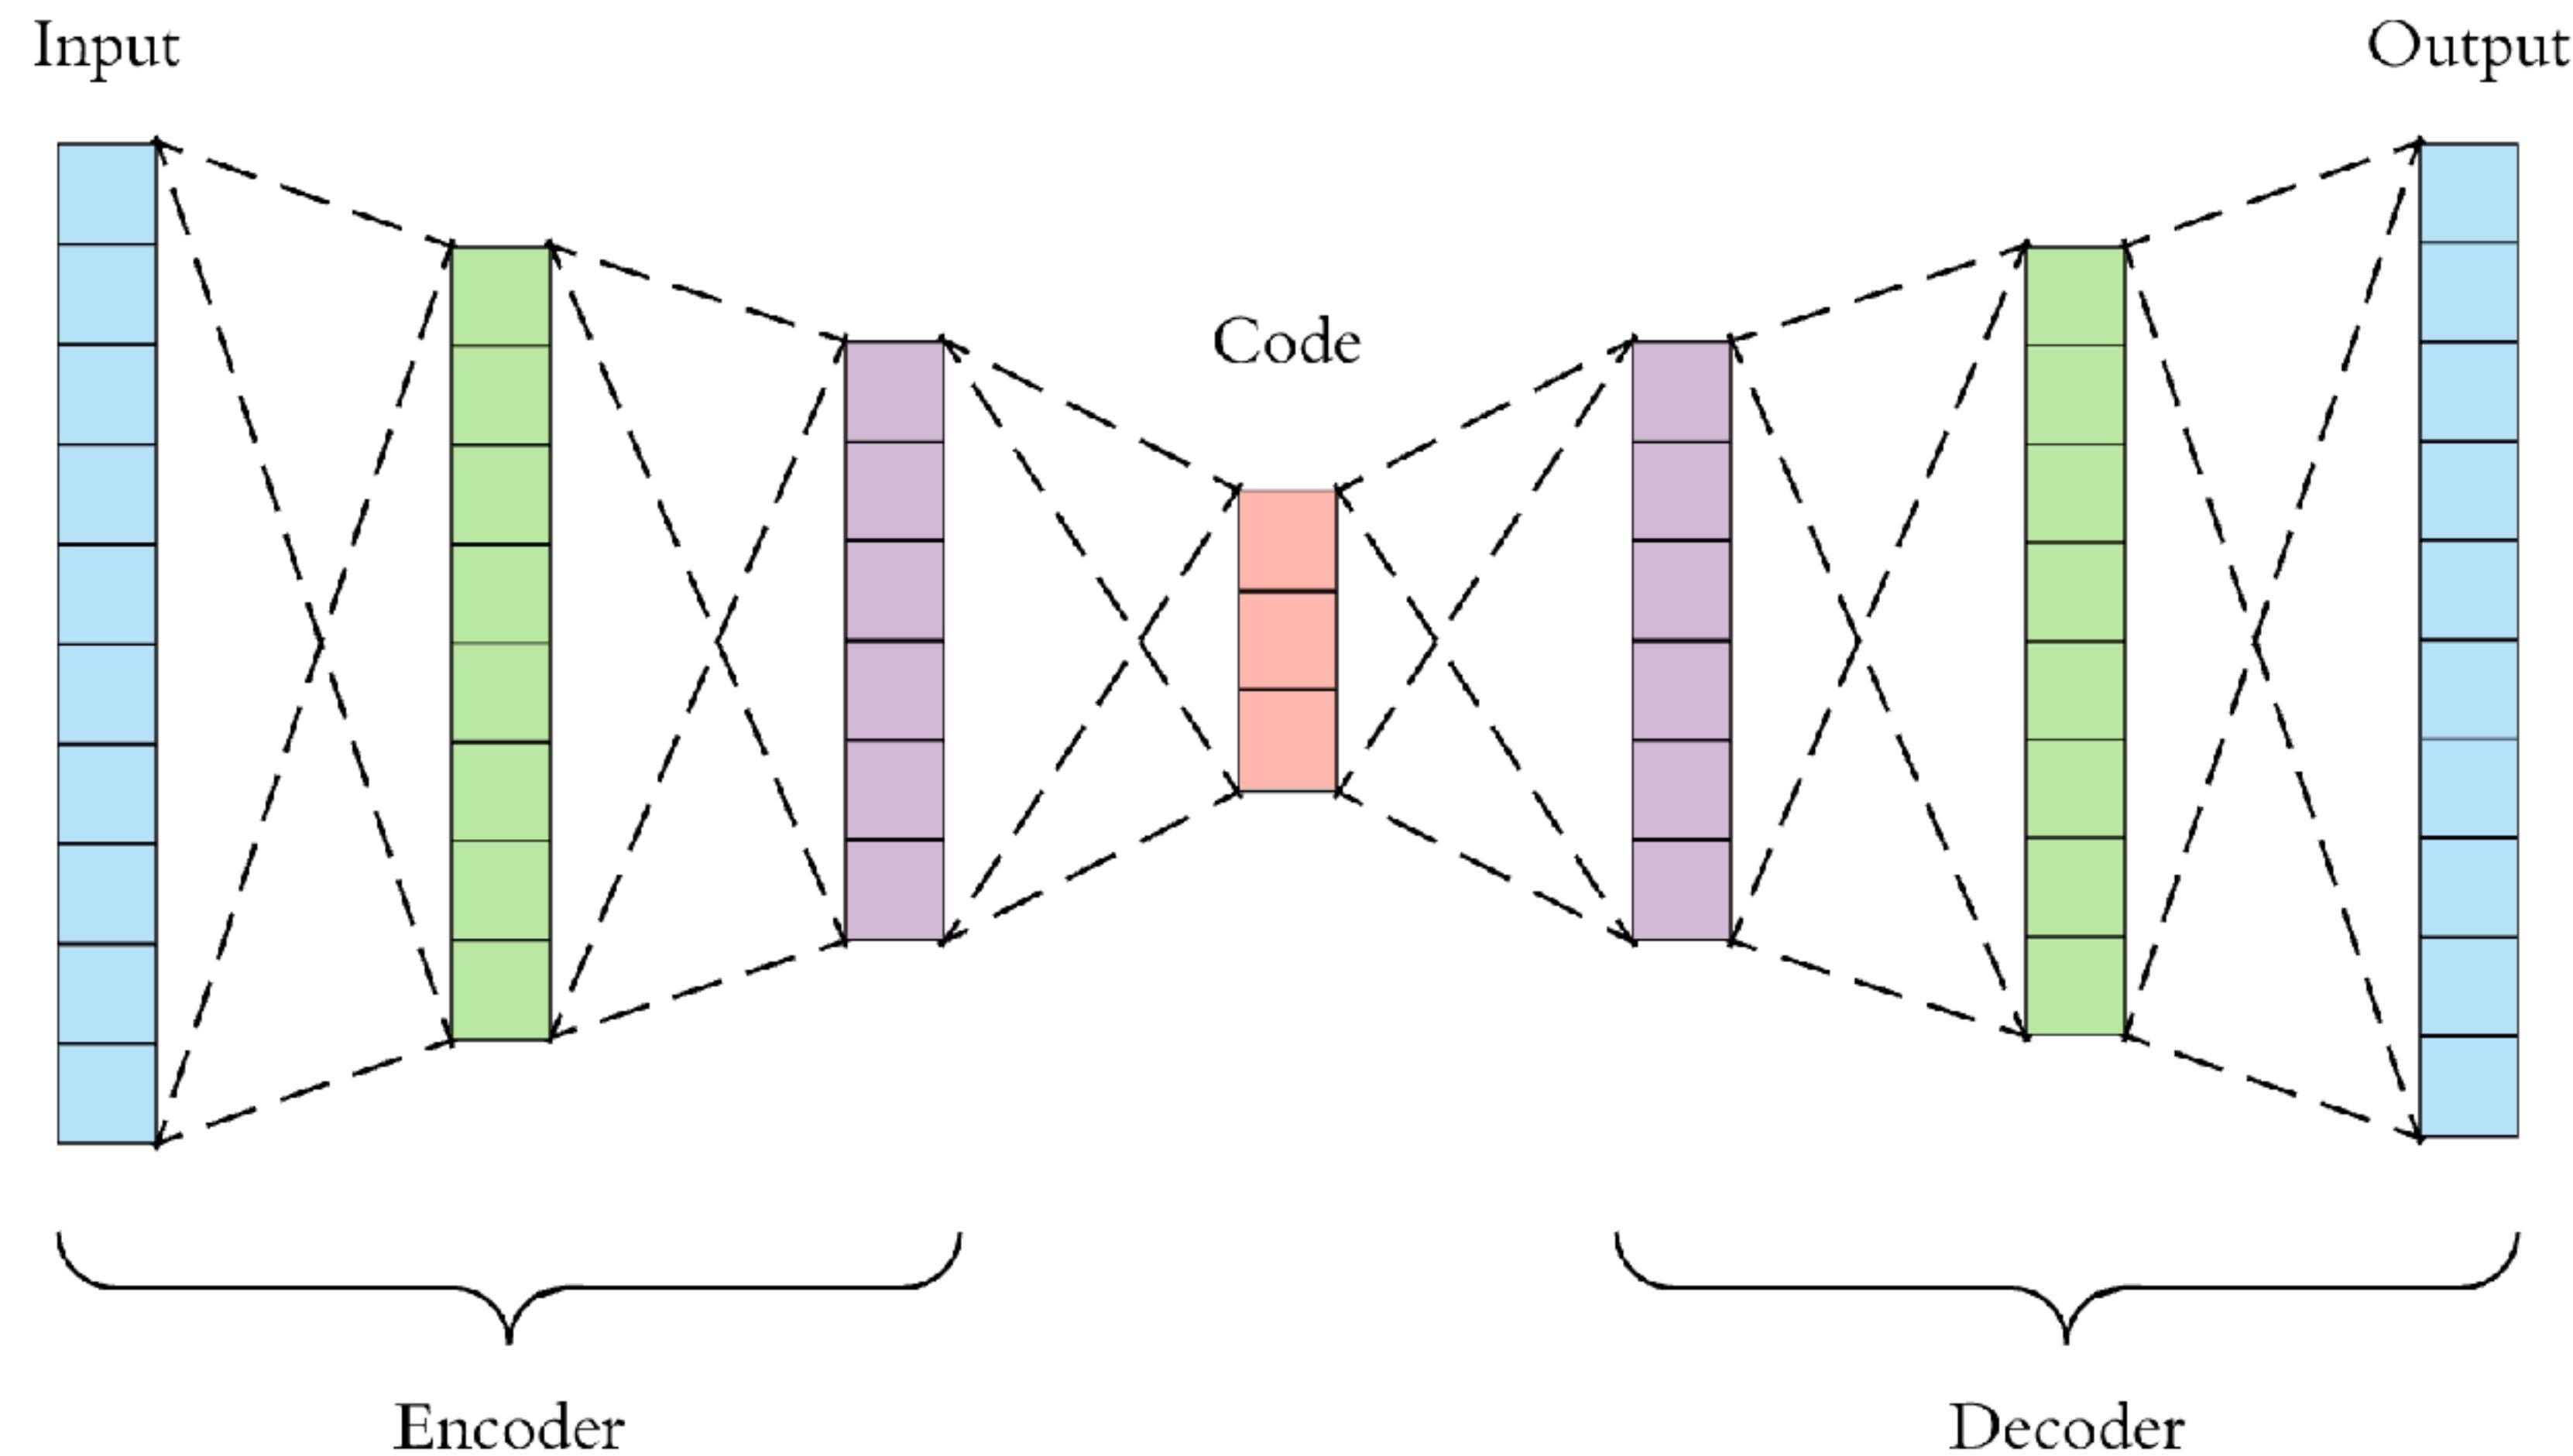
Recurrent neural network

```
model <- keras_model_sequential()
model %>%
  layer_embedding(input_dim = max_features, output_dim =
128, input_length = maxlen) %>%
  bidirectional(layer_lstm(units = 64)) %>%
  layer_dropout(rate = 0.5) %>%
  layer_dense(units = 1, activation = 'sigmoid')
```

Generating new

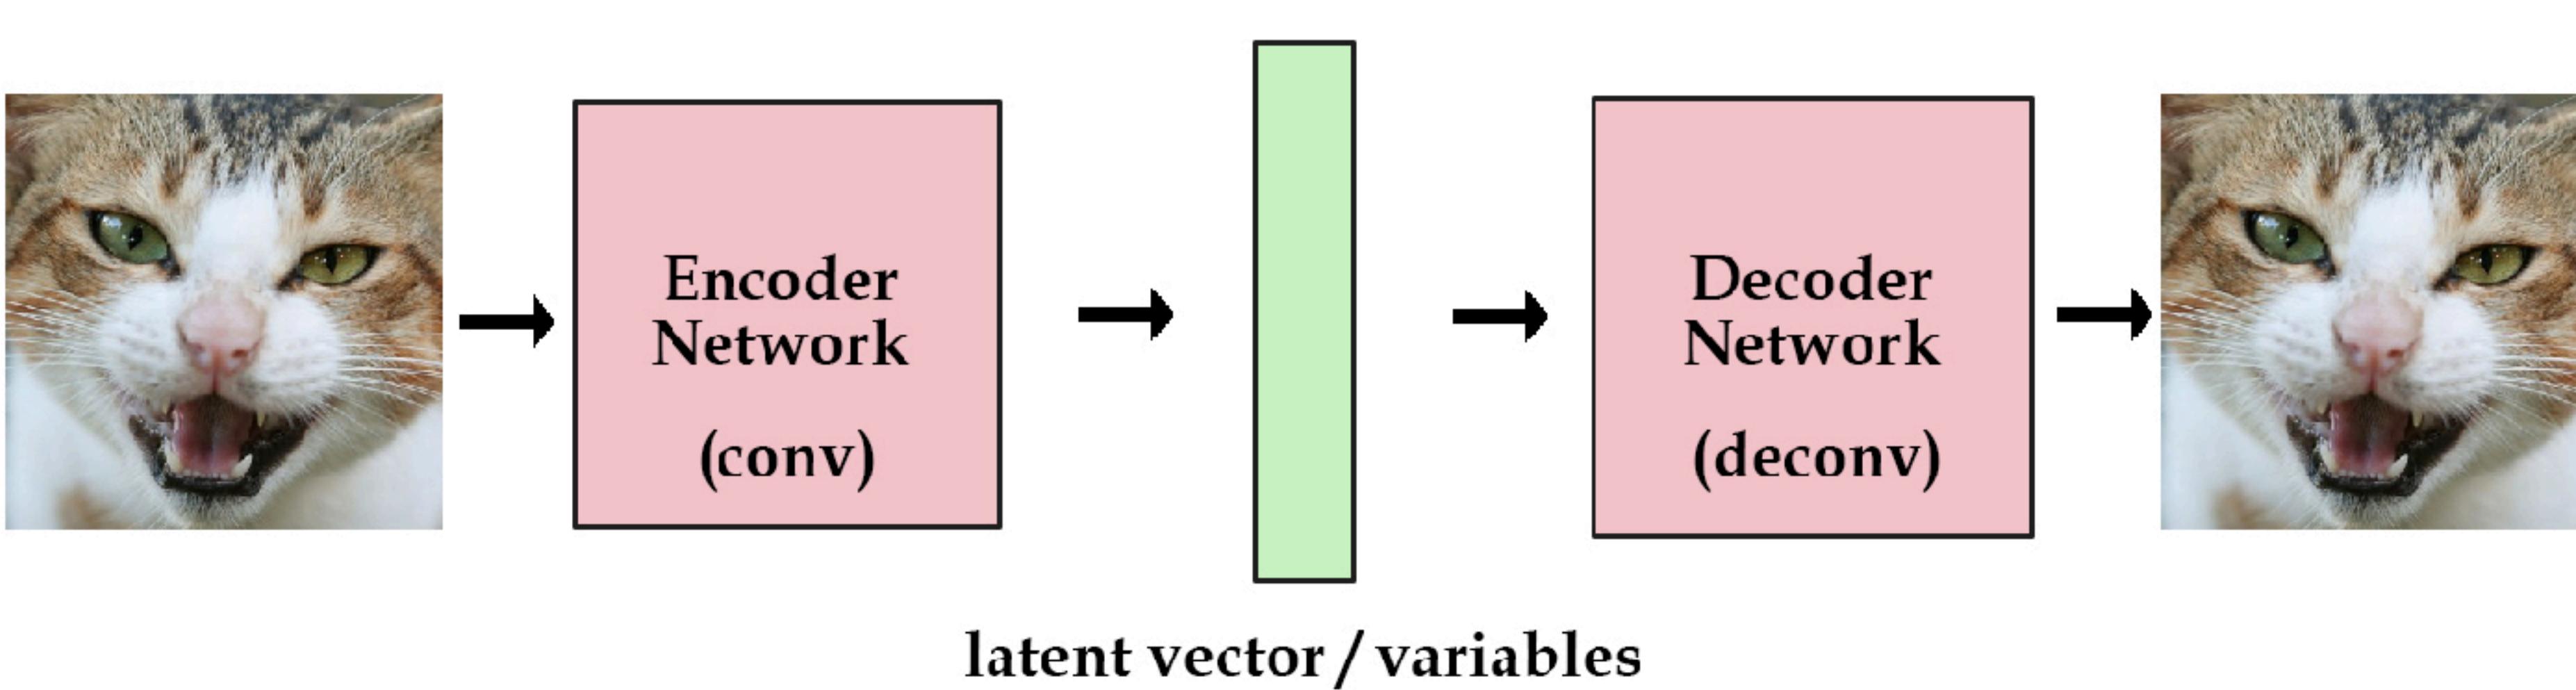


autoencoder



autoencoder

<http://kvfrans.com/variational-autoencoders-explained/>

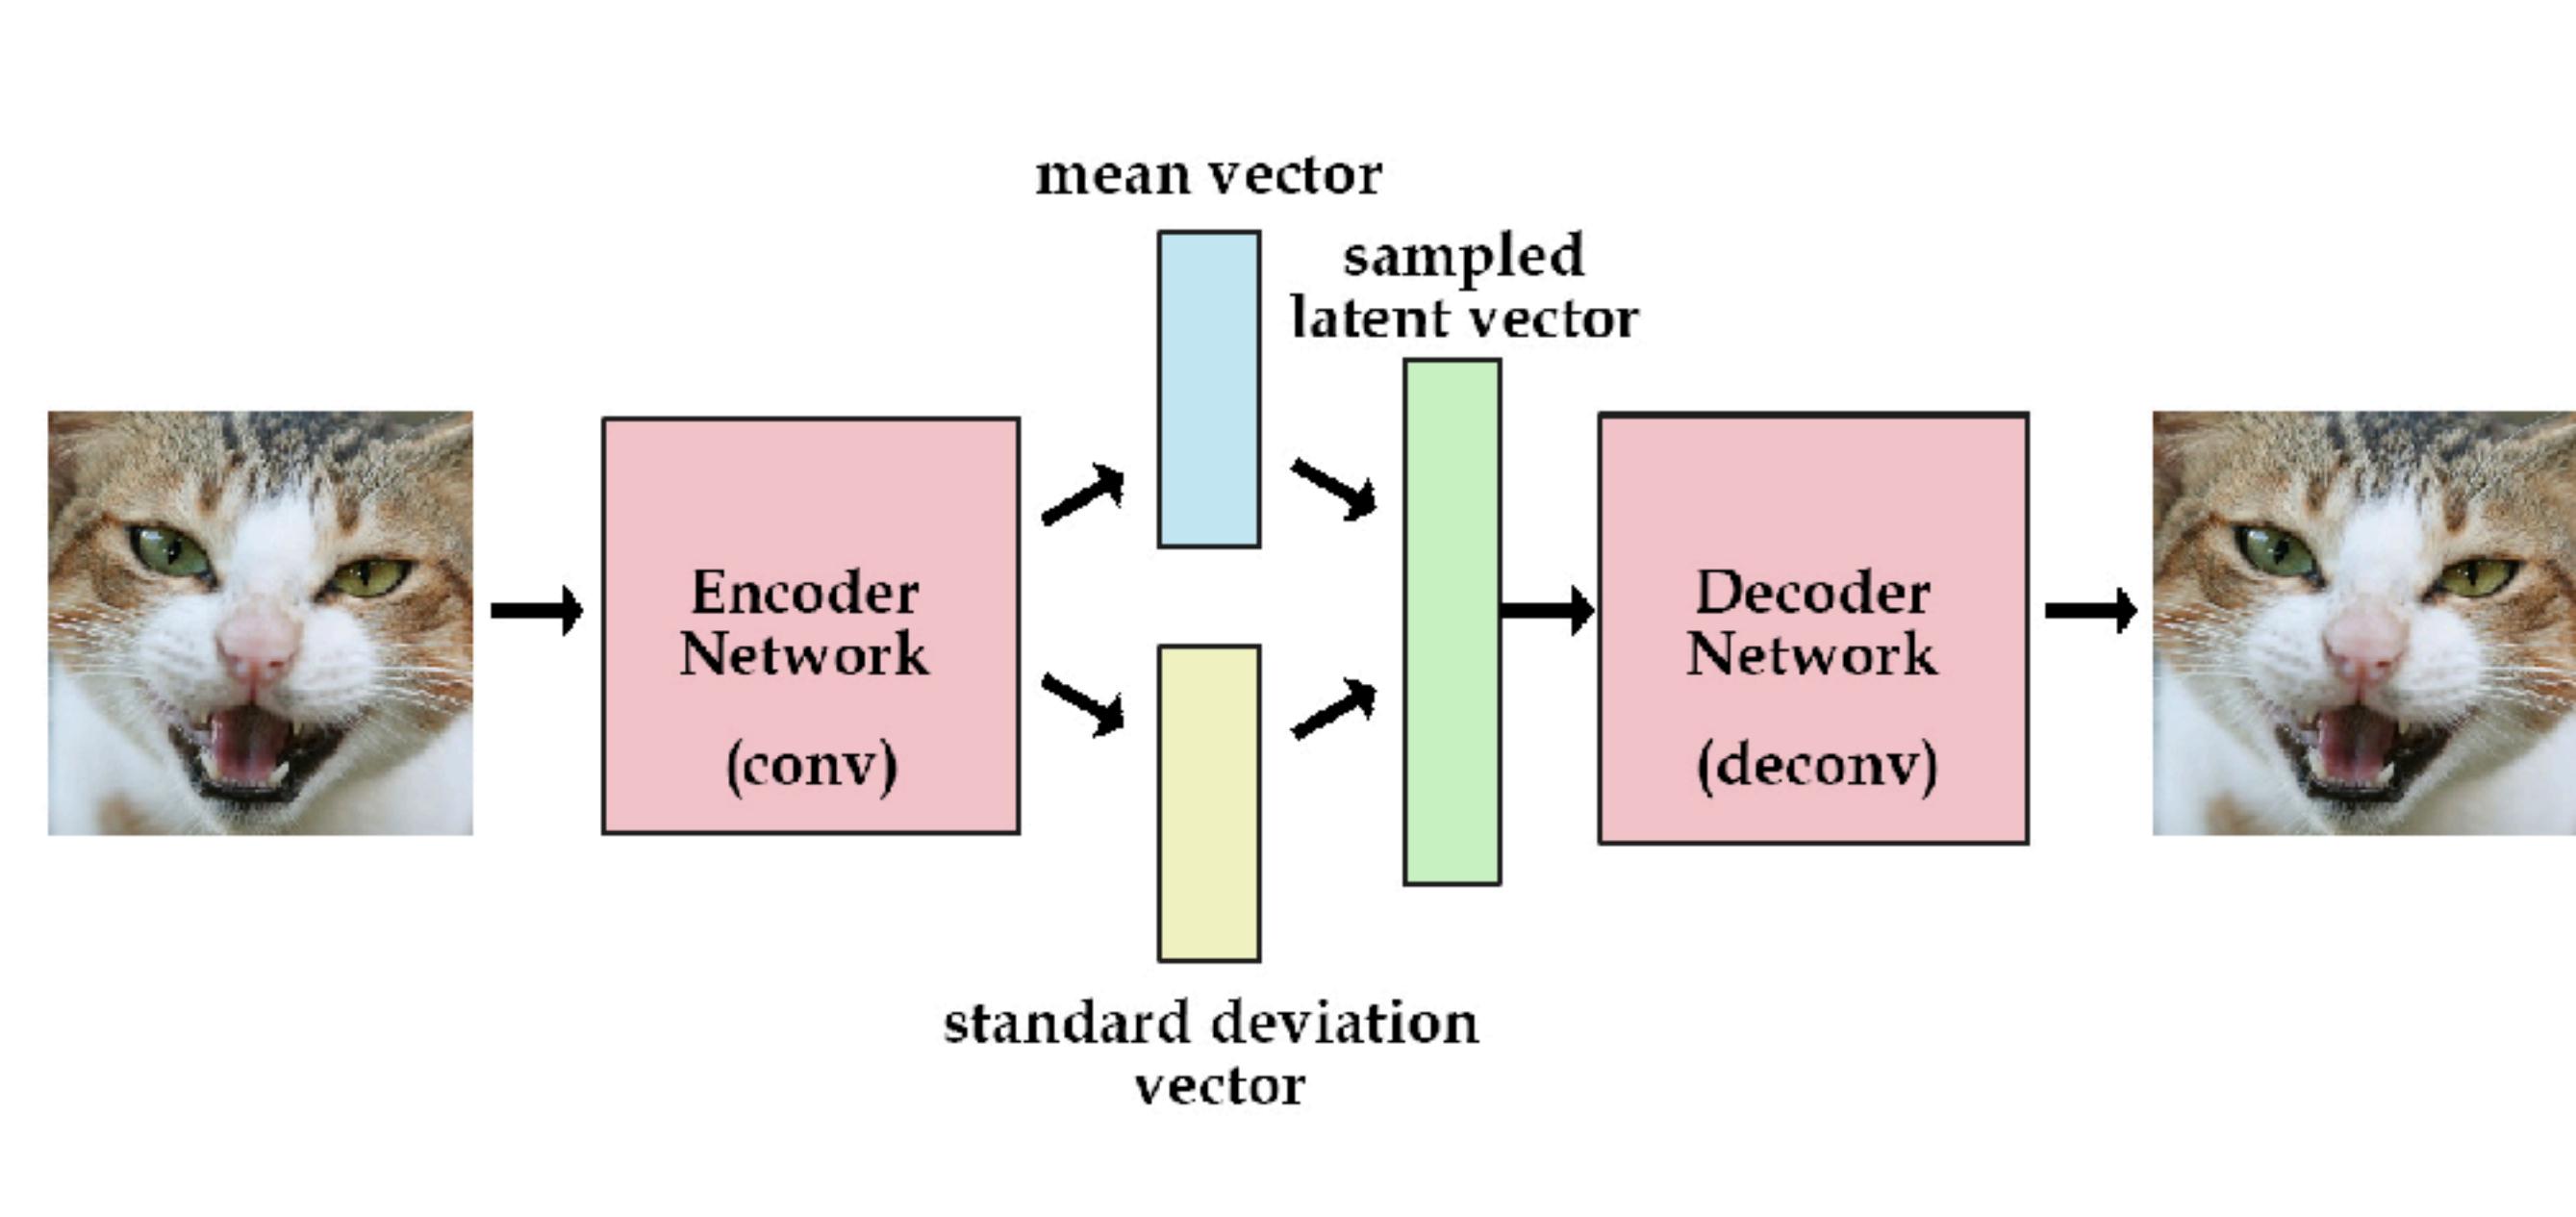


Deep Autoencoder

```
input_img = Input(shape=(784,))
encoded = Dense(128, activation='relu')(input_img)
encoded = Dense(64, activation='relu')(encoded)
encoded = Dense(32, activation='relu')(encoded)

decoded = Dense(64, activation='relu')(encoded)
decoded = Dense(128, activation='relu')(decoded)
decoded = Dense(784, activation='sigmoid')(decoded)
```

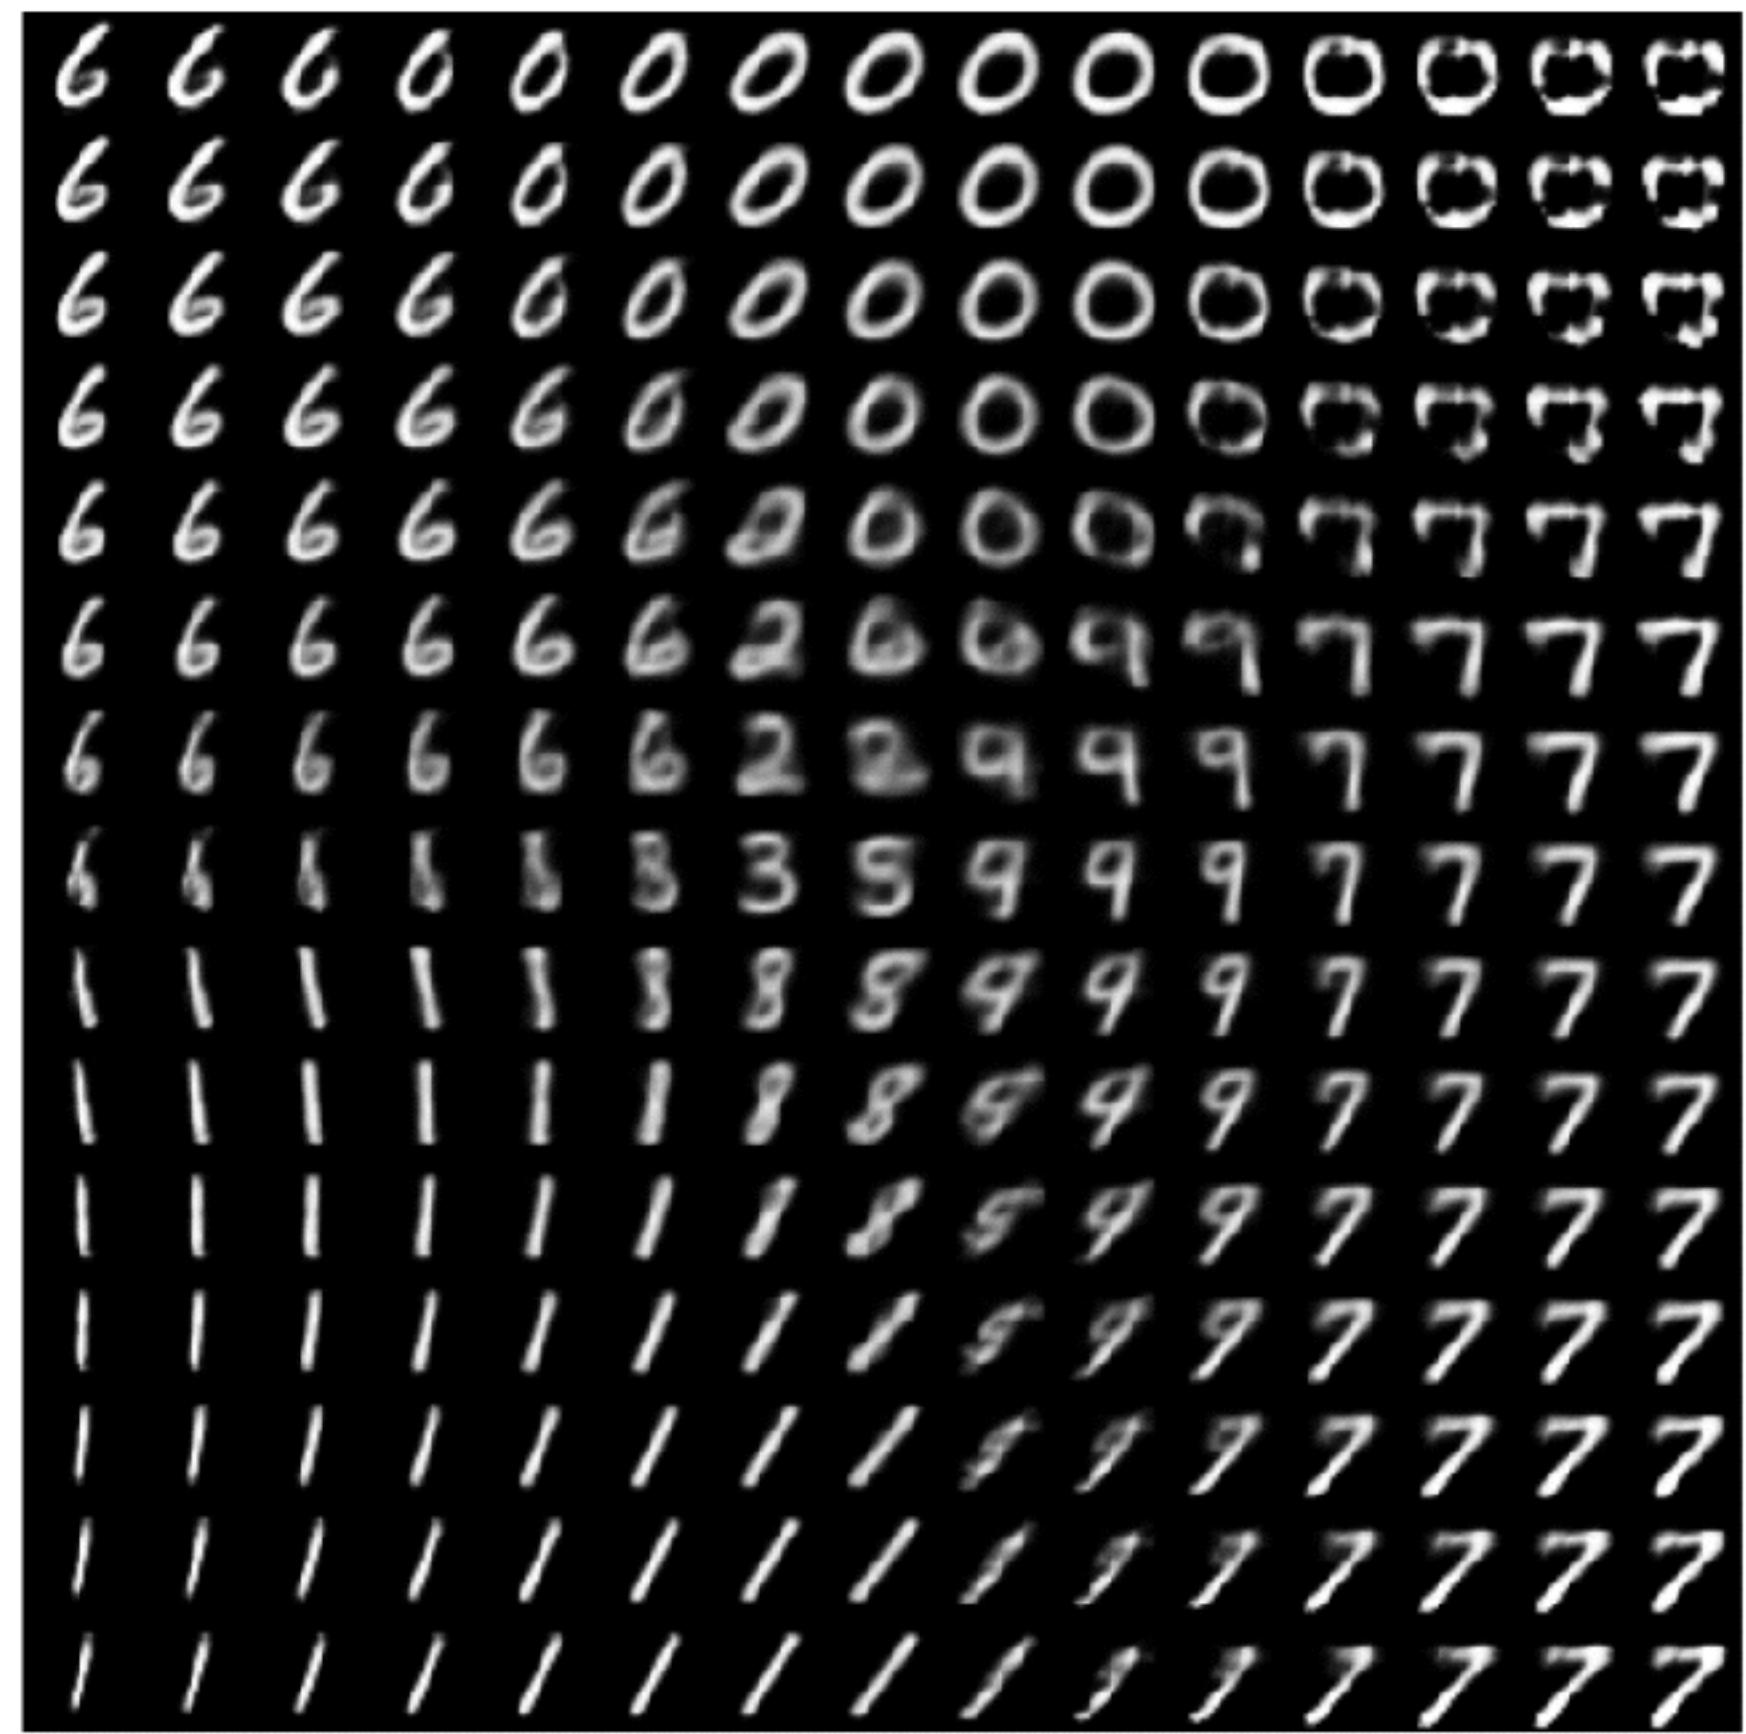
variational autoencoder



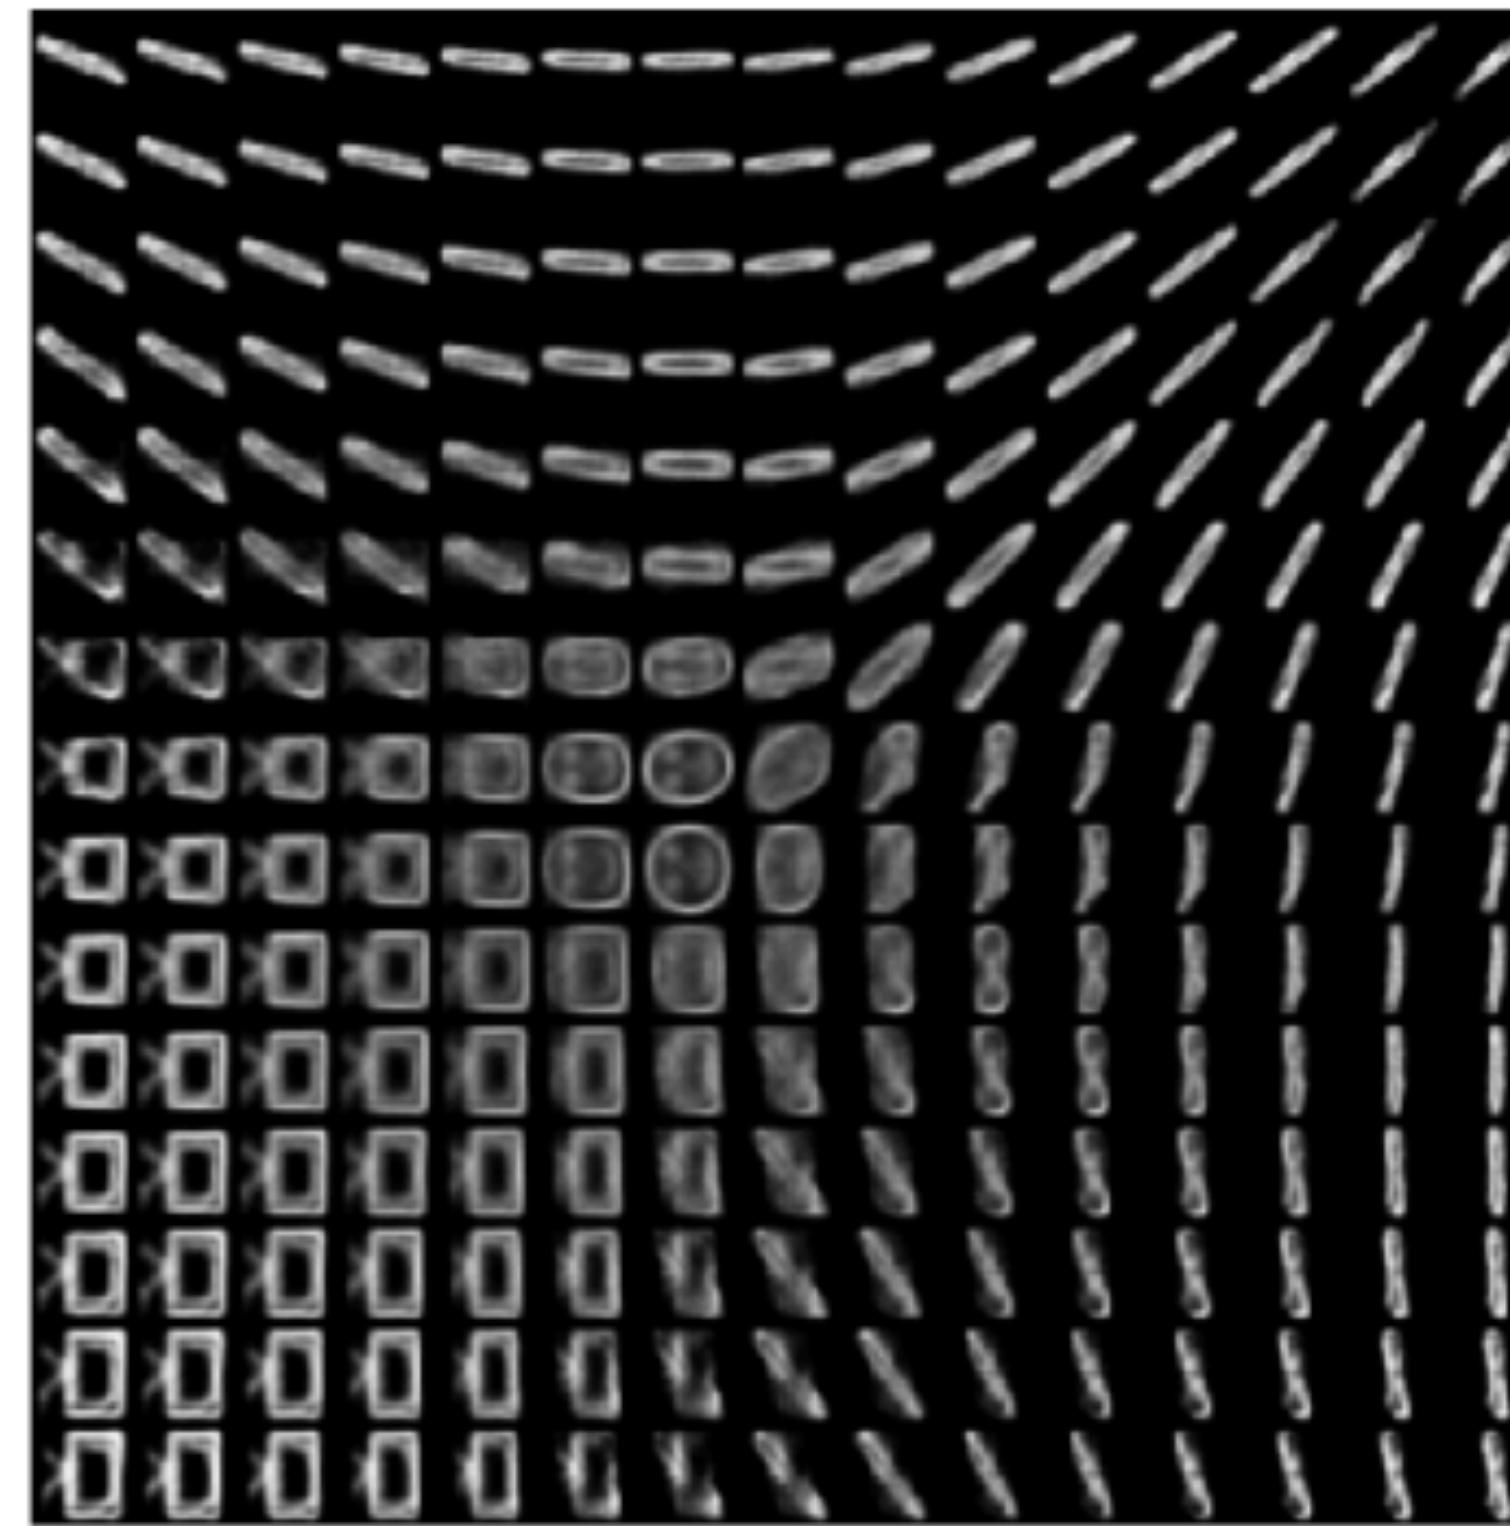
VAE loss

```
vae_loss <- function(x, x_decoded_mean){  
  xent_loss <- (original_dim/1.0)*loss_binary_crossentropy(x, x_decoded_mean)  
  kl_loss <- -0.5*K$mean(1 + z_log_var - K$square(z_mean) - K$exp(z_log_var),  
  axis = -1L)  
  xent_loss + kl_loss  
}
```

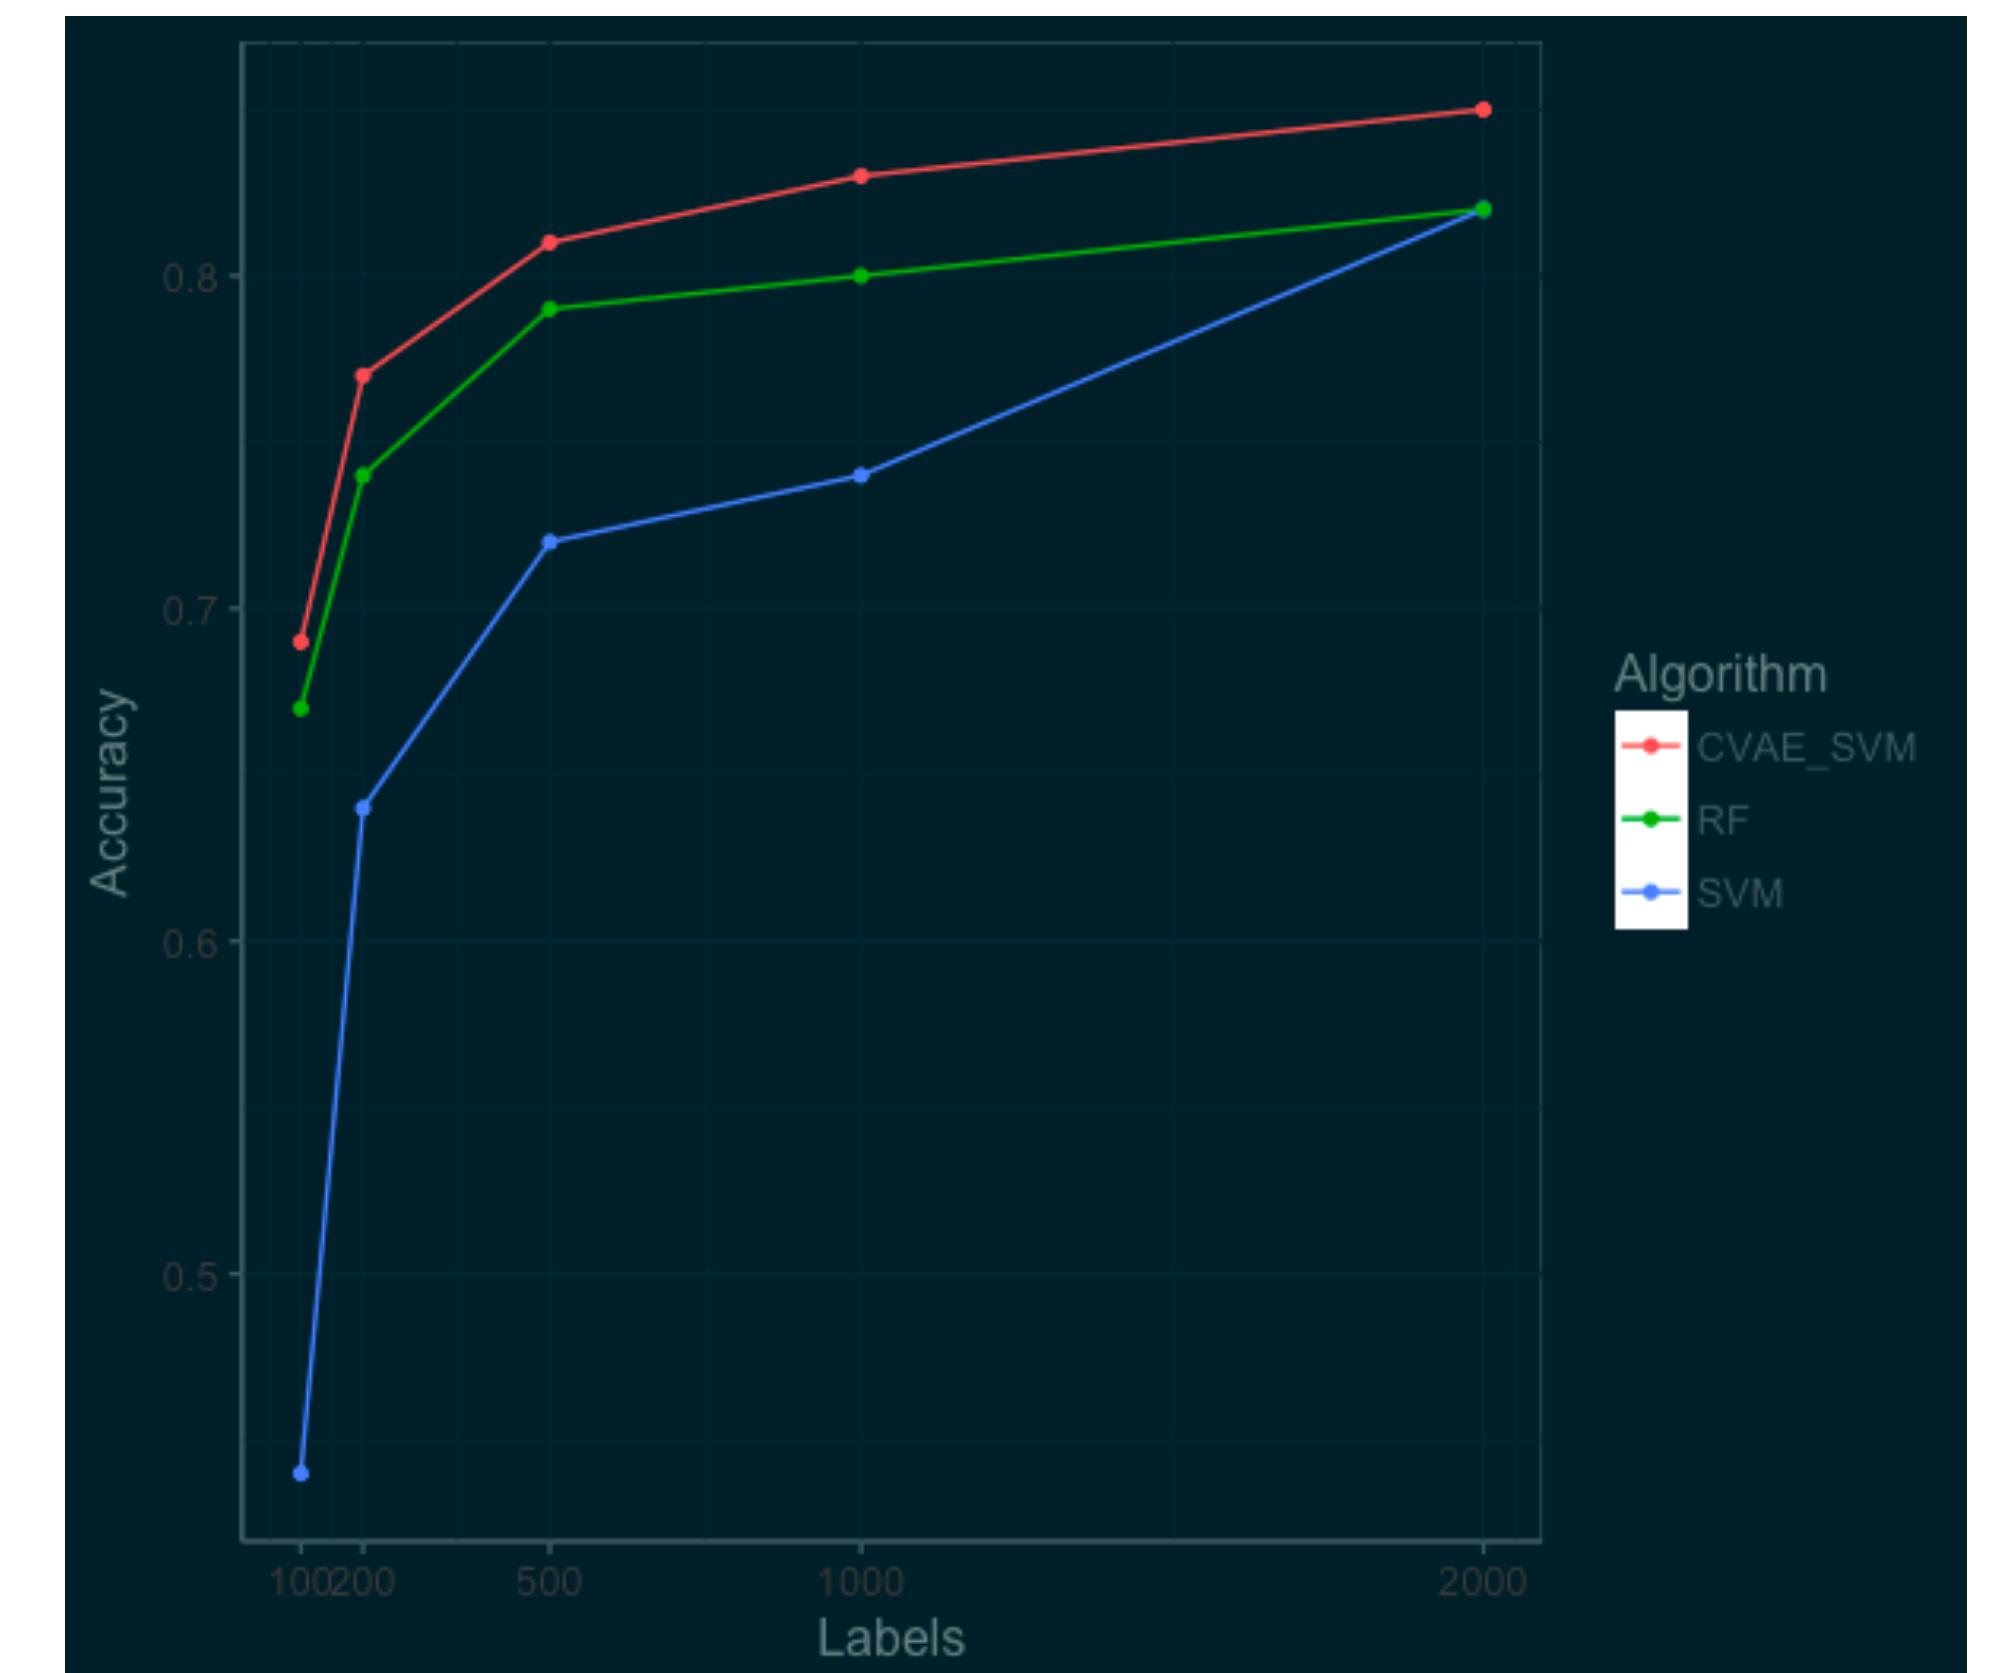
MNIST



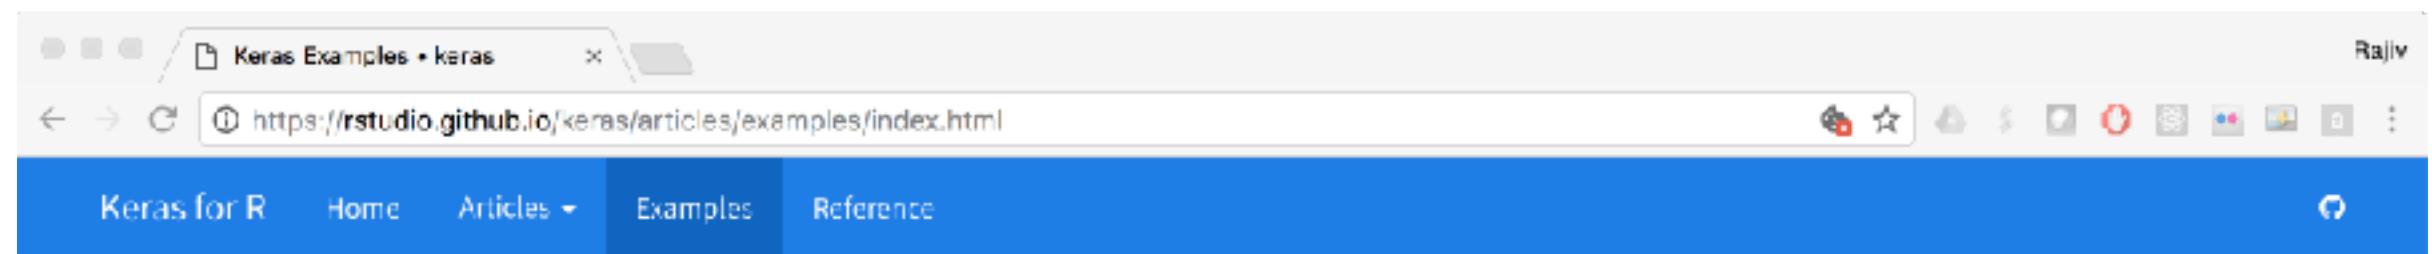
QuickDraw



Semi supervised learning



More examples . . .



Keras Examples

Example	Description
addition_rnn.R	Implementation of sequence to sequence learning for performing addition of two numbers (as strings).
babl_memnn.R	Trains a memory network on the bAbI dataset for reading comprehension.
babl_rnn.R	Trains a two-branch recurrent network on the bAbI dataset for reading comprehension.
cifar10_cnn.R	Trains a simple deep CNN on the CIFAR10 small images dataset.
deep_dream.R	Deep Dreams in Keras.
imdb_bidirectional_lstm.R	Trains a Bidirectional LSTM on the IMDB sentiment classification task.
imdb_cnn.R	Demonstrates the use of Convolution1D for text classification.
imdb_cnn_lstm.R	Trains a convolutional stack followed by a recurrent stack network on the IMDB sentiment classification task.
imdb_fasttext.R	Trains a FastText model on the IMDB sentiment classification task.
imdb_lstm.R	Trains a LSTM on the IMDB sentiment classification task.
lstm_text_generation.R	Generates text from Nietzsche's writings.
mnist_conv.R	Trains a simple convnet on the MNIST dataset.
mnist_irnn.R	Reproduction of the IRNN experiment with pixel-by-pixel sequential MNIST in "A Simple Way to Initialize Recurrent Networks of Rectified Linear Units" by Le et al.
mnist_mlp.R	Trains a simple deep multi-layer perceptron on the MNIST dataset.
mnist_transfer_cnn.R	Transfer learning toy example.
reuters_mlp.R	Trains and evaluates a simple MLP on the Reuters newswire topic classification task.
stateful_lstm.R	Demonstrates how to use stateful RNNs to model long sequences efficiently.

whew . . .

The development of Keras/R

How to use Keras in R for:
convolutional neural networks
recurrent neural networks
variational autoencoders

Deep Learning with Keras and R

Feb 15, 2018

RajivShah.com

rshah@pobox.com

 github.com/rajshah4/image_keras

 rajcs4